

Chapter 5: Advanced SQL

Outline

- n Accessing SQL From a Programming Language
- n Functions and Procedures
- n *Triggers
- n **Recursive Queries
- n **Advanced Aggregation Features

Accessing SQL from a Programming Language

A database programmer must have access to a general-purpose programming language for at least two reasons

- n Not all queries can be expressed in SQL, since SQL does not provide the full expressive power of a general-purpose language.
- n Non-declarative actions -- such as printing a report, interacting with a user, or sending the results of a query to a graphical user interface -- cannot be done from within SQL.

Accessing SQL from a Programming Language (Cont.)

There are two approaches to accessing database from a general-purpose programming language

- n **API** (Application Program Interface) - - A general-purpose program can connect to and communicate with a database server using a collection of functions.
- n **Embedded SQL** -- provides a means by which a program can interact with a database server.
 - | The SQL statements are translated at compile time into function calls.
 - | At runtime, these function calls connect to the database using an API that provides dynamic SQL facilities.

JDBC and ODBC

- n **API (application-program interface)** for a program to interact with a database server
- n Application makes calls to
 - | Connect with the database server
 - | Send SQL commands to the database server
 - | Fetch tuples of result one-by-one into program variables
- n **ODBC** (Open Database Connectivity) works with C, C++, C#
- n **JDBC** (Java Database Connectivity) works with Java
- n **Embedded SQL** in C
- n **SQLJ** - embedded SQL in Java
- n **JPA(Java Persistence API)** - OR mapping of Java

JDBC

- n **JDBC** is a Java API for communicating with database systems supporting SQL.
- n JDBC supports a variety of features for querying and updating data, and for retrieving query results.
- n JDBC also supports **metadata retrieval**, such as querying about relations present in the database and the names and types of relation attributes.
- n Model for communicating with the database:
 - | **Open a connection**
 - | **Create a “statement” object**
 - | **Execute queries** using the Statement object to send queries and fetch results
 - | **Exception mechanism** to handle errors

JDBC Code

```
public static void JDBCexample(String dbid, String userid, String passwd)
{
    try {
        Connection conn = DriverManager.getConnection(
            "jdbc:oracle:thin:@db.yale.edu:2000:univdb", userid, passwd);
        Statement stmt = conn.createStatement();
        ... Do Actual Work ....
        stmt.close();
        conn.close();
    }
    catch (SQLException sqle) {
        System.out.println("SQLException : " + sqle);
    }
}
```

JDBC Code (Cont.)

n Update to database

```
try {  
    stmt.executeUpdate(  
        "insert into instructor values('77987', 'Kim', 'Physics', 98000)");  
} catch (SQLException sqle)  
{  
    System.out.println("Could not insert tuple. " + sqle);  
}
```

n Execute query and fetch and print results

```
ResultSet rset = stmt.executeQuery(  
    "select dept_name, avg (salary)  
    from instructor  
    group by dept_name");  
while (rset.next()) {  
    System.out.println(rset.getString("dept_name") + " " +  
        rset.getFloat(2));  
}
```


JDBC Code Details

n Getting result fields:

| **rset.getString("dept_name")** and **rset.getString(1)**
equivalent if dept_name is the first argument of select
result.

n Dealing with Null values

| **int a = rset.getInt("a");**
if (rset.isNull()) Systems.out.println("Got null value");

Prepared Statement

- n `PreparedStatement pStmt = conn.prepareStatement(
 "insert into instructor values(?,?,?,?)");
pStmt.setString(1, "88877"); pStmt.setString(2, "Perry");
pStmt.setString(3, "Finance"); pStmt.setInt(4, 125000);
pStmt.executeUpdate();
pStmt.setString(1, "88878");
pStmt.executeUpdate();`
- n WARNING: always use prepared statements when taking an input from the user and adding it to a query
 - | **NEVER create a query by concatenating strings which you get as inputs**
 - | `"insert into instructor values(' " + ID + " ', ' " + name + " ', " +
 " ' " + dept name + " ', " + salary + ")"`
 - | What if name is "D'Souza"?

SQL Injection(SQL 注入)

- n Suppose query is constructed using
 - | "select * from instructor where name = '" + name + "'"
- n Suppose the user, instead of entering a name, enters:
 - | X' or 'Y' = 'Y'
- n then the resulting string of the statement becomes:
 - | "select * from instructor where name = '" + "X' or 'Y' = 'Y'" + "'"
 - | which is:
 - ▶ select * from instructor where name = 'X' or 'Y' = 'Y'
 - | User could have even used
 - ▶ X'; update instructor set salary = salary + 10000;
 - then →
 - ▶ select * from instructor where name = 'X';
 - update instructor set salary = salary + 10000;
- n **Always use prepared statements, with user inputs as parameters**

Metadata Features

- n **ResultSet metadata**

- n E.g., after executing query to get a **ResultSet** rs:

- | **ResultSetMetaData** rsmd = rs.**getMetaData()**;
for(int i = 1; i <= rsmd.**getColumnCount()**; i++) {
 System.out.println(rsmd.**getColumnName(i)**);
 System.out.println(rsmd.**getColumnTypeName(i)**);
}

Metadata (Cont)

- n **Database metadata**

- n `DatabaseMetaData dbmd = conn.getMetaData();`

- `ResultSet rs = dbmd.getColumns(null, "univdb", "department", "%");`

- `// Arguments to getColumns: Catalog, Schema-pattern, Table-pattern,`
`// and Column-Pattern`

- `// Returns: One row for each column; row has a number of attributes`

- `// such as COLUMN_NAME, TYPE_NAME`

- `while(rs.next()) {`

- `System.out.println(rs.getString("COLUMN_NAME"),`

- `rs.getString("TYPE_NAME");`

- `}`

Transaction Control in JDBC

- n By default, each SQL statement is treated as a separate transaction that is committed automatically
 - | bad idea for transactions with multiple updates
- n Can turn off automatic commit on a connection
 - | `conn.setAutoCommit(false);`
- n Transactions must then be committed or rolled back explicitly
 - | `conn.commit();` or
 - | `conn.rollback();`
- n `conn.setAutoCommit(true)` turns on automatic commit.

Other JDBC Features

- n Calling functions and procedures
 - | `CallableStatement cStmt1 = conn.prepareCall("{? = call some function(?)})");`
 - | `CallableStatement cStmt2 = conn.prepareCall("{call some procedure(?,?)})");`
- n Handling large object types
 - | `getBlob()` and `getClob()` that are similar to the `getString()` method, but return objects of type `Blob` and `Clob`, respectively
 - | get data from these objects by `getBytes()`
 - | associate an open stream with Java `Blob` or `Clob` object to update large objects
 - ▶ `blob.setBlob(int parameterIndex, InputStream inputStream).`

JDBC Resources

- n JDBC Basics Tutorial

- | <https://docs.oracle.com/javase/tutorial/jdbc/index.html>

SQLJ

- n JDBC is overly dynamic, errors cannot be caught by compiler
- n SQLJ: embedded SQL in Java
 - | `#sql iterator deptInfolter (String dept name, int avgSal);`
`deptInfolter iter = null;`
`#sql iter = { select dept_name, avg(salary) as avgSal from instructor`
`group by dept name };`
`while (iter.next()) {`
`String deptName = iter.dept_name();`
`int avgSal = iter.avgSal();`
`System.out.println(deptName + " " + avgSal);`
`}`
`iter.close();`

SQLJ Resources

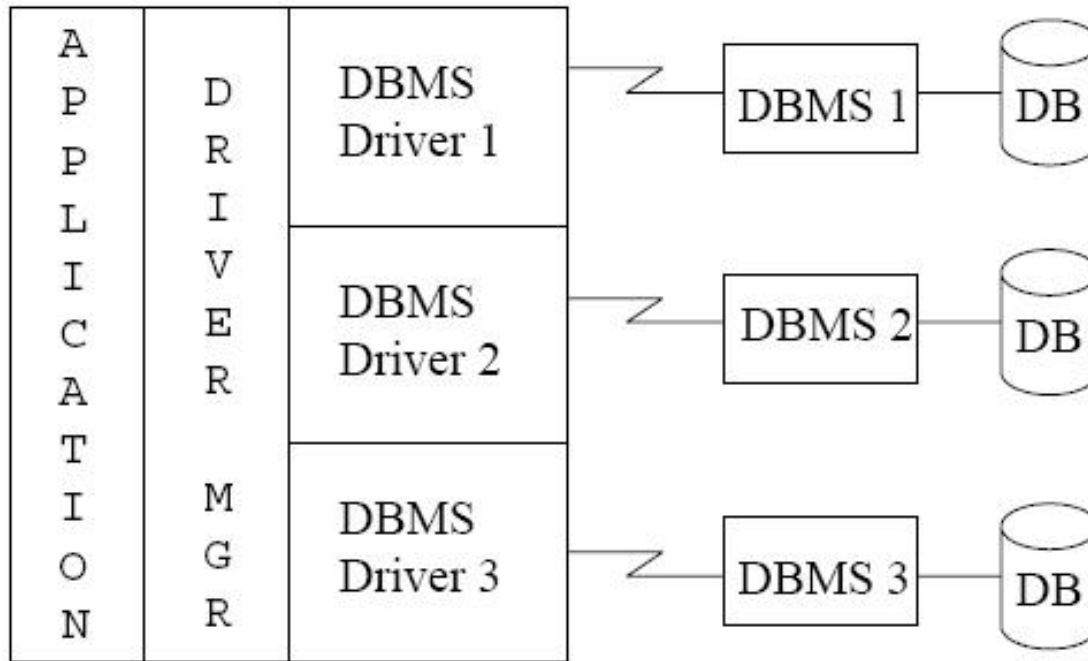
- n SQLJ Developer's Guide
- n <https://docs.oracle.com/en/database/oracle/oracle-database/20/jsqlj/toc.htm>

ODBC

- n **Open DataBase Connectivity(ODBC)** standard
 - | standard for application program to communicate with a database server.
 - | application program interface (API) to
 - ▶ open a connection with a database,
 - ▶ send queries and updates,
 - ▶ get back results.
- n Applications such as **GUI, spreadsheets**, etc. can use ODBC
- n Was defined originally for **Basic** and **C**, versions available for many languages.

ODBC (Cont.)

- Each database system supporting ODBC provides a "driver" library that must be linked with the client program.



ODBC (Cont.)

- n Each database system supporting ODBC provides a "driver" library that must be linked with the client program.
- n When client program makes an ODBC API call, the code in the library communicates with the server to carry out the requested action, and fetch results.
- n ODBC program first allocates an SQL environment, then a database connection handle.
- n Opens database connection using SQLConnect(). Parameters for SQLConnect:
 - | connection handle,
 - | the server to which to connect
 - | the user identifier,
 - | password
- n Must also specify types of arguments:
 - | SQL_NTS denotes previous argument is a null-terminated string.

ODBC Code

```
n int ODBCexample()
{
    RETCODE error;
    HENV  env;  /* environment */
    HDBC  conn; /* database connection */
    SQLAllocEnv(&env);
    SQLAllocConnect(env, &conn);
    SQLConnect(conn, "db.yale.edu", SQL_NTS, "avi", SQL_NTS,
               "avipasswd", SQL_NTS);
    { .... Do actual work ... }

    SQLDisconnect(conn);
    SQLFreeConnect(conn);
    SQLFreeEnv(env);
}
```

ODBC Code (Cont.)

- n Program sends SQL commands to database by using `SQLExecDirect`
- n Result tuples are fetched using `SQLFetch()`
- n `SQLBindCol()` binds C language variables to attributes of the query result
 - | When a tuple is fetched, its attribute values are automatically stored in corresponding C variables.
 - | Arguments to `SQLBindCol()`
 - ▶ ODBC stmt variable, attribute position in query result
 - ▶ The type conversion from SQL to C.
 - ▶ The address of the variable.
 - ▶ For variable-length types like character arrays,
 - The maximum length of the variable
 - Location to store actual length when a tuple is fetched.
 - Note: A negative value returned for the length field indicates null value
- n Good programming requires checking results of every function call for errors; we have omitted most checks for brevity.

ODBC Code (Cont.)

n Main body of program

```
char deptname[80];
float salary;
int lenOut1, lenOut2;
HSTMT stmt;
char * sqlquery = "select dept_name, sum (salary)
                  from instructor
                  group by dept_name";
SQLAllocStmt(conn, &stmt);
error = SQLExecDirect(stmt, sqlquery, SQL_NTS);
if (error == SQL_SUCCESS) {
    SQLBindCol(stmt, 1, SQL_C_CHAR, deptname , 80, &lenOut1);
    SQLBindCol(stmt, 2, SQL_C_FLOAT, &salary, 0 , &lenOut2);
    while (SQLFetch(stmt) == SQL_SUCCESS) {
        printf (" %s %g\n", deptname, salary);
    }
}
SQLFreeStmt(stmt, SQL_DROP);
```


ODBC Prepared Statements

n Prepared Statement

- | SQL statement prepared: compiled at the database
- | Can have placeholders: E.g. insert into account values(?,?,?)
- | Repeatedly executed with actual values for the placeholders

n To prepare a statement

`SQLPrepare(stmt, <SQL String>);`

n To bind parameters

`SQLBindParameter(stmt, <parameter#>, ... type information and value omitted for simplicity..)`

n To execute the statement

`retcode = SQLExecute(stmt);`

SQL injection(SQL注入)

n An SQL injection example in Java:

n “**select * from instructor where name= '** ” + **name + ” ’ ”**

n malicious hacker enters **name** : **X' or 'Y' ='Y**

→

select * from instructor where name= ' X' or 'Y' ='Y ’

n In the resulting SQL statement, the **where** clause is always true and the entire instructor relation is returned.

n Use of **prepared statement** would prevent this problem because the input string would have **escape characters(转义符)** inserted, so the resulting query becomes:

n **select * from instructor where name= ' X\' or \'Y\' =\'Y ’**

n This query is harmless and returns the empty relation.

n To avoid SQL injection security risk, do not create SQL strings directly using user input; instead use prepared statements to bind user inputs.

More ODBC Features

n Metadata features

- | finding all the relations in the database and
- | finding the names and types of columns of a query result or a relation in the database.

n By default, each SQL statement is treated as a separate transaction that is committed automatically.

- | Can turn off automatic commit on a connection
 - ▶ `SQLSetConnectOption(conn, SQL_AUTOCOMMIT, 0)`
- | Transactions must then be committed or rolled back explicitly by
 - ▶ `SQLTransact(conn, SQL_COMMIT)` or
 - ▶ `SQLTransact(conn, SQL_ROLLBACK)`

ODBC Conformance Levels

- n Conformance levels specify subsets of the functionality defined by the standard.
 - | Core
 - | Level 1 requires support for metadata querying
 - | Level 2 requires ability to send and retrieve arrays of parameter values and more detailed catalog information.
- n SQL Call Level Interface (CLI) standard similar to ODBC interface, but with some minor differences.

ODBC Resources

- n MySQL Connector/ODBC Developer Guide
- n https://docs.oracle.com/cd/E17952_01/connector-odbc-en/index.html

Embedded SQL

- n The SQL standard defines embeddings of SQL in a variety of programming languages such as **C**, **C++**, **Java**, **Fortran**, and **PL/1**,
- n A language to which SQL queries are embedded is referred to as a **host language**, and the SQL structures permitted in the host language comprise *embedded SQL*.
- n The basic form of these languages follows that of the System R embedding of SQL into PL/1.
- n **EXEC SQL** statement is used in the host language to identify embedded SQL request to the preprocessor

EXEC SQL <embedded SQL statement >;

Note: this varies by language:

- | In some languages, like **COBOL**, the semicolon is replaced with END-EXEC
- | In **Java** embedding uses **# SQL { }**;

Embedded SQL (Cont.)

- n Before executing any SQL statements, the program must first connect to the database. This is done using:
`EXEC-SQL connect to server user user-name using password;`
Here, *server* identifies the server to which a connection is to be established.
- n Variables of the **host language** can be used within embedded SQL statements. They are preceded by a colon (:) to distinguish from SQL variables (e.g., `:credit_amount`)
- n **Host Variables** used as above must be declared within **DECLARE** section, as illustrated below. The syntax for declaring the variables, however, follows the usual host language syntax.

EXEC-SQL BEGIN DECLARE SECTION

int *credit-amount* ;

EXEC-SQL END DECLARE SECTION;

Embedded SQL (Cont.)

- n To write an embedded SQL query, we use the

declare *c* cursor for <SQL query>

statement. The variable *c* is used to identify the query

- n Example:

- | From within a host language, find the ID and name of students who have completed more than the number of credits stored in variable *credit_amount* in the host language
- | Specify the query in SQL as follows:

EXEC SQL

declare *c* cursor for

select *ID, name*

from *student*

where *tot_cred* > :*credit_amount*;

Embedded SQL (Cont.)

- n The **open** statement for our example is as follows:

```
EXEC SQL open c ;
```

This statement causes the database system to execute the query and to save the results within a temporary relation. The query uses the value of the host-language variable *credit-amount* at the time the **open** statement is executed.

- n The fetch statement causes the values of one tuple in the query result to be placed on host language variables.

```
EXEC SQL fetch c into :si, :sn;
```

Repeated calls to fetch get successive tuples in the query result

Embedded SQL (Cont.)

- n A variable called **SQLSTATE** in the SQL communication area (**SQLCA**) gets set to '02000' to indicate no more data is available
- n The **close** statement causes the database system to delete the temporary relation that holds the result of the query.

EXEC SQL close c ;

Updates Through Embedded SQL

- n Embedded SQL expressions for database modification (**update**, **insert**, and **delete**)
- n Can update tuples fetched by cursor by declaring that the cursor is for update

EXEC SQL

```
declare c cursor for  
select *  
from instructor  
where dept_name = 'Music'  
for update
```

- n We then iterate through the tuples by performing **fetch** operations on the cursor , and after fetching each tuple we execute the following code:

```
update instructor  
set salary = salary + 1000  
where current of c
```

Embedded SQL Resources

- n Oracle Database Programmer's Guide - Embedded SQL
- n <https://docs.oracle.com/en/database/oracle/oracle-database/20/lnpcc/embedded-SQL.html#GUID-C671CABF-202A-4503-A16B-DC78D3F1AB13>

Procedural Constructs in SQL

Procedural Extensions and Stored Procedures

- n SQL provides a **module** language
 - | Permits definition of procedures in SQL, with **if-then-else** statements, **for** and **while** loops, etc.
- n **Stored Procedures**
 - | Can store procedures in the database
 - | then execute them using the **call** statement
 - | permit external applications to operate on the database without knowing about internal details

Functions and Procedures

- n Functions and procedures allow “business logic” to be stored in the database and executed from SQL statements.
- n These can be defined either by the procedural component of SQL or by an external programming language such as Java, C, or C++.
- n The syntax we present here is defined by the SQL standard.
 - | Most databases implement nonstandard versions of this syntax.

SQL Functions

- n Define a function that, given the name of a department, returns the count of the number of instructors in that department.

```
create function dept_count (dept_name varchar(20))  
returns integer  
begin  
    declare d_count integer;  
    select count ( * ) into d_count  
    from instructor  
    where instructor.dept_name = dept_name  
    return d_count;  
end
```

- n Find the department name and budget of all departments with more that 12 instructors.

```
select dept_name, budget  
from department  
where dept_count (dept_name) > 12
```


Table Functions

- n SQL:2003 added functions that return a relation as a result
- n Example: Return all accounts owned by a given customer

```
create function instructors_of (dept_name char(20) )
```

```
    returns table (   ID varchar(5),  
                    name varchar(20),  
                    dept_name varchar(20),  
                    salary numeric(8,2))
```

```
return table
```

```
    (select ID, name, dept_name, salary  
    from instructor  
    where instructor.dept_name = instructors_of.dept_name)
```

- n Usage

```
select *  
from table (instructors_of ('Music'))
```

SQL Procedures

- n The *dept_count* function could instead be written as procedure:

```
create procedure dept_count_proc (in dept_name varchar(20),  
                                out d_count integer)
```

```
begin
```

```
    select count(*) into d_count
```

```
    from instructor
```

```
    where instructor.dept_name = dept_count_proc.dept_name
```

```
end
```

- n Procedures can be invoked either from an SQL procedure or from embedded SQL, using the **call** statement.

```
    declare d_count integer;
```

```
    call dept_count_proc( 'Physics', d_count);
```

Procedures and functions can be invoked also from dynamic SQL

Procedural Constructs

- n Warning: most database systems implement their own variant of the standard syntax below
 - | read your system manual to see what works on your system
- n Compound statement: **begin ... end**,
 - | May contain multiple SQL statements between **begin** and **end**.
 - | Local variables can be declared within a compound statements
- n **While** and **repeat** statements :

```
declare n integer default 0;
while n < 10 do
    set n = n + 1
end while

repeat
    set n = n - 1
until n = 0
end repeat
```

Procedural Constructs (Cont.)

n **For** loop

- | Permits iteration over all results of a query
- | Example:

```
declare n integer default 0;  
for r as  
    select budget from department  
    where dept_name = 'Music'  
do  
    set n = n - r.budget  
end for
```

Procedural Constructs (cont.)

- n Conditional statements (**if-then-else**)

```
if boolean expression
    then statement or compound statement
elseif boolean expression
    then statement or compound statement
else statement or compound statement
end if
```

- n SQL:1999 also supports a **case** statement similar to C case statement

Example procedure

```
create function registerStudent(  
    in s_id varchar(5),  
    in s_courseid varchar (8),  
    in s_secid varchar (8),  
    in s_semester varchar (6),  
    in s_year numeric (4,0),  
    out errorMsg varchar(100)  
  
returns integer  
begin  
    declare currEnrol int;  
    select count(*) into currEnrol  
        from takes  
        where course_id = s_courseid and sec_id = s_secid  
            and semester = s_semester and year = s_year;  
    declare limit int;  
    select capacity into limit  
        from classroom natural join section  
        where course_id = s_courseid and sec_id = s_secid  
            and semester = s_semester and year = s_year;  
    if (currEnrol < limit)  
        begin  
            insert into takes values  
                (s_id, s_courseid, s_secid, s_semester, s_year, null);  
            return(0);  
        end  
    -- Otherwise, section capacity limit already reached  
    set errorMsg = 'Enrollment limit reached for course ' || s_courseid  
        || ' section ' || s_secid;  
    return(-1);  
end;
```

- Registers a student after ensuring
- classroom capacity is not exceeded.
- Returns 0 on success, and -1 if
- capacity is exceeded.

External Language Functions/Procedures

- n SQL:1999 permits the use of functions and procedures written in other languages such as C or C++
- n Declaring external language procedures and functions

```
create procedure dept_count_proc(in dept_name varchar(20),  
                                out count integer)
```

```
language C
```

```
external name ' /usr/avi/bin/dept_count_proc'
```

```
create function dept_count(dept_name varchar(20))
```

```
returns integer
```

```
language C
```

```
external name '/usr/avi/bin/dept_count'
```

External Language Routines (Cont.)

- n Benefits of external language functions/procedures:
 - | more efficient for many operations, and more expressive power.
- n Drawbacks
 - | Code to implement function may need to be loaded into database system and executed in the database system's address space.
 - ▶ risk of accidental corruption of database structures
 - ▶ security risk, allowing users access to unauthorized data
 - | There are alternatives, which give good security at the cost of potentially worse performance.
 - | Direct execution in the database system's space is used when efficiency is more important than security.

Security with External Language Routines

- n To deal with security problems
 - | Use **sandbox** techniques
 - ▶ that is use a safe language like Java, which cannot be used to access/damage other parts of the database code.
 - | Or, run external language functions/procedures in a separate process, with no access to the database process' memory.
 - ▶ Parameters and results communicated via inter-process communication
- n Both have performance overheads
- n Many database systems support both above approaches as well as direct executing in database system address space.

Triggers

Triggers

- n A **trigger** is a statement that is executed automatically by the system as a side effect of a modification to the database.
- n **Trigger - ECA rule**
 - | **E: Event** (**insert, delete , update**)
 - | **C: Condition**
 - | **A: Action**
- n To design a trigger mechanism, we must:
 - | Specify the conditions under which the trigger is to be executed.
 - | Specify the actions to be taken when the trigger executes.
- n Triggers introduced to SQL standard in SQL:1999, but supported even earlier using non-standard syntax by most databases.

Trigger Example

```
n  account_log(account, amount, datetime)

n  create trigger account_trigger after update of account on balance
    referencing new row as nrow
    referencing old row as orow
    for each row
    when nrow.balance - orow.balance >=200000 or
        orow.balance -nrow.balance >=50000
    begin
        insert into account_log values (nrow.account-number,
                                           nrow.balance-orow.balance , current_time() )
    end
```

Trigger Example

- n E.g. *time_slot_id* is not a primary key of *timeslot*, so we cannot create a foreign key constraint from *section* to *timeslot*.
- n Alternative: use triggers on *section* and *timeslot* to enforce integrity constraints

```
create trigger timeslot_check1 after insert on section  
referencing new row as nrow  
for each row  
when (nrow.time_slot_id not in (  
    select time_slot_id  
    from time_slot)) /* time_slot_id not present in time_slot */  
begin  
    rollback  
end;
```

Trigger Example Cont.

```
create trigger timeslot_check2 after delete on timeslot  
  referencing old row as orow  
  for each row  
  when (orow.time_slot_id not in (  
    select time_slot_id  
    from time_slot)  
    /* last tuple for time slot id deleted from time slot */  
    and orow.time_slot_id in (  
      select time_slot_id  
      from section))  
    /* and time_slot_id still referenced from section */  
  begin  
    rollback  
  end;
```

Triggering Events and Actions in SQL

- n Triggering event can be **insert**, **delete** or **update**
- n Triggers on update can be restricted to specific attributes
 - | **E.g., after update of *takes* on *grade***
- n Values of attributes before and after an update can be referenced
 - | **referencing old row as** : for deletes and updates
 - | **referencing new row as** : for inserts and updates
- n Triggers can be activated before an event, which can serve as extra constraints. E.g. convert blank grades to null.

```
create trigger setnull_trigger before update of takes  
referencing new row as nrow  
for each row  
when (nrow.grade = ' ')  
begin atomic  
    set nrow.grade = null;  
end;
```

Trigger to Maintain credits_earned value

```
n create trigger credits_earned after update of takes on grade
referencing new row as nrow
referencing old row as orow
for each row
when nrow.grade <> 'F' and nrow.grade is not null
and (orow.grade = 'F' or orow.grade is null)
begin atomic
  update student
  set tot_cred= tot_cred +
    (select credits
     from course
     where course.course_id= nrow.course_id)
  where student.id = nrow.id;
end;
```


Statement Level Triggers

- n Instead of executing a separate action for each affected row, a single action can be executed for all rows affected by a transaction
 - | Use **for each statement** instead of **for each row**
 - | Use **referencing old table** or **referencing new table** to refer to temporary tables (called *transition tables*) containing the affected rows
 - | Can be more efficient when dealing with SQL statements that update a large number of rows

Statement Level Triggers

```
n create trigger grade_trigger after update of takes on grade  
referencing new table as new_table  
for each statement  
when exists( select avg(grade)  
              from new_table  
              group by course_id, sec_id, semester, year  
              having avg(grade)< 60 )  
begin  
    rollback  
end
```

When Not To Use Triggers

- n Triggers were used earlier for tasks such as
 - | maintaining summary data (e.g., total salary of each department)
 - | Replicating databases by recording changes to special relations (called **change** or **delta** relations) and having a separate process that applies the changes over to a replica
- n There are better ways of doing these now:
 - | Databases today provide built in **materialized view** facilities to maintain summary data
 - | Databases provide built-in support for replication
- n Encapsulation facilities can be used instead of triggers in many cases
 - | Define methods to update fields
 - | Carry out actions as part of the update methods instead of through a trigger

When Not To Use Triggers

- n Risk of unintended execution of triggers, for example, when
 - | loading data from a backup copy
 - | replicating updates at a remote site
 - | Trigger execution can be disabled before such actions.
- n Other risks with triggers:
 - | Error leading to failure of critical transactions that set off the trigger
 - | Cascading execution

Recursive Queries

Recursion in SQL

- n SQL:1999 permits recursive view definition
- n Example: find which courses are a prerequisite, whether directly or indirectly, for a specific course

```
with recursive rec_prereq(course_id, prereq_id) as (  
    select course_id, prereq_id  
    from prereq  
    union  
    select rec_prereq.course_id, prereq.prereq_id,  
    from rec_prereq, prereq  
    where rec_prereq.prereq_id = prereq.course_id  
    )  
select *  
from rec_prereq;
```

This example view, *rec_prereq*, is called the *transitive closure* of the *prereq* relation

The Power of Recursion

- n Recursive views make it possible to write queries, such as transitive closure queries, that cannot be written without recursion or iteration.
 - | Intuition: Without recursion, a non-recursive non-iterative program can perform only a fixed number of joins of *prereq* with itself
 - ▶ This can give only a fixed number of levels of prerequisites.
 - ▶ Given a fixed non-recursive query, we can construct a database with a greater number of levels of prerequisites on which the query will not work
 - ▶ Alternative: write a procedure to iterate as many times as required
 - See procedure *findAllPrereqs* in book on Page.189

The Power of Recursion

```
create function findAllPrereqs(cid varchar(8))
-- Finds all courses that are prerequisite (directly or indirectly) for cid
returns table (course_id varchar(8))
-- The relation prereq(course_id, prereq_id) specifies which course is
-- directly a prerequisite for another course.
begin
  create temporary table c_prereq (course_id varchar(8));
  -- table c_prereq stores the set of courses to be returned
  create temporary table new_c_prereq (course_id varchar(8));
  -- table new_c_prereq contains courses found in the previous iteration
  create temporary table temp (course_id varchar(8));
  -- table temp is used to store intermediate results
  insert into new_c_prereq
  select prereq_id
  from prereq
  where course_id = cid;
  repeat
    insert into c_prereq
    select course_id
    from new_c_prereq;

    insert into temp
    (select prereq.course_id
     from new_c_prereq, prereq
     where new_c_prereq.course_id = prereq.prereq_id
    )
    except (
      select course_id
      from c_prereq
    );
    delete from new_c_prereq;
    insert into new_c_prereq
    select *
    from temp;
    delete from temp;

  until not exists (select * from new_c_prereq)
  end repeat;
  return table c_prereq;
end
```


Advanced Aggregation Features

Ranking

- n Ranking is done in conjunction with an order by specification.

- n Suppose we are given a relation
student_grades(*ID*, *GPA*)
giving the grade-point average of each student

- n Find the rank of each student.

```
select ID, rank() over (order by GPA desc) as s_rank  
from student_grades
```

- n An extra **order by** clause is needed to get them in sorted order

```
select ID, rank() over (order by GPA desc) as s_rank  
from student_grades  
order by s_rank
```

- n Ranking may leave gaps: e.g. if 2 students have the same top GPA, both have rank 1, and the next rank is 3

- | **dense_rank** does not leave gaps, so next dense rank would be 2

Ranking

- n Ranking can be done using basic SQL aggregation, but resultant query is very inefficient

```
select ID, (1 + (select count(*)  
                from student_grades B  
                where B.GPA > A.GPA)) as s_rank  
from student_grades A  
order by s_rank;
```

Ranking (Cont.)

- n Ranking can be done within partition of the data.

- n “Find the rank of students within each department.”

```
select ID, dept_name,  
       rank () over (partition by dept_name order by GPA desc)  
       as dept_rank  
from dept_grades  
order by dept_name, dept_rank;
```

- n Multiple **rank** clauses can occur in a single **select** clause.

- n Ranking is done *after* applying **group by** clause/aggregation

- n Can be used to find top-n results

- | More general than the **limit** *n* clause supported by many databases, since it allows top-n within each partition

Ranking (Cont.)

n Other ranking functions:

- | **percent_rank** (within partition, if partitioning is done)

- | **cume_dist** (cumulative distribution)

 - ▶ fraction of tuples with preceding values

- | **row_number** (non-deterministic in presence of duplicates)

n SQL:1999 permits the user to specify **nulls first** or **nulls last**

select *ID*,

rank () over (order by *GPA* desc nulls last) as *s_rank*

from *student_grades*

Ranking (Cont.)

- n For a given constant n , the ranking the function $ntile(n)$ takes the tuples in each partition in the specified order, and divides them into n buckets with equal numbers of tuples.
- n E.g.,
select *ID*, **ntile**(4) **over** (**order by** *GPA desc*) **as** *quartile*
from *student_grades*;

Windowing

- n Used to smooth out random variations.
- n E.g., **moving average**: “Given sales values for each date, calculate for each date the average of the sales on that day, the previous day, and the next day”
- n **Window specification** in SQL:
 - | Given relation *sales(date, value)*
select *date*, **sum**(*value*) **over**
 (**order by** *date* **between** rows 1 **preceding** and 1 **following**)
from *sales*

Windowing

- n Examples of other window specifications:
 - | **between rows unbounded preceding and current**
 - | **rows unbounded preceding**
 - | **range between 10 preceding and current row**
 - ▶ All rows with values between current row value –10 to current value
 - | **range interval 10 day preceding**
 - ▶ Not including current row

Windowing (Cont.)

- n Can do windowing within partitions
- n E.g., Given a relation *transaction* (*account_number*, *date_time*, *value*), where *value* is positive for a deposit and negative for a withdrawal
 - | “Find total balance of each account after each transaction on the account”

```
select account_number, date_time,  
       sum (value) over  
         (partition by account_number  
          order by date_time  
          rows unbounded preceding)  
       as balance  
from transaction  
order by account_number, date_time
```

OLAP**

Data Analysis and OLAP

n Online Analytical Processing (OLAP)

- | Interactive analysis of data, allowing data to be summarized and viewed in different ways in an online fashion (with negligible delay)
- n Data that can be modeled as dimension attributes and measure attributes are called **multidimensional data**.
 - | **Measure attributes**
 - ▶ measure some value
 - ▶ can be aggregated upon
 - ▶ e.g., the attribute *number* of the *sales* relation
 - | **Dimension attributes**
 - ▶ define the dimensions on which measure attributes (or aggregates thereof) are viewed
 - ▶ e.g., attributes *item_name*, *color*, and *size* of the *sales* relation

Example sales relation

<i>item_name</i>	<i>color</i>	<i>clothes_size</i>	<i>quantity</i>
skirt	dark	small	2
skirt	dark	medium	5
skirt	dark	large	1
skirt	pastel	small	11
skirt	pastel	medium	9
skirt	pastel	large	15
skirt	white	small	2
skirt	white	medium	5
skirt	white	large	3
dress	dark	small	2
dress	dark	medium	6
dress	dark	large	12
dress	pastel	small	4
dress	pastel	medium	3
dress	pastel	large	3
dress	white	small	2
dress	white	medium	3
dress	white	large	0
shirt	dark	small	2
shirt	dark	medium	4
...
...

Cross Tabulation of sales by *item_name* and *color*


clothes_size **all**

		<i>color</i>			
		dark	pastel	white	total
<i>item_name</i>	skirt	8	35	10	53
	dress	20	10	5	35
	shirt	14	7	28	49
	pants	20	2	5	27
	total	62	54	48	164

- n The table above is an example of a **cross-tabulation** (**cross-tab**), also referred to as a **pivot-table**.
 - | Values for one of the dimension attributes form the row headers
 - | Values for another dimension attribute form the column headers
 - | Other dimension attributes are listed on top
 - | Values in individual cells are (aggregates of) the values of the dimension attributes that specify the cell.


Data Cube

- n A **data cube** is a multidimensional generalization of a cross-tab
- n Can have n dimensions; we show 3 below
- n Cross-tabs can be used as views on a data cube




Hierarchies on Dimensions

- n **Hierarchy** on dimension attributes: lets dimensions to be viewed at different levels of detail
 - H E.g., the dimension DateTime can be used to aggregate by hour of day, date, day of week, month, quarter or year



a) Time Hierarchy



b) Location Hierarchy

Cross Tabulation With Hierarchy

- n Cross-tabs can be easily extended to deal with hierarchies
 - Can drill down or roll up on a hierarchy

clothes_size: **all**

		<i>color</i>				
<i>category</i>	<i>item_name</i>	dark	pastel	white	total	
womenswear	skirt	8	8	10	53	88
	dress	20	20	5	35	
	subtotal	28	28	15		
menswear	pants	14	14	28	49	76
	shirt	20	20	5	27	
	subtotal	34	34	33		
total		62	62	48		164

Relational Representation of Cross-tabs

- n Cross-tabs can be represented as relations
- We use the value **all** is used to represent aggregates.
 - The SQL standard actually uses null values in place of **all** despite confusion with regular null values.

<i>item_name</i>	<i>color</i>	<i>clothes_size</i>	<i>quantity</i>
skirt	dark	all	8
skirt	pastel	all	35
skirt	white	all	10
skirt	all	all	53
dress	dark	all	20
dress	pastel	all	10
dress	white	all	5
dress	all	all	35
shirt	dark	all	14
shirt	pastel	all	7
shirt	White	all	28
shirt	all	all	49
pant	dark	all	20
pant	pastel	all	2
pant	white	all	5
pant	all	all	27
all	dark	all	62
all	pastel	all	54
all	white	all	48
all	all	all	164

Extended Aggregation to Support OLAP

- n The **cube** operation computes union of **group by**'s on every subset of the specified attributes
- n Example relation for this section
sales(item_name, color, clothes_size, quantity)
- n E.g. consider the query

```
select item_name, color, size, sum(number)  
from sales  
group by cube(item_name, color, size)
```

This computes the union of eight different groupings of the *sales* relation:

```
{ (item_name, color, size), (item_name, color),  
  (item_name, size),      (color, size),  
  (item_name),            (color),  
  (size),                 ( ) }
```

where () denotes an empty **group by** list.

- n For each grouping, the result contains the null value for attributes not present in the grouping.

Online Analytical Processing Operations

- n Relational representation of cross-tab that we saw earlier, but with *null* in place of **all**, can be computed by

```
select item_name, color, sum(number)  
from sales  
group by cube(item_name, color)
```

- n The function **grouping()** can be applied on an attribute
 - | Returns 1 if the value is a null value representing all, and returns 0 in all other cases.

```
select item_name, color, size, sum(number),  
      grouping(item_name) as item_name_flag,  
      grouping(color) as color_flag,  
      grouping(size) as size_flag,  
from sales  
group by cube(item_name, color, size)
```

Online Analytical Processing Operations

- n Can use the function **decode()** in the **select** clause to replace such nulls by a value such as **all**
 - | E.g., replace *item_name* in first query by
decode(grouping(item_name), 1, 'all', item_name)

Extended Aggregation (Cont.)

- n The **rollup** construct generates union on every prefix of specified list of attributes

- n E.g.,

```
select item_name, color, size, sum(number)  
from sales  
group by rollup(item_name, color, size)
```

Generates union of four groupings:

{ (*item_name*, *color*, *size*), (*item_name*, *color*), (*item_name*), () }

- n Rollup can be used to generate aggregates at multiple levels of a hierarchy.
- n E.g., suppose table *itemcategory*(*item_name*, *category*) gives the category of each item. Then

```
select category, item_name, sum(number)  
from sales, itemcategory  
where sales.item_name = itemcategory.item_name  
group by rollup(category, item_name)
```

would give a hierarchical summary by *item_name* and by *category*.

Extended Aggregation (Cont.)

- n Multiple rollups and cubes can be used in a single group by clause
 - | Each generates set of group by lists, cross product of sets gives overall set of group by lists

n E.g.,

```
select item_name, color, size, sum(number)  
from sales  
group by rollup(item_name), rollup(color, size)
```

generates the groupings

$$\{item_name, ()\} \times \{(color, size), (color), ()\}$$
$$= \{ (item_name, color, size), (item_name, color), (item_name), (color, size), (color), () \}$$

Online Analytical Processing Operations

- n **Pivoting:** changing the dimensions used in a cross-tab is called
- n **Slicing:** creating a cross-tab for fixed values only
 - | Sometimes called **dicing**, particularly when values for multiple dimensions are fixed.
- n **Rollup:** moving from finer-granularity data to a coarser granularity
- n **Drill down:** The opposite operation - that of moving from coarser-granularity data to finer-granularity data

OLAP Implementation

- n The earliest OLAP systems used multidimensional arrays in memory to store data cubes, and are referred to as **multidimensional OLAP (MOLAP)** systems.
- n OLAP implementations using only relational database features are called **relational OLAP (ROLAP)** systems
- n Hybrid systems, which store some summaries in memory and store the base data and other summaries in a relational database, are called **hybrid OLAP (HOLAP)** systems.

OLAP Implementation (Cont.)

- n Early OLAP systems precomputed *all* possible aggregates in order to provide online response
 - | Space and time requirements for doing so can be very high
 - ▶ 2^n combinations of **group by**
 - | It suffices to precompute some aggregates, and compute others on demand from one of the precomputed aggregates
 - ▶ Can compute aggregate on *(item_name, color)* from an aggregate on *(item_name, color, size)*
 - For all but a few “non-decomposable” aggregates such as *median*
 - is cheaper than computing it from scratch
- n Several optimizations available for computing multiple aggregates
 - | Can compute aggregate on *(item_name, color)* from an aggregate on *(item_name, color, size)*
 - | Can compute aggregates on *(item_name, color, size)*, *(item_name, color)* and *(item_name)* using a single sorting of the base data

End of Chapter

Figure 5.22

<i>item_name</i>	<i>clothes_size</i>	<i>dark</i>	<i>pastel</i>	<i>white</i>
skirt	small	2	11	2
skirt	medium	5	9	5
skirt	large	1	15	3
dress	small	2	4	2
dress	medium	6	3	3
dress	large	12	3	0
shirt	small	2	4	17
shirt	medium	6	1	1
shirt	large	6	2	10
pant	small	14	1	3
pant	medium	6	0	0
pant	large	0	1	2

Figure 5.23

<i>item_name</i>	<i>quantity</i>
skirt	53
dress	35
shirt	49
pant	27

Figure 5.24

<i>item_name</i>	<i>color</i>	<i>quantity</i>
skirt	dark	8
skirt	pastel	35
skirt	white	10
dress	dark	20
dress	pastel	10
dress	white	5
shirt	dark	14
shirt	pastel	7
shirt	white	28
pant	dark	20
pant	pastel	2
pant	white	5

Another Recursion Example

- n Given relation
manager(employee_name, manager_name)
- n Find all employee-manager pairs, where the employee reports to the manager directly or indirectly (that is manager's manager, manager's manager's manager, etc.)

```
with recursive empl (employee_name, manager_name) as (  
    select employee_name, manager_name  
    from   manager  
    union  
    select manager.employee_name, empl.manager_name  
    from   manager, empl  
    where manager.manager_name = empl.employee_name)  
select *  
from   empl
```

This example view, *empl*, is the *transitive closure* of the *manager* relation

Merge statement (now in Chapter 24)

- n Merge construct allows batch processing of updates.
- n Example: relation *funds_received* (*account_number*, *amount*) has batch of deposits to be added to the proper account in the *account* relation

```
merge into account as A
  using (select *
        from funds_received as F)
on (A.account_number = F.account_number)
when matched then
  update set balance = balance + F.amount
```

End of Chapter 5