



浙江大学
ZheJiang University

Computer Organization & Design

—The Hardware/Software Interface

Chapter 4-1 Processor Design





Contents

- **4.1 Introduction**
- 4.2 Logic Design Conventions
- 4.3 Building a datapath
- 4.4 A Simple Implementation Scheme
 - Single cycle CPU
 - Multiple cycle CPU
- 4.5 Pipelining





4.1 Introduction

■ CPU performance factors

- Instruction count
 - Determined by ISA and compiler
- CPI and Cycle time
 - Determined by CPU hardware

■ We will examine two RISC-V implementations

- A simplified version
- A more realistic pipelined version

■ Simple subset, shows most aspects

- Memory reference: ld, sd
- Arithmetic/logical: add, sub, and, or, slt
- Control transfer: beq, jal

□ 实现不少于下列指令

R-Type: add, sub, and, or, xor, srl, sra;

I-Type: addi, andi, ori, xori, slti, slli, srli,
lb, lh, lw, lbu, lhu, lwu; Jalr

S-Type: sb, sh, sw;

B-Type: beq, bne, blt, bge, bltu, bgeu

UJ-Type: Jal;

U-type: Lui;





- Generic Implementation:
 - use the program counter (PC) to supply instruction address
 - get the instruction from memory
 - read registers
 - use the instruction to decide exactly what to do





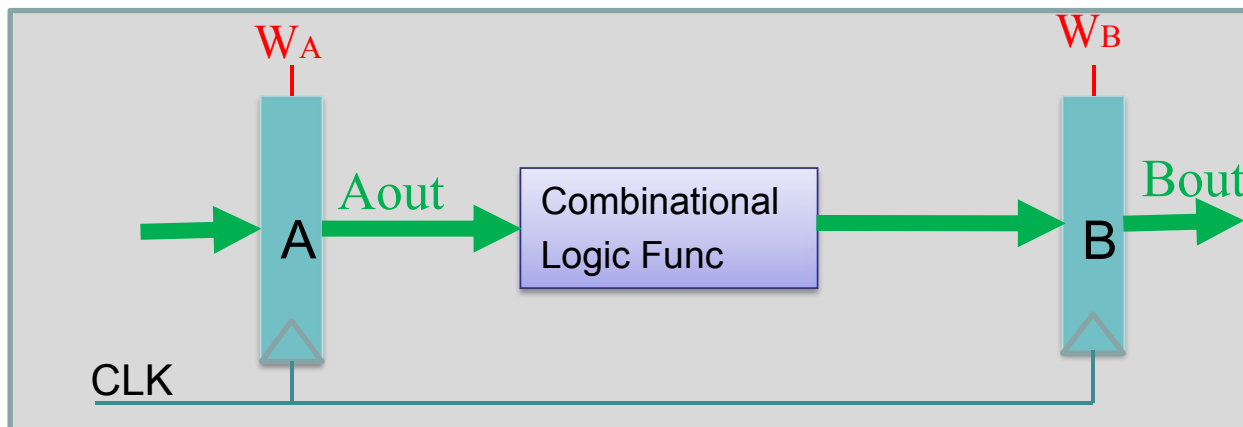
Instruction Execution Overview

- **For every instruction, the first two steps are identical**
 - Fetch the instruction from the memory
 - Decode and read the registers
- **Next steps depend on the instruction class**
 - Memory-reference Arithmetic-logical branches
- **Depending on instruction class**
 - Use ALU to calculate
 - Arithmetic result
 - Memory address for load/store
 - Branch comparison
 - Access data memory for load/store
 - $PC \leftarrow \text{target address or } PC + 4$










Abstraction Model of Datapath



Function(A) \rightarrow B



Timing of D-flip-flop

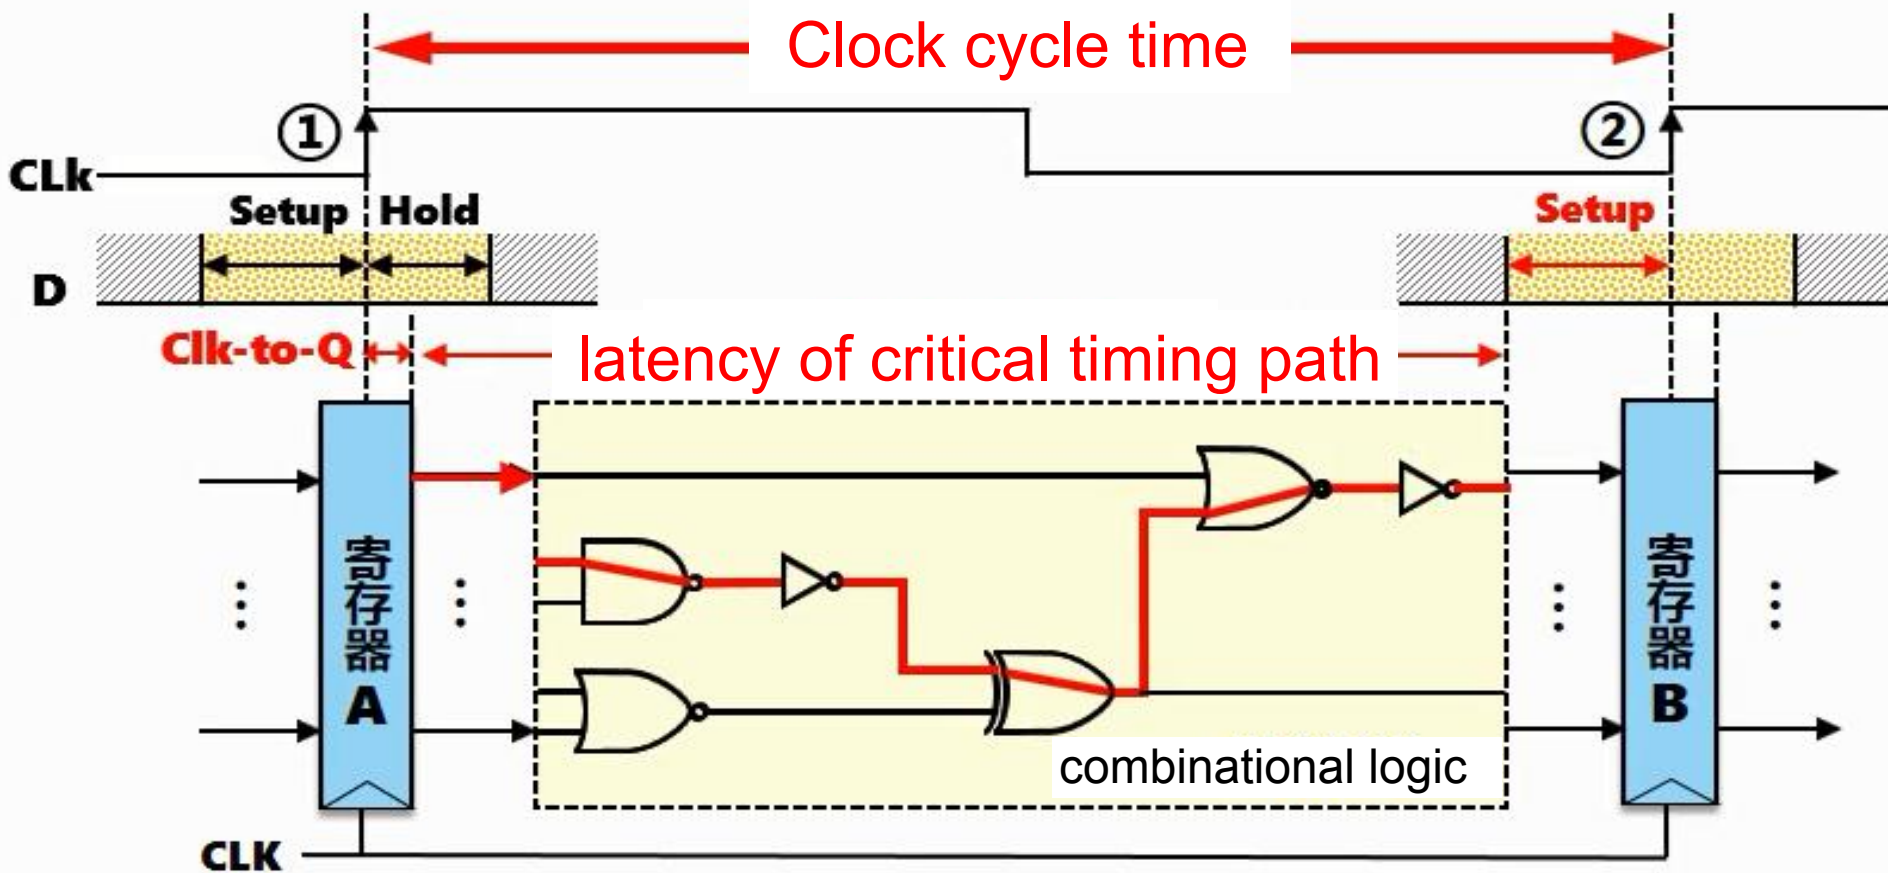


- **Setup time:** Input data keep stable before rising edge of CLK
- **Hold time:** Input data keep stable after rising edge of CLK
- **Latency of D-flip-flop:** time of CLK-to-Q





Datapath & Clock cycle time

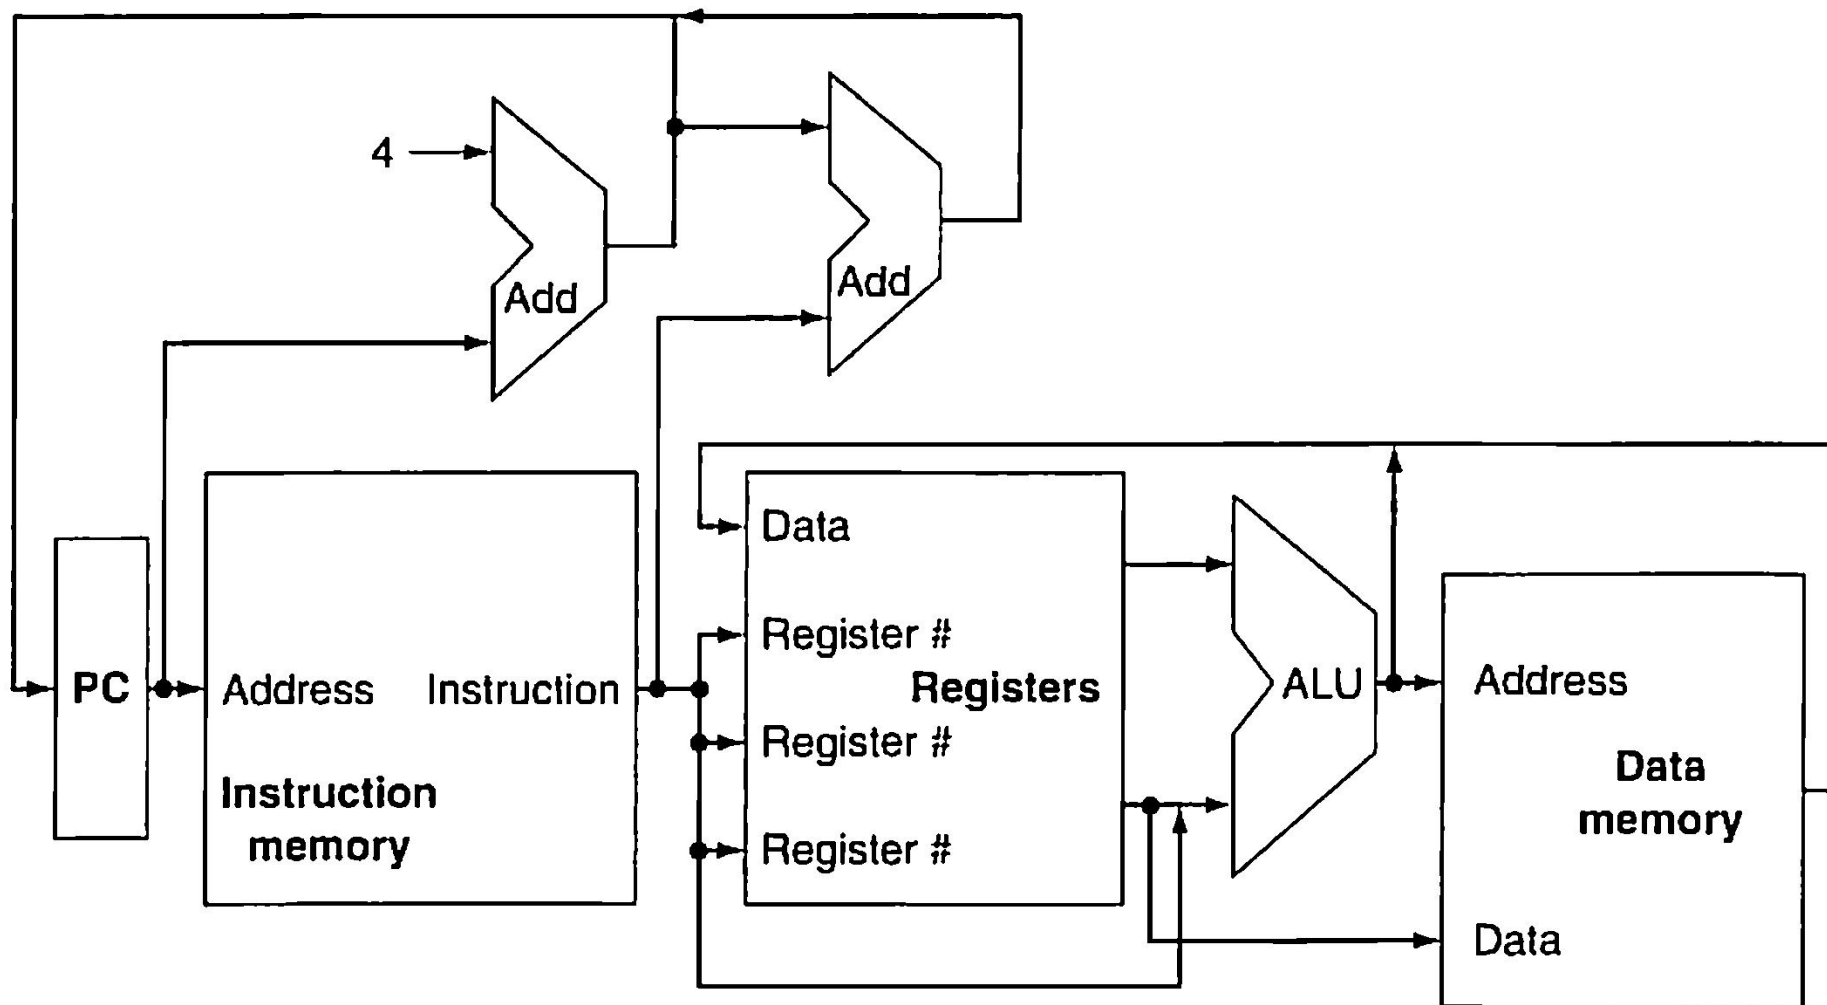


■ 时钟周期 > Clk_to_Q + 关键路径时延 + Setup Time





An abstract view of the implementation of RISC-V






Contents


- 4.1 Introduction
- **4.2 Logic Design Conventions**
- 4.3 Building a datapath
- 4.4 A Simple Implementation Scheme
- 4.5 An Overview of Pipelining




Single implementation





a. Instruction memory



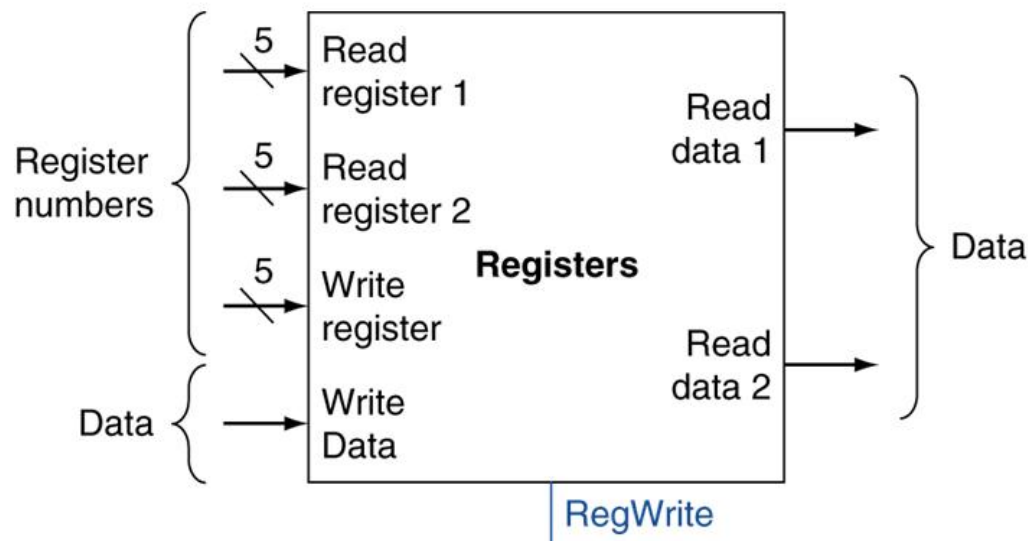

b. Program counter




c. Adder



b. Immediate generation unit



a. Registers




a. Data memory unit





State Elements

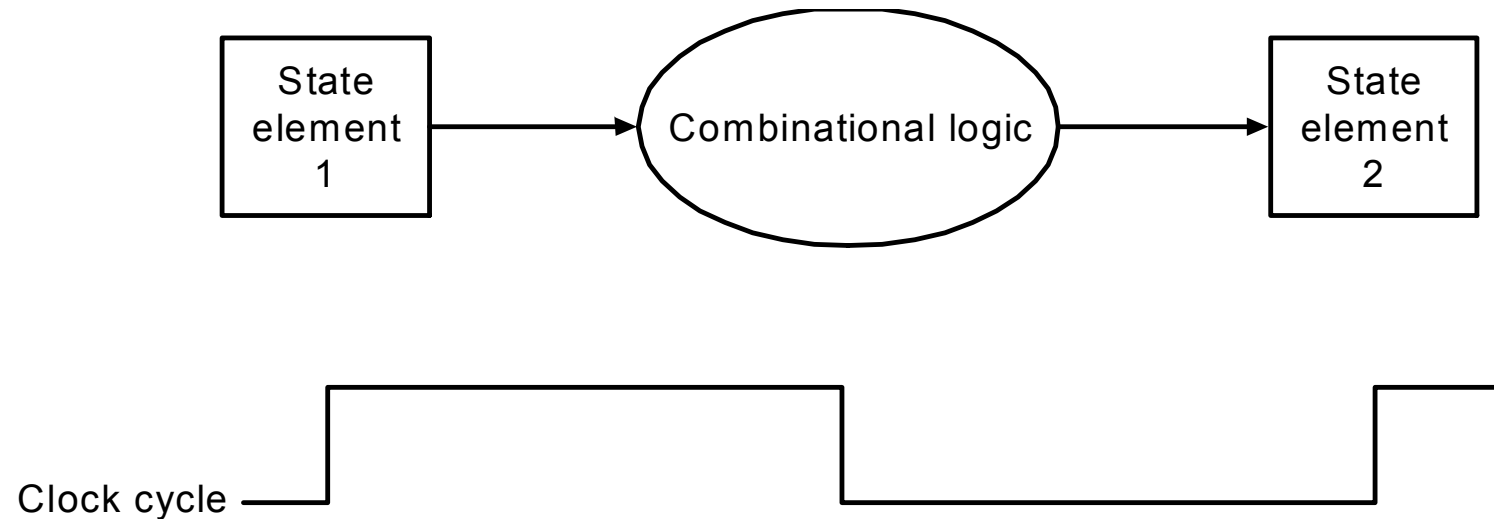
- Unclocked vs. Clocked
- Clocks used in synchronous logic
 - when should an element that contains state be updated?





Our Implementation

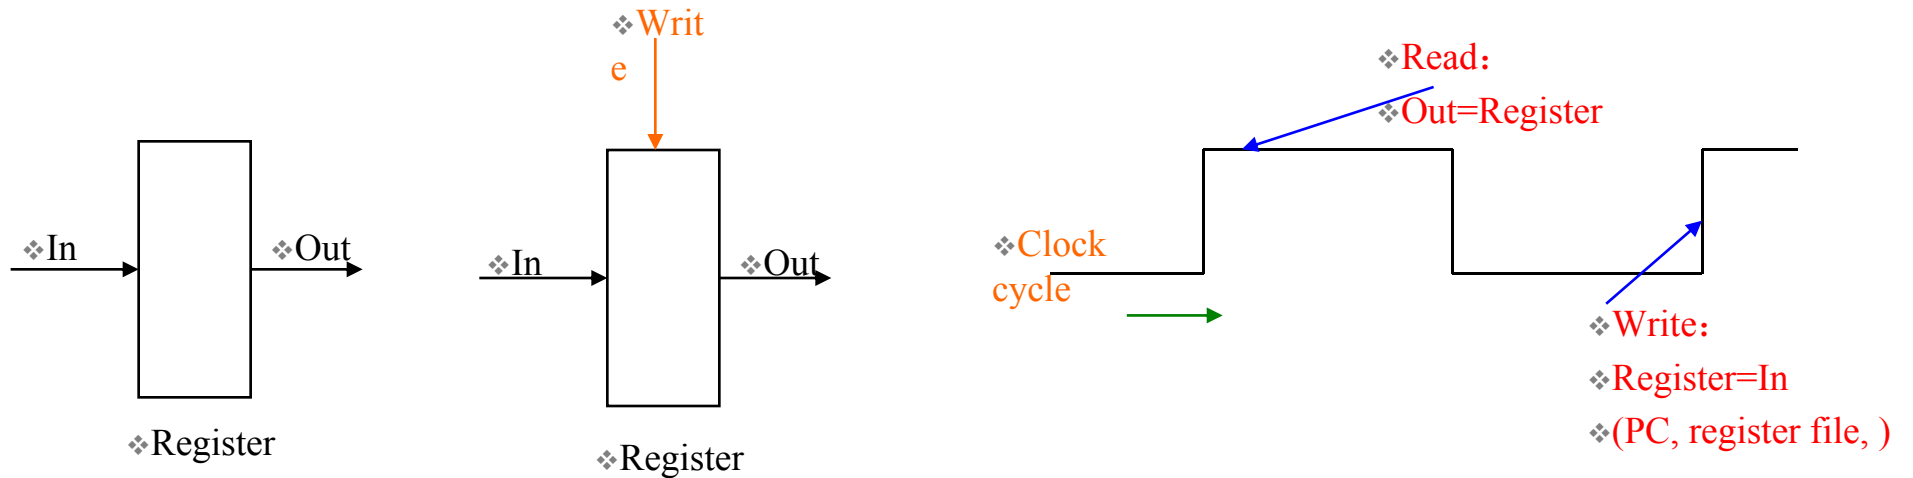
- An edge triggered methodology
- Typical execution:
 - read contents of some state elements,
 - send values through some combinational logic
 - write results to one or more state elements





REGISTER



- Register
 - State element.
 - Can be controlled by **Write** signal.





Memory

- **Instruction memory** : ROM, input address, get instruction
- **Data memory**: RAM.
 - Can read or write controlled by signal MemRead and MemWrite.
 - Read: input address, get instruction if MemRead=1
 - Write: input address & datain, write into memory if MemWrite =1







D 触发器

| C | D | S | R | Q(t + 1) |
|---|---|---|---|------------|
| 0 | X | X | X | hold |
| 1 | 0 | 0 | 1 | Q=0: reset |
| 1 | 1 | 1 | 0 | Q=1: reset |

| C | D | Q(t + 1) |
|---|---|------------|
| 0 | X | hold |
| 1 | 0 | Q=0: reset |
| 1 | 1 | Q=1: set |

Truth table

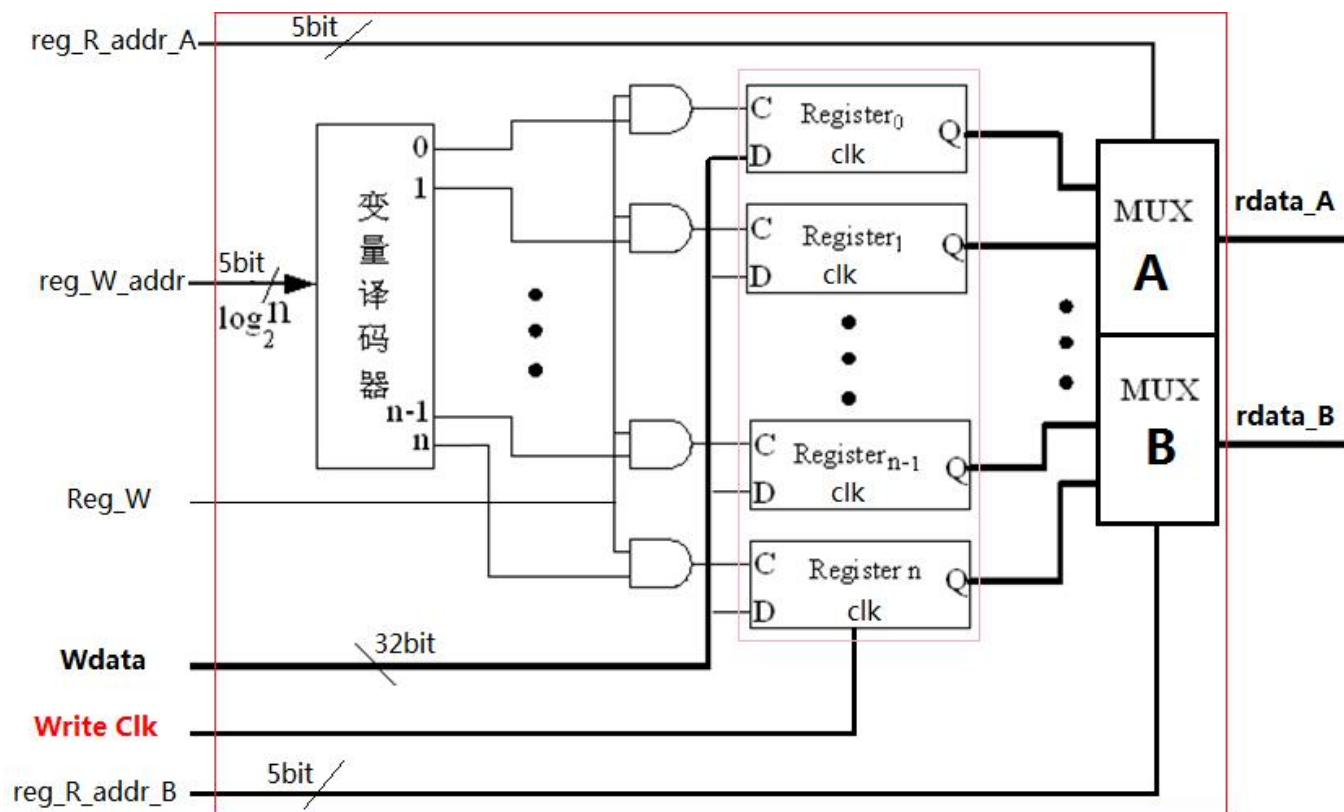


Symbol of Latch



Register File

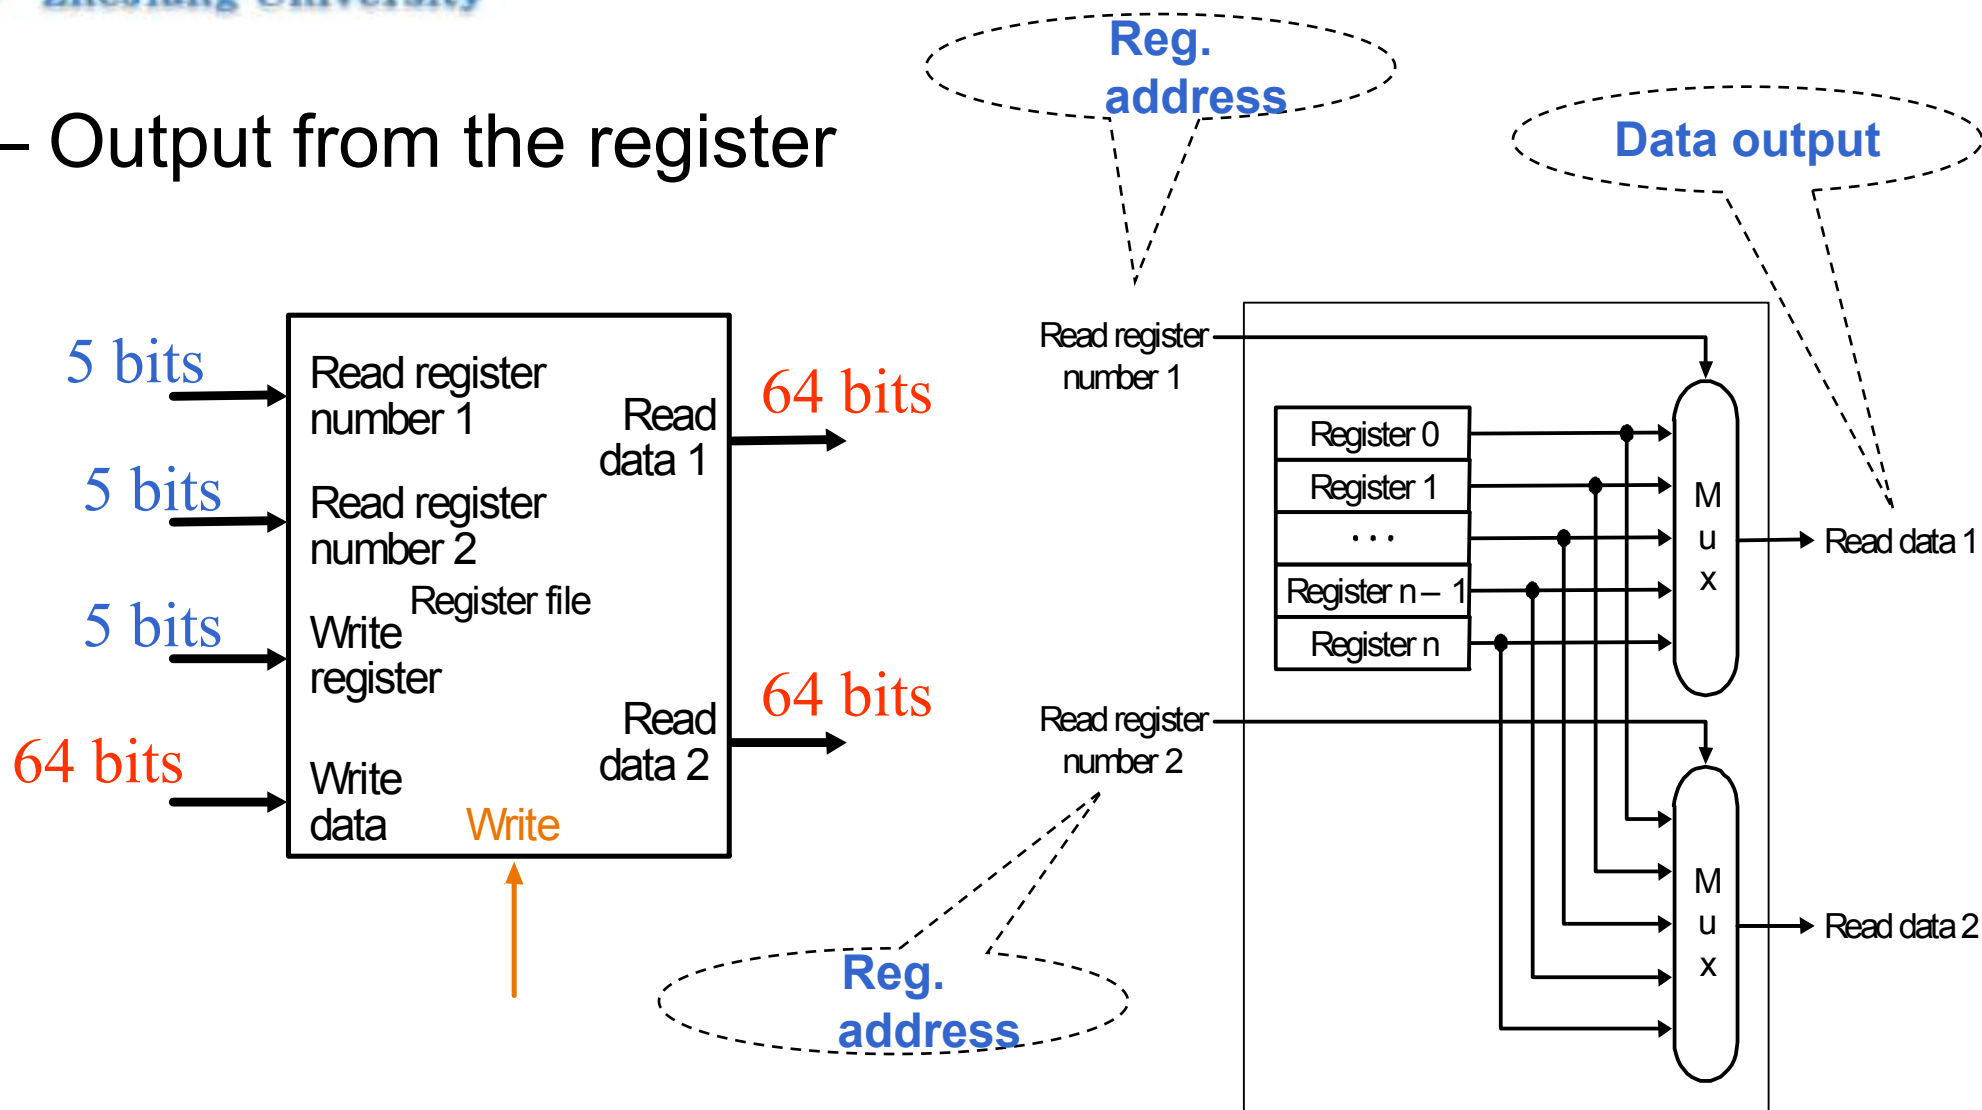
- 32 64-bit Registers;
- Input: 2 5-bit register number/ one 5-bit register number and 64-bit data;
- Output: 64-bit data;
- Register write control。





Register File: Read-Output

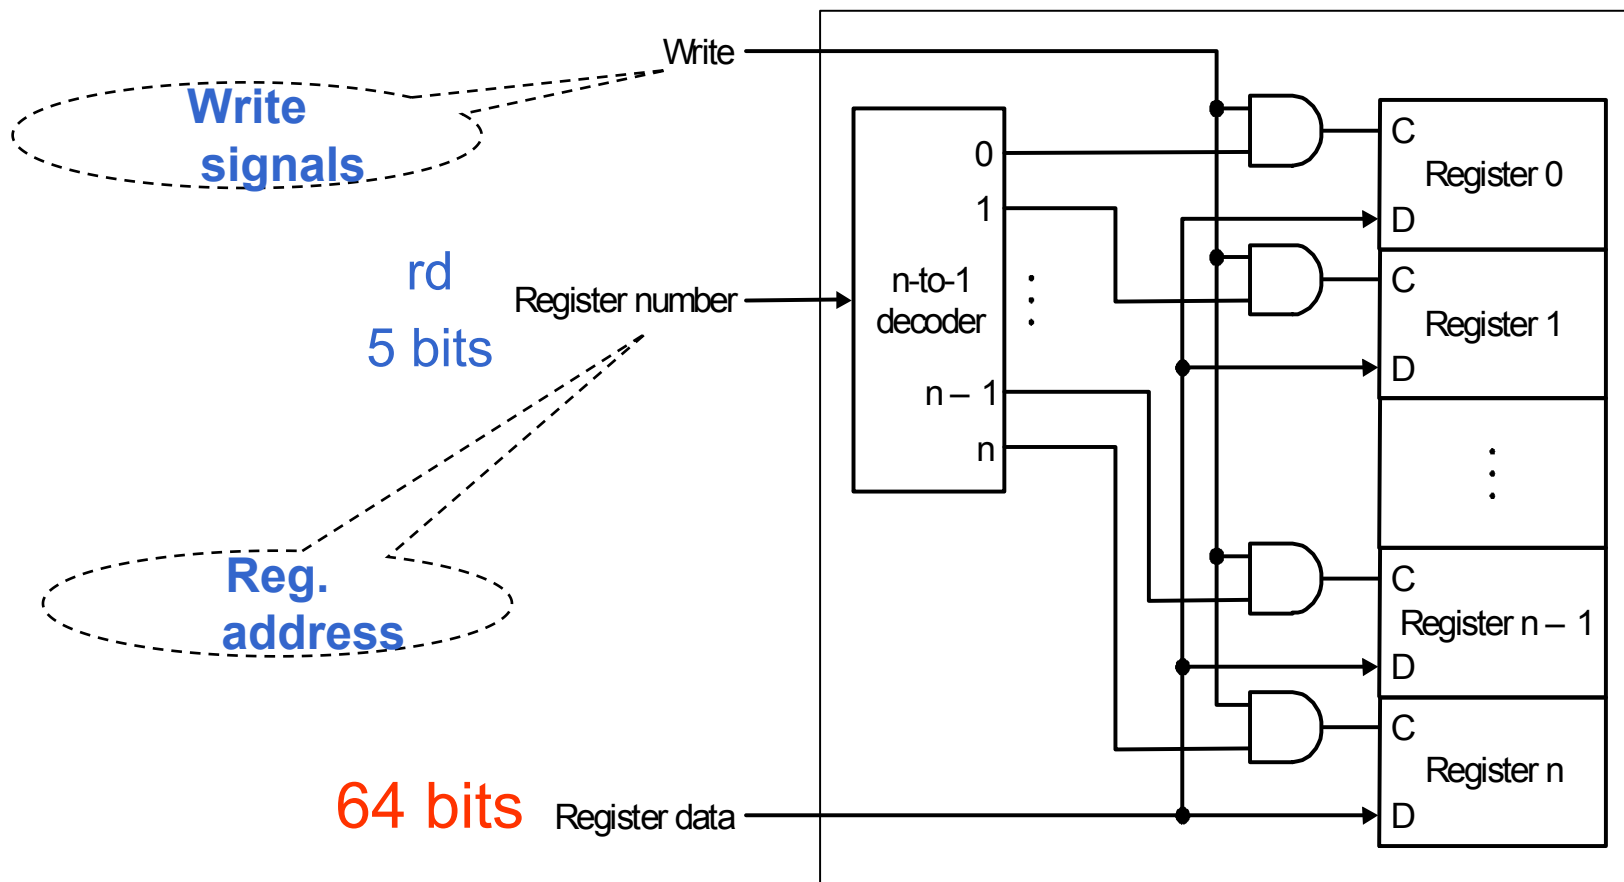
– Output from the register





Register File: Write-Input

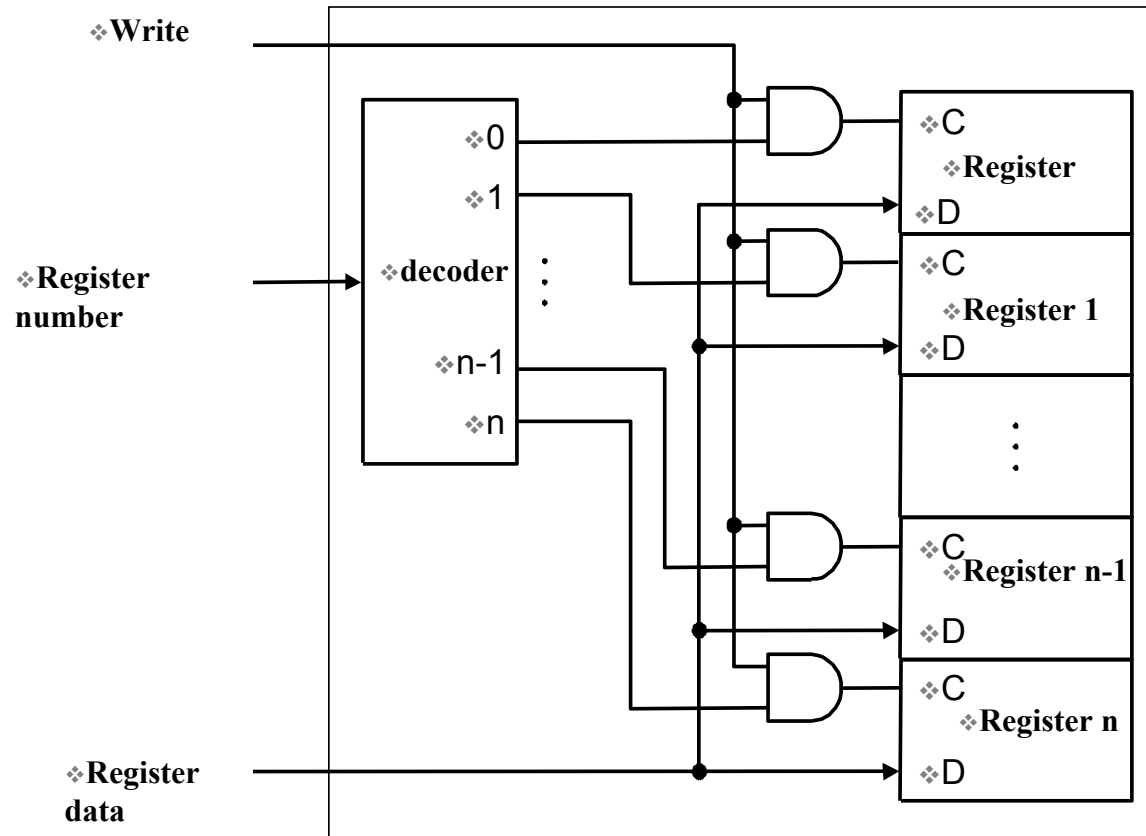
– Written to the register





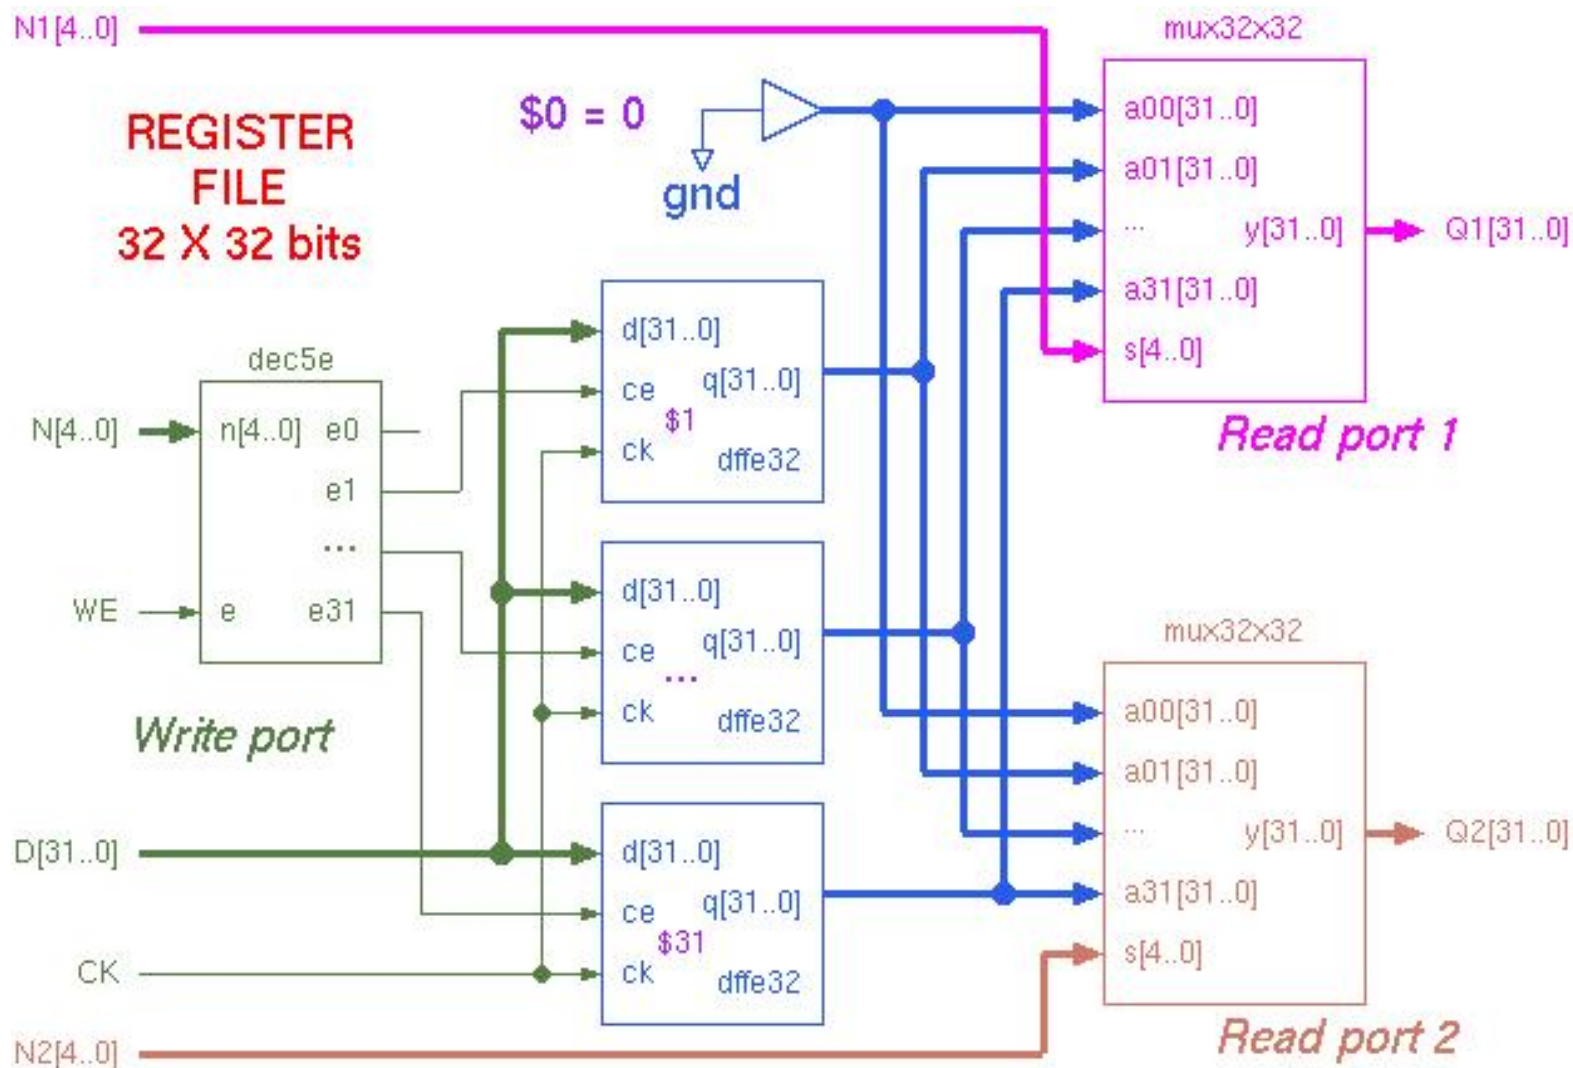
Register File

- Write Register:
 - we still use the real clock to determine when to write





Register File 32 X 32 bits





Description: 32×64 bits Register files

```
Module regs( input      clk, rst, RegWrite,
              input  [4:0] Rs1_addr, Rs2_addr, Wt_addr,
              input  [63:0] Wt_data,
              output [63:0] Rs1_data, Rs2_data
            );
    reg [63:0] register [1:31];           // r1 - r31
    integer i;


    assign rdata_A = (Rs1_addr== 0) ? 0 : register[Rs1_addr];           // read
    assign rdata_B = (Rs2_addr== 0) ? 0 : register[Rs2_addr];           // read
    always @(posedge clk or posedge rst)
    begin  if (rst==1)
            for (i=1; i<32; i=i+1) register[i] <= 0;                 // reset
        else if ((Wt_addr != 0) && (RegWrite == 1))
            register[Wt_addr] <= Wt_data;                             // write
    end
endmodule
```





The other elements

- Immediate generation unit:
 - 输入指令产生立即数的逻辑功能
 - 根据指令类型（加载，存储或者分支指令），产生相应的立即数
 - 转移指令偏移量左移位的功能
 - 立即数字段符号扩展为64位结果输出



b. Immediate generation unit

- Immediate generation

- Load:

0000011

L_imm = {{52{inst[31]}}, Inst[31:20]};

- Save:

0100011

S_imm = {{52{inst[31]}}, Inst[31:25], inst[11:7]};

- Branch:

1100011

SB_imm = {{51{inst [31]}}, inst[31], inst [7], inst[30:25], inst[11:8],1'b0};

- Jal:

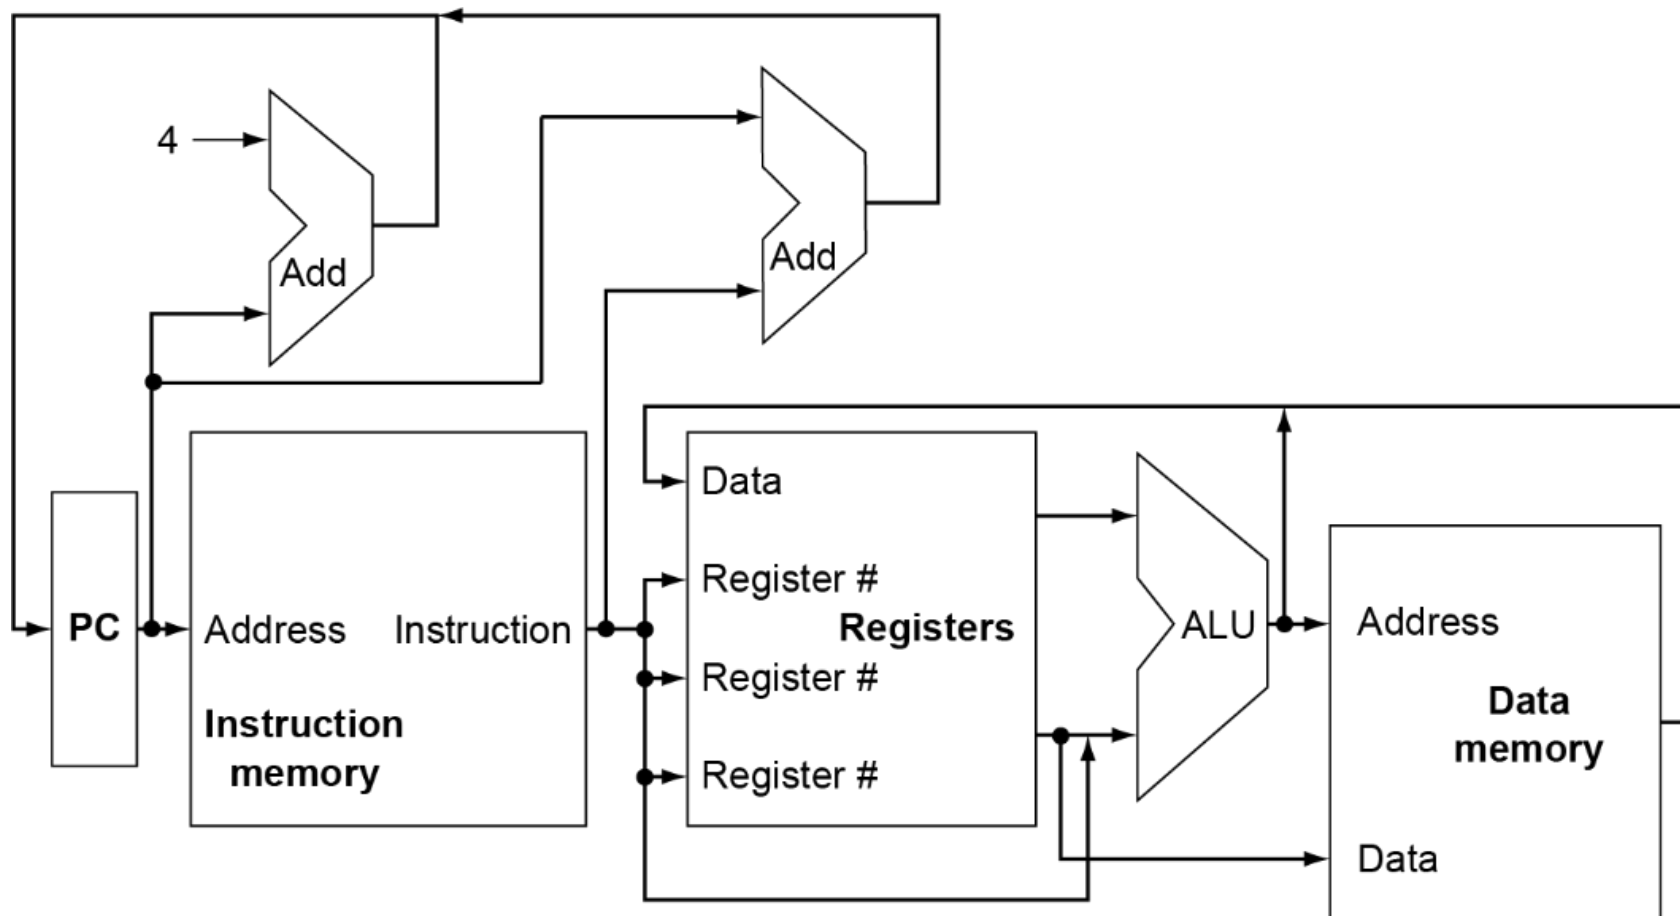
1101111

UJ = {{43{inst[31]}}, inst [31], inst[19:12], inst[20], inst[30:21],1'b0};



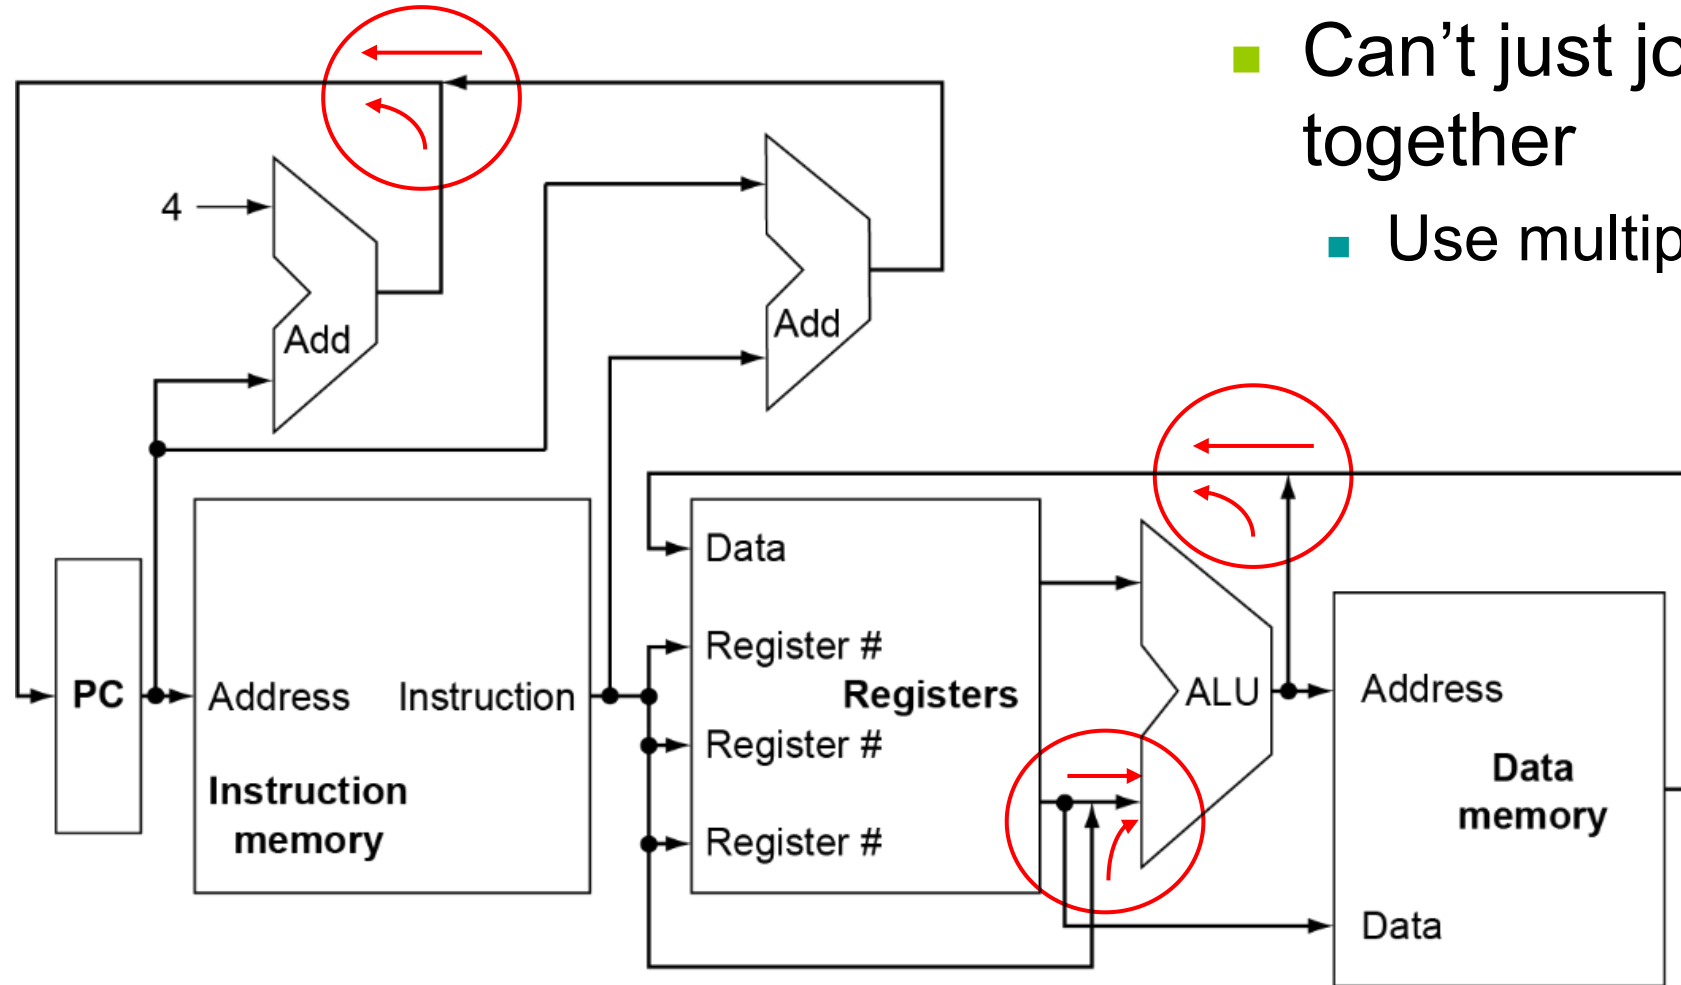


CPU Overview





Multiplexers

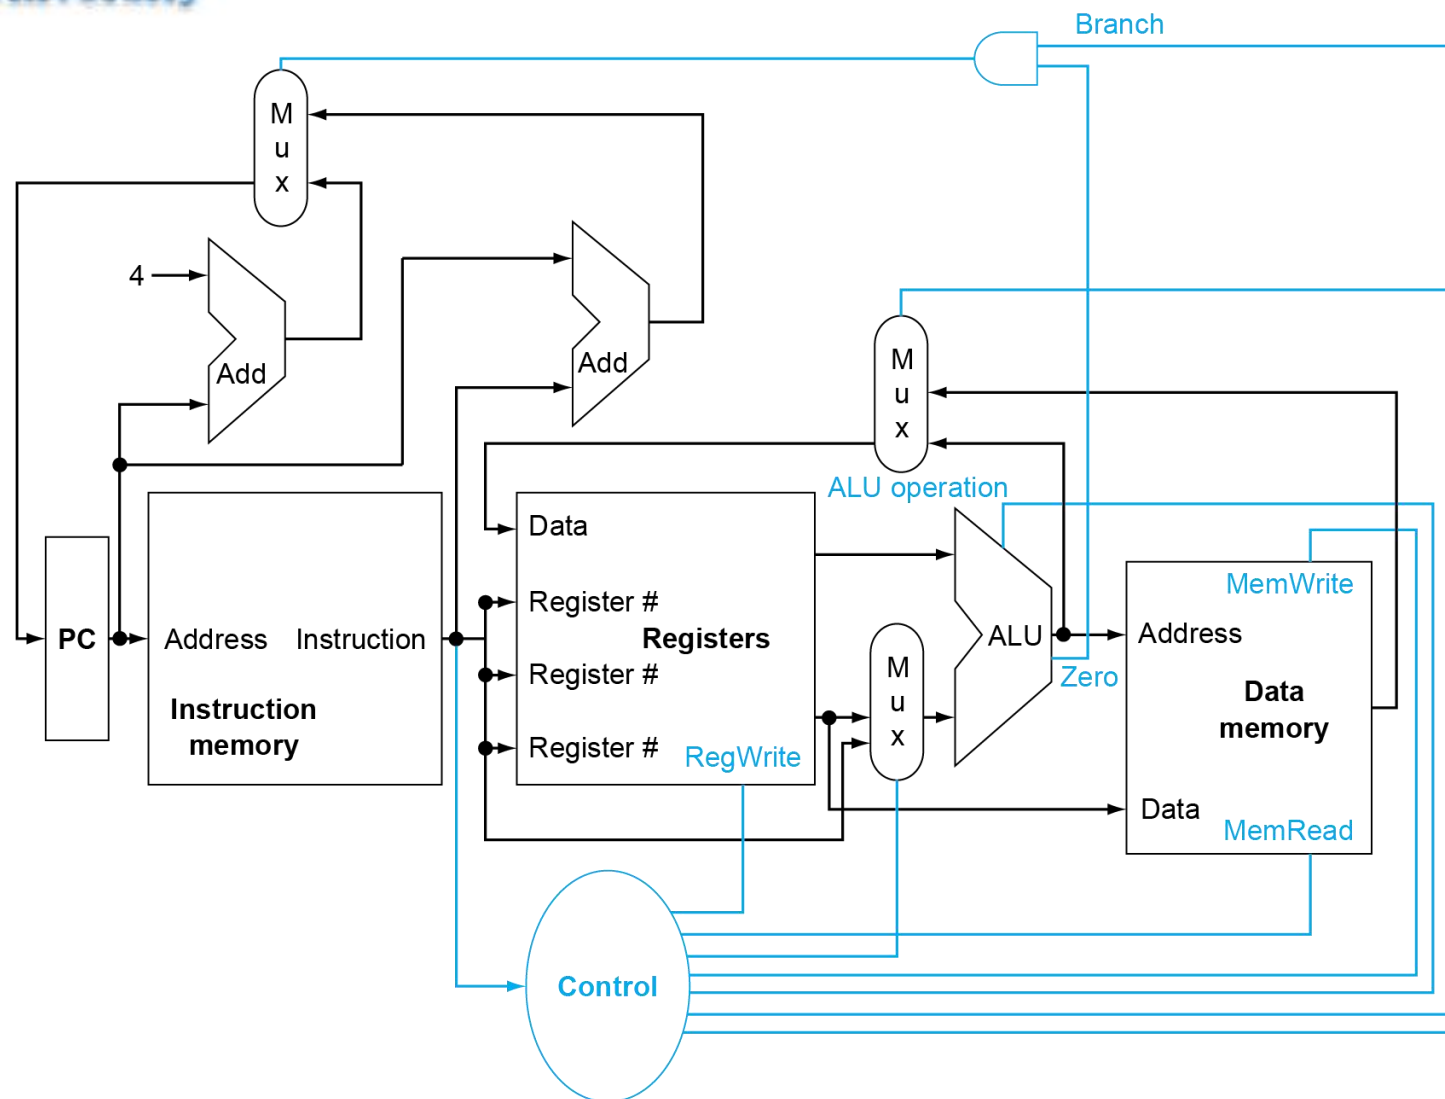


- Can't just join wires together
 - Use multiplexers





Control





Logic Design Conventions



- Information encoded in binary
 - Low voltage = 0, High voltage = 1
 - One wire per bit
 - Multi-bit data encoded on multi-wire buses
- Combinational element
 - Operate on data
 - Output is a function of input
- State (sequential) elements
 - Store information





Clocking Methodology

- An edge triggered methodology
- Typical execution:
 - read contents of some state elements,
 - send values through some combinational logic
 - write results to one or more state elements





Contents

- 4.1 Introduction
- 4.2 Logic Design Conventions
- **4.3 Building a datapath**
- 4.4 A Simple Implementation Scheme
 - Single cycle CPU
 - Multiple cycle CPU
- 4.5 Pipelining





Contents

- Introduction & Logic Design Conventions
- Building a datapath
- A Simple Implementation Scheme
- Pipelining
- Exceptions





4.3 Building a datapath

- Datapath
 - Elements that process data and addresses in the CPU
 - Registers, ALUs, mux's, memories, ...
- We will build a RISC-V datapath incrementally
 - Refining the overview design



RISC-V fields (format)

| Name (Field Size) | 7 bits | 5 bits | 5 bits | 3 bits | 5 bits | 7 bits | Comments |
|----------------------|-----------------------------|--------|--------|--------|---------------|--------|-------------------------------|
| R-type | funct7 | rs2 | rs1 | funct3 | rd | opcode | Arithmetic instruction format |
| I-type | immediate[11:0] | | rs1 | funct3 | rd | opcode | Loads & immediate arithmetic |
| S-type | immed[11:5] | rs2 | rs1 | funct3 | immed[4:0] | opcode | Stores |
| SB-type | immed[12,10:5] | rs2 | rs1 | funct3 | immed[4:1,11] | opcode | Conditional branch format |
| UJ-type | immediate[20,10:1,11,19:12] | | | | rd | opcode | Unconditional jump format |
| U-type | immediate[31:12] | | | | rd | opcode | Upper immediate format |

- **opcode:** basic operation of the instruction.
- **rs1:** the first register source operand.
- **rs2:** the second register source operand.
- **rd:** the register destination operand.
- **funct:** function, this field selects the specific variant of the operation in the op field.
- **Immediate:** address or immediate



| Name | Example | Comments |
|-----------------------|---|--|
| 32 registers | x0-x31 | Fast locations for data. In RISC-V, data must be in registers to perform arithmetic. Register x0 always equals 0. |
| 2^{61} memory words | Memory[0], Memory[8] , , Memory[18446744073709551608] | Accessed only by data transfer instructions. RISC-V uses byte addresses, so sequential doubleword accesses differ by 8. Memory holds data structures, arrays, and spilled registers. |

| Name | Register name | Usage | Preserved On call? |
|---------|---------------|-------------------------------|--------------------|
| x0 | 0 | The constant value 0 | n.a. |
| x1(ra) | 1 | Return address(link register) | yes |
| x2(sp) | 2 | Stack pointer | yes |
| x3(gp) | 3 | Global pointer | yes |
| x4(tp) | 4 | Thread pointer | yes |
| x5-x7 | 5-7 | Temporaries | no |
| x8-x9 | 8-9 | Saved | yes |
| x10-x17 | 10-17 | Arguments/results | no |
| x18-x27 | 18-27 | Saved | yes |
| x28-x31 | 28-31 | Temporaries | no |





RISC-V assembly language

| Category | Instruction | Example | Meaning | Comments |
|----------------------|------------------------|------------------|--|--|
| Arithmetic | add | add x5,x6,x7 | $x5 = x6 + x7$ | Add two source register operands |
| | subtract | sub x5,x6,x7 | $x5 = x6 - x7$ | First source register subtracts second one |
| | add immediate | addi x5,x6,20 | $x5 = x6 + 20$ | Used to add constants |
| Data transfer | load doubleword | ld x5, 40(x6) | $x5 = \text{Memory}[x6 + 40]$ | doubleword from memory to register |
| | store doubleword | sd x5, 40(x6) | $\text{Memory}[x6 + 40] = x5$ | doubleword from register to memory |
| Logical | and | and x5, x6, 3 | $x5 = x6 \& 3$ | Arithmetic shift right by register |
| | inclusive or | or x5,x6,x7 | $x5 = x6 x7$ | Bit-by-bit OR |
| Conditional Branch | branch if equal | beq x5, x6, 100 | if($x5 == x6$) go to PC+100 | PC-relative branch if registers equal |
| | branch if not equal | bne x5, x6, 100 | if($x5 != x6$) go to PC+100 | PC-relative branch if registers not equal |
| Unconditional Branch | jump and link | jal x1, 100 | $x1 = \text{PC} + 4$; go to PC+100 | PC-relative procedure call |
| | jump and link register | jalr x1, 100(x5) | $x1 = \text{PC} + 4$; go to $x5 + 100$ | procedure return; indirect call |





Instruction execution in RISC-V

- **Fetch :**
 - Take instructions from the instruction memory
 - Modify PC to point the next instruction
- **Instruction decoding & Read Operand:**
 - Will be translated into machine control command
 - Reading Register Operands, whether or not to use
- **Executive Control:**
 - Control the implementation of the corresponding ALU operation
- **Memory access:**
 - Write or Read data from memory
 - Only ld/sd
- **Write results to register:**
 - If it is R-type instructions, ALU results are written to rd
 - If it is I-type instructions, memory data are written to rd
- **Modify PC** for branch instructions

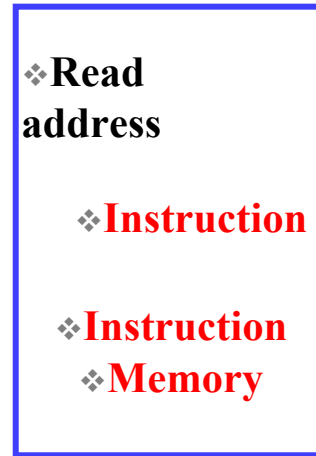





Instruction fetching three elements

Data Stream of Instruction fetching


How to connect? Who?



Instruction memory



Program counter



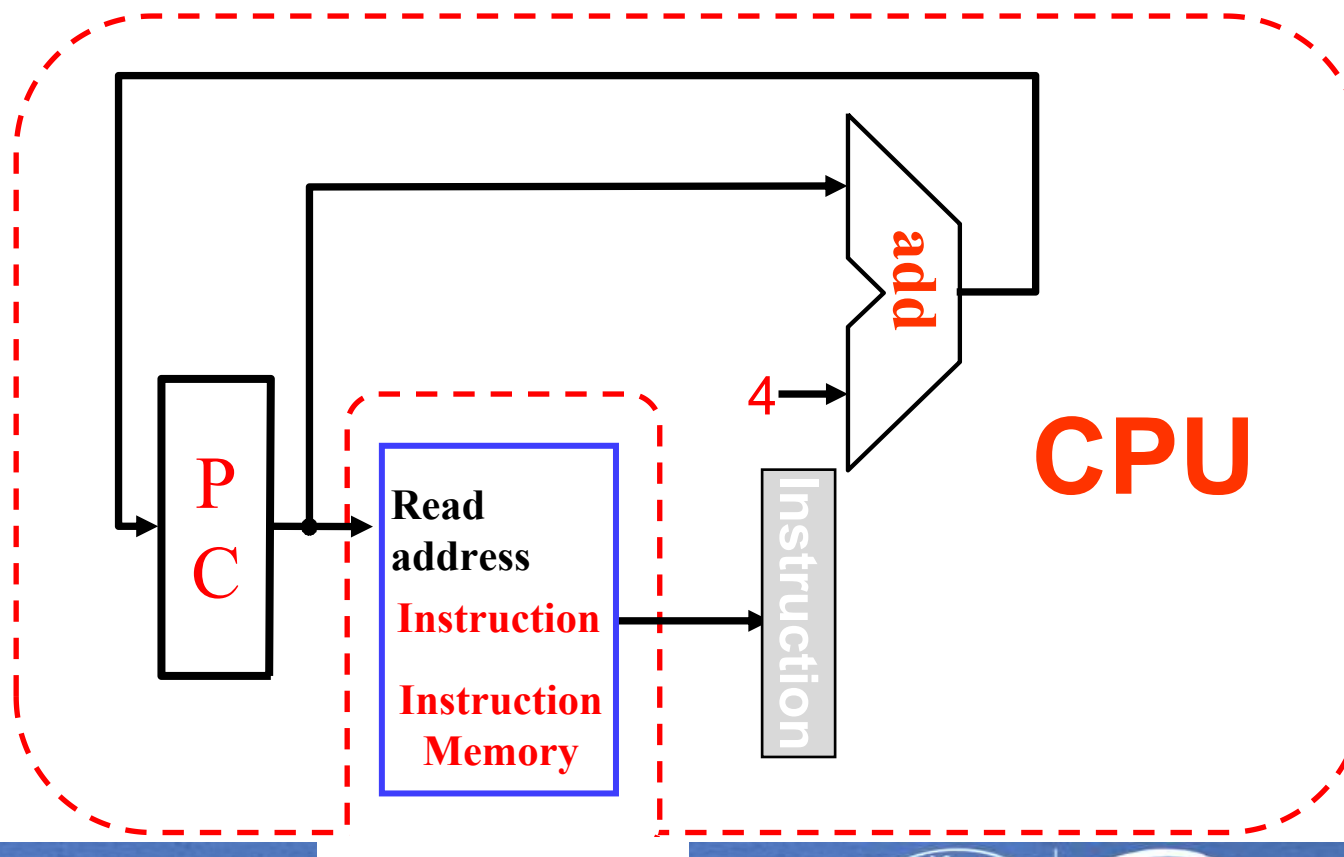
Adder





Instruction fetching unit

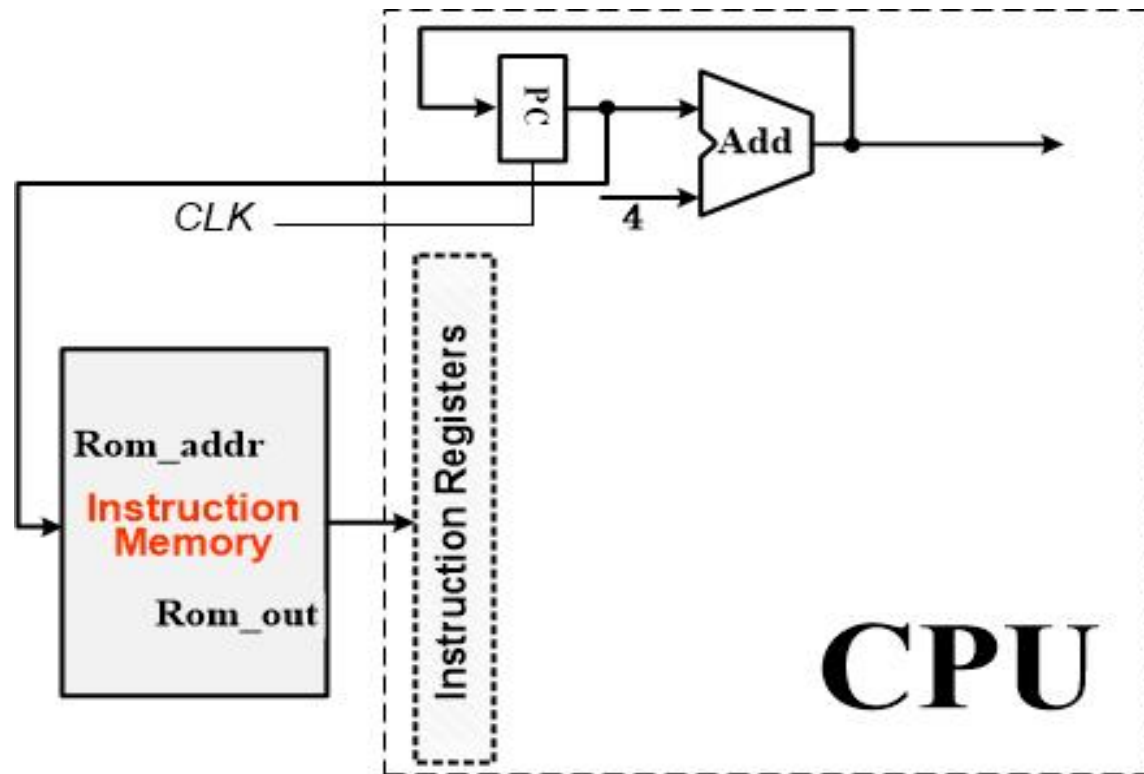
- Instruction Register
 - Can you omit it?





How simple is!

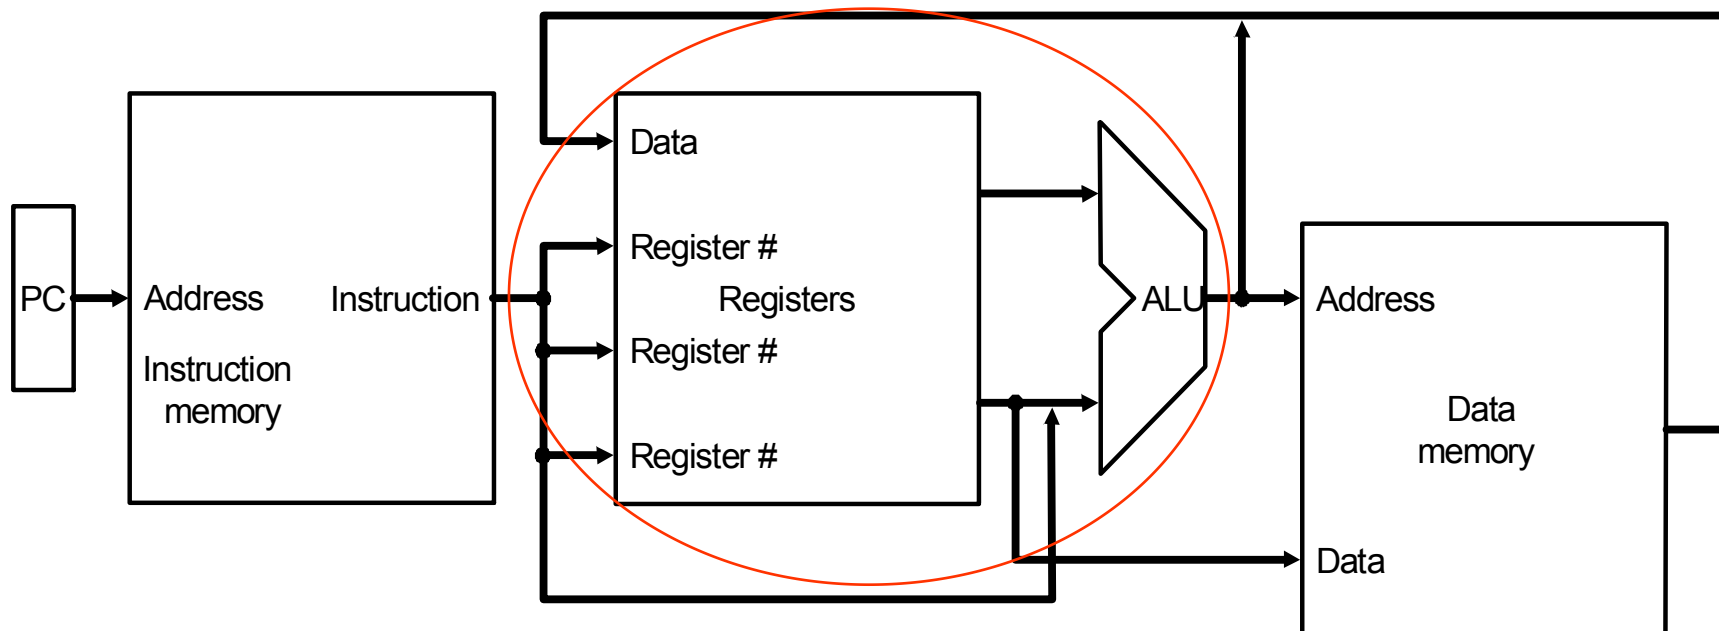
- Why PC+4?





More Implementation Details

- Abstract / Simplified View:





Path Built using Multiplexer

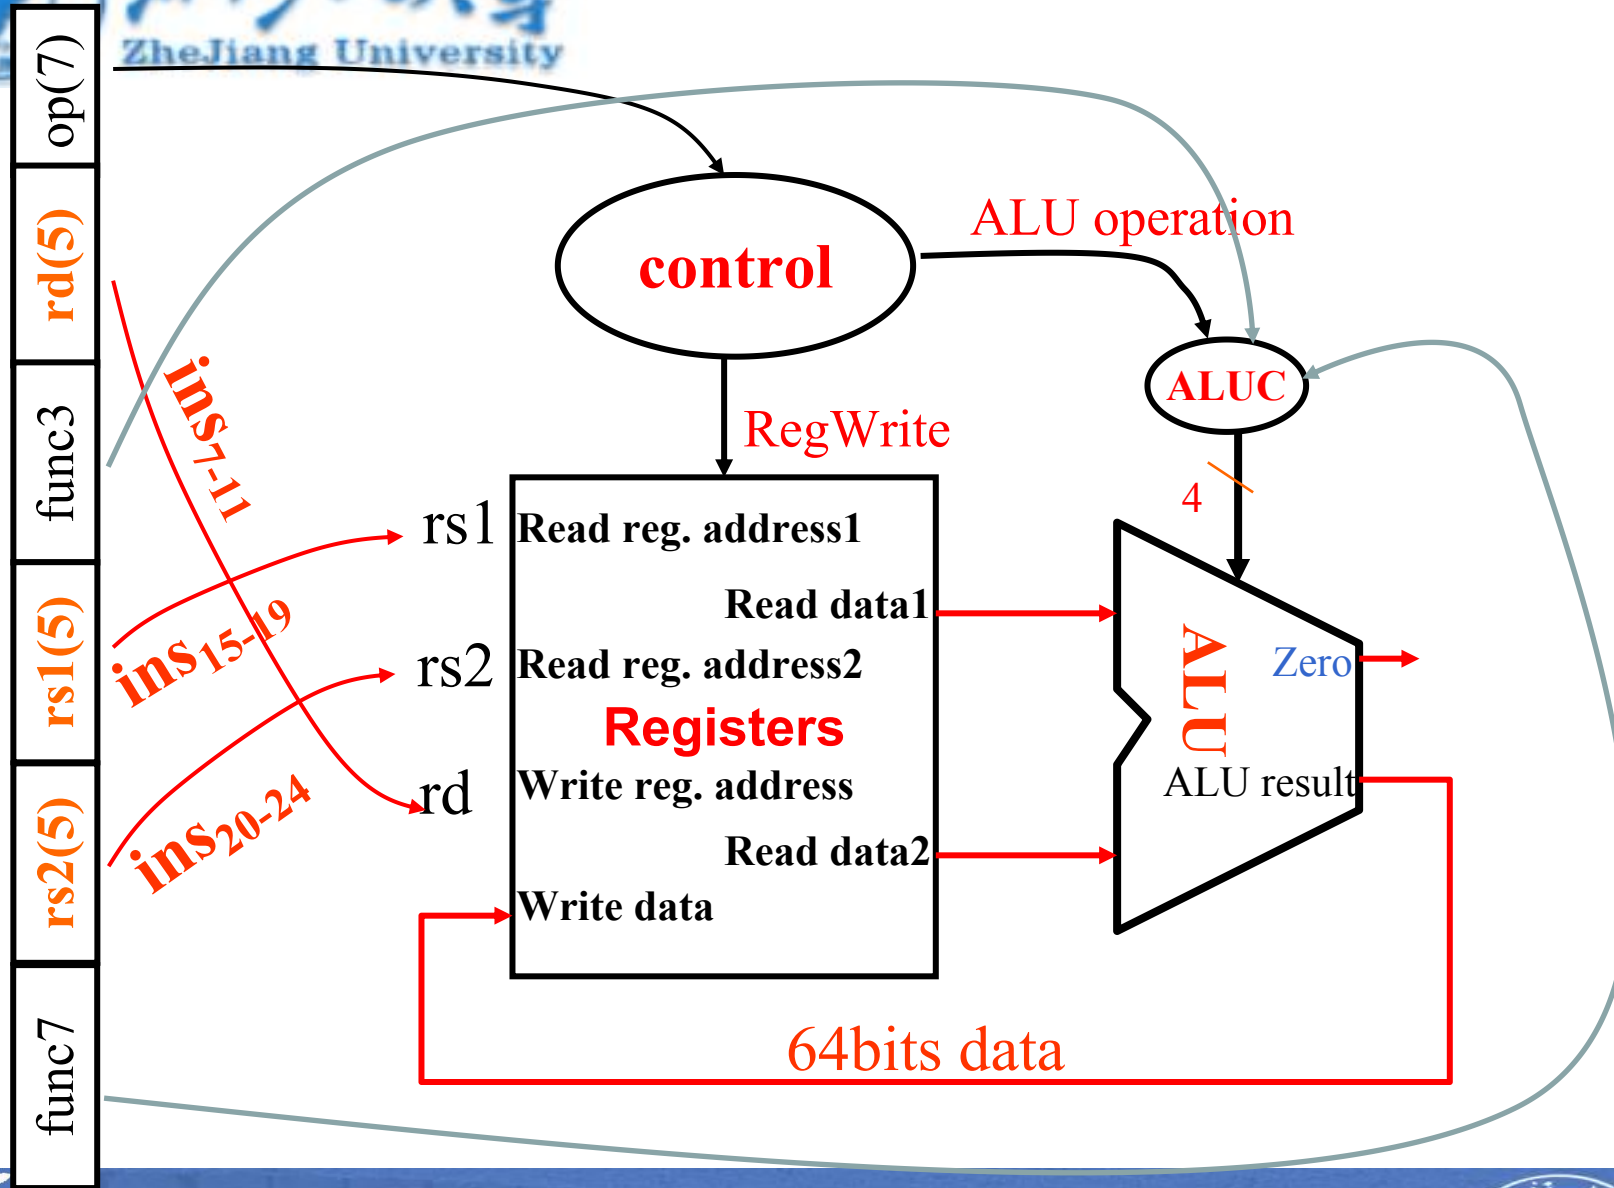
Data Stream of Instruction executing

- R-type instruction Datapath
- I-type instruction Datapath
 - For ALU
 - For load
- S-type (store) instruction Datapath
- SB-type (branch) instruction Datapath
- UJ-type instruction Datapath
 - For Jump
- First, Look at the data flow within instruction execution





R type Instruction & Data stream



add x9, x20, x21

1. Read two register operands
2. Perform arithmetic/logical operation
3. Write register result

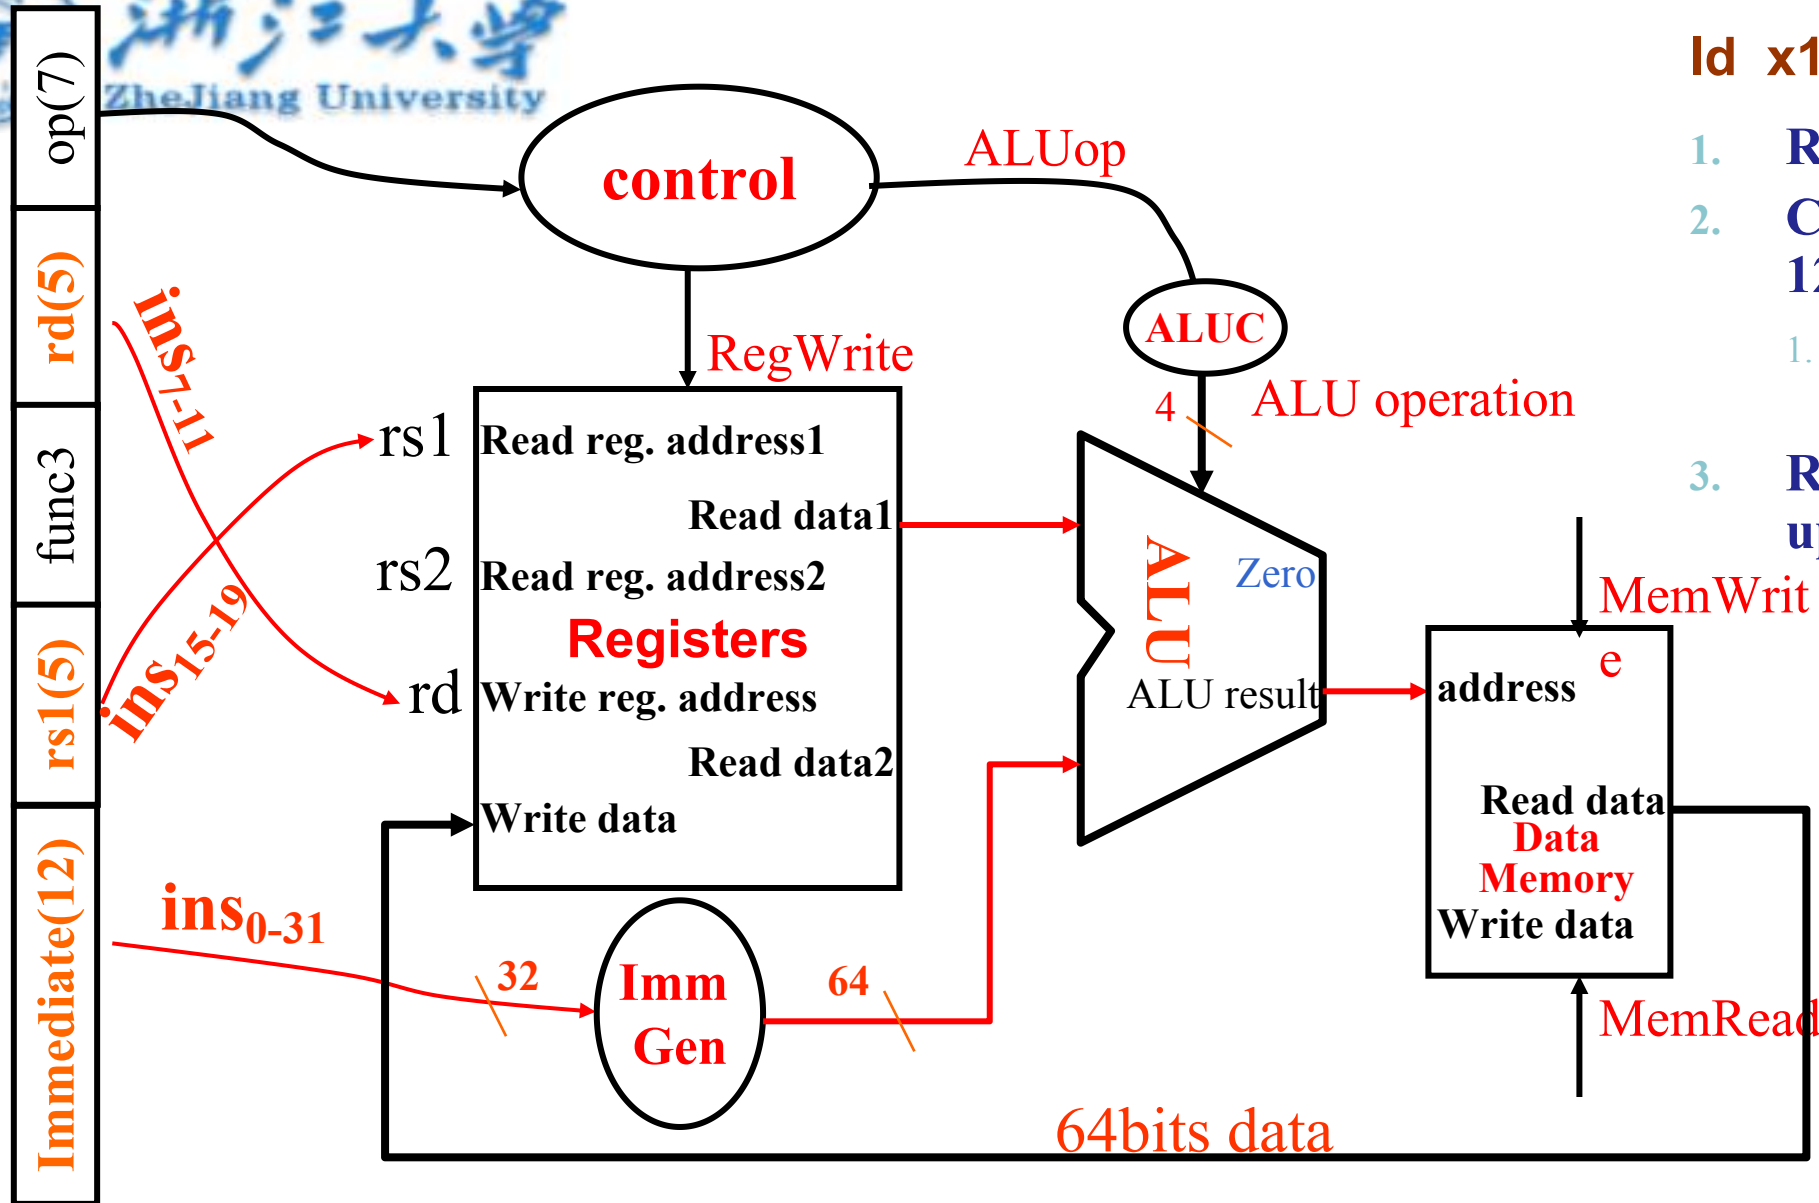
I type (load) Instruction & Data stream



浙江大学
Zhejiang University

ld x1, 200(x2)

1. Read register operands
2. Calculate address using 12-bit offset
 1. Use ALU, but sign-extend offset
3. Read memory and update register



addi x1, x2, 4?

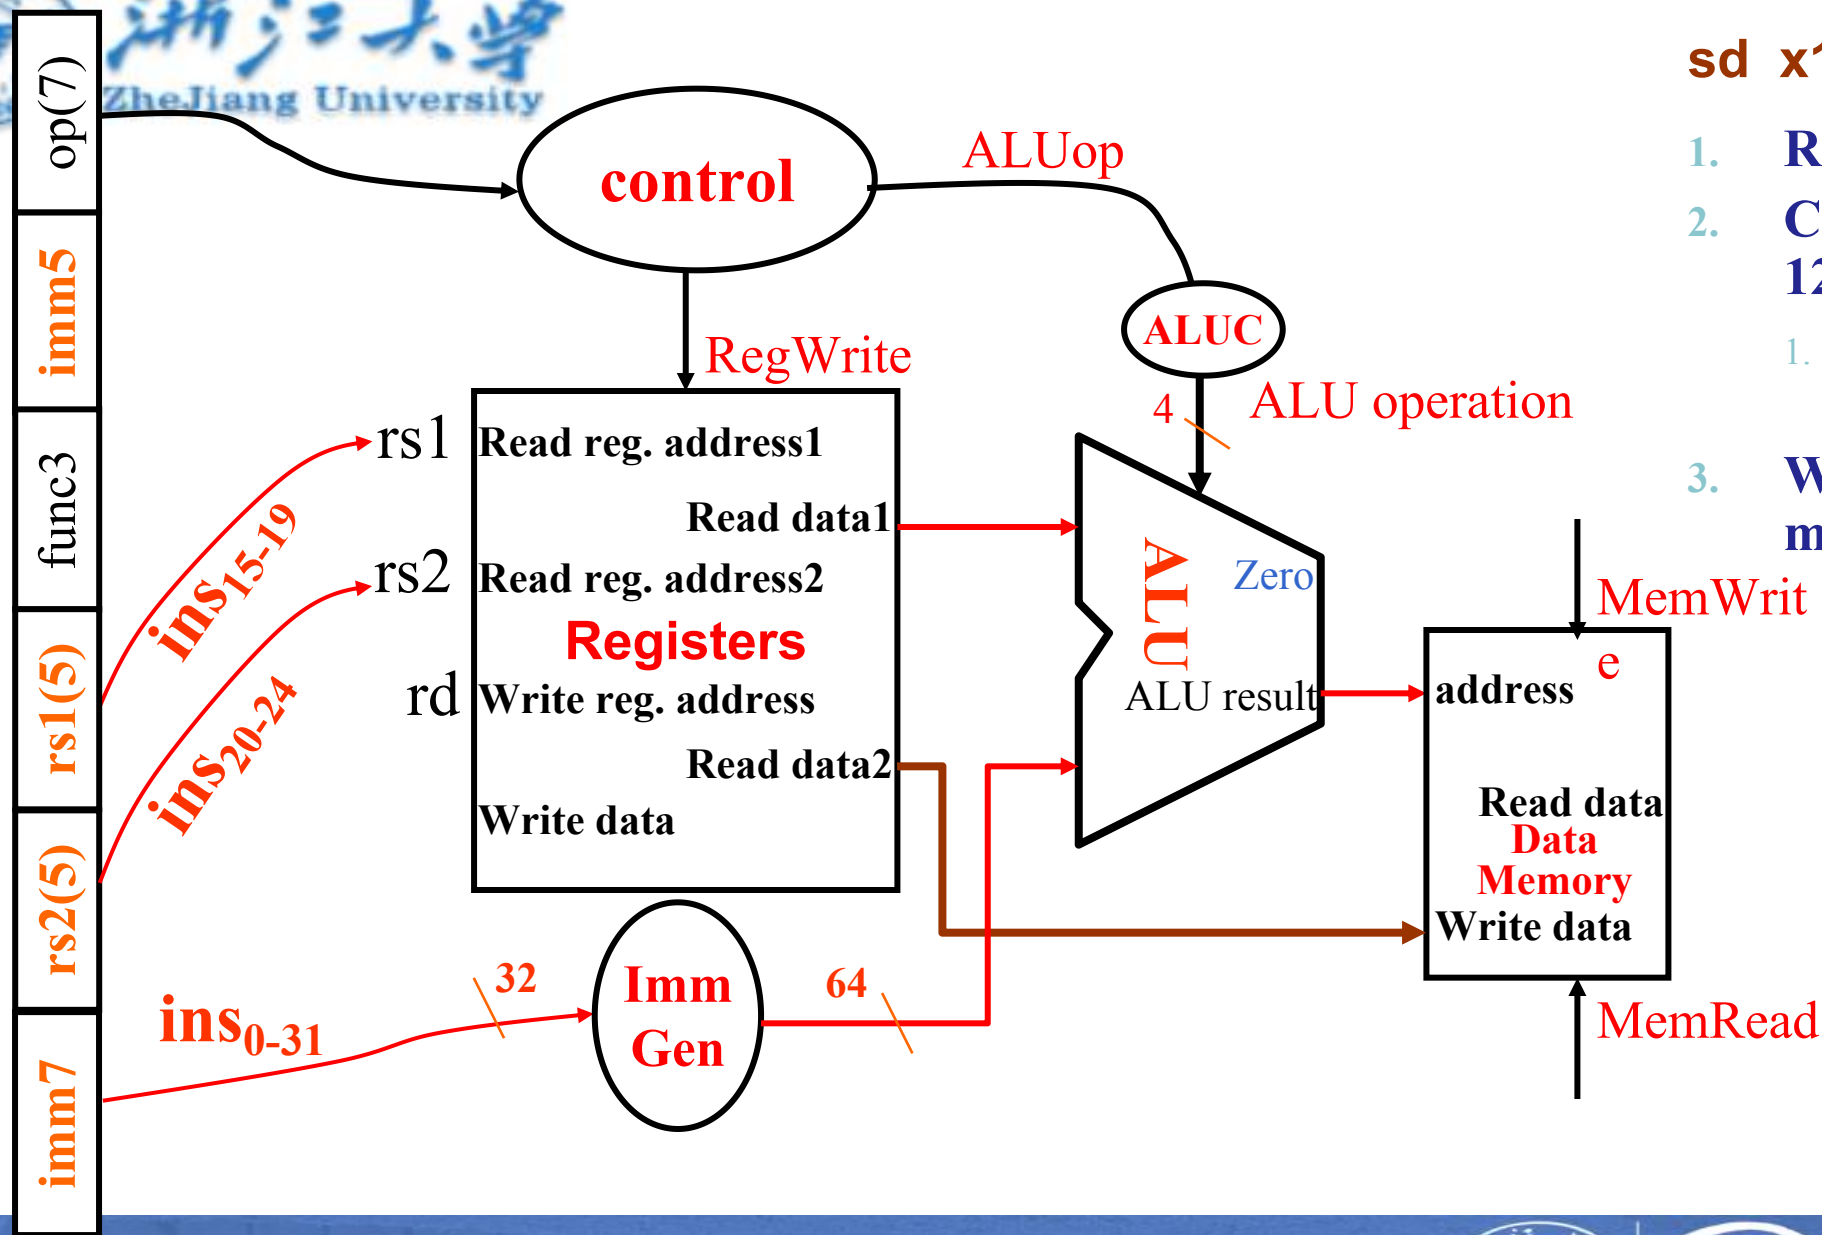




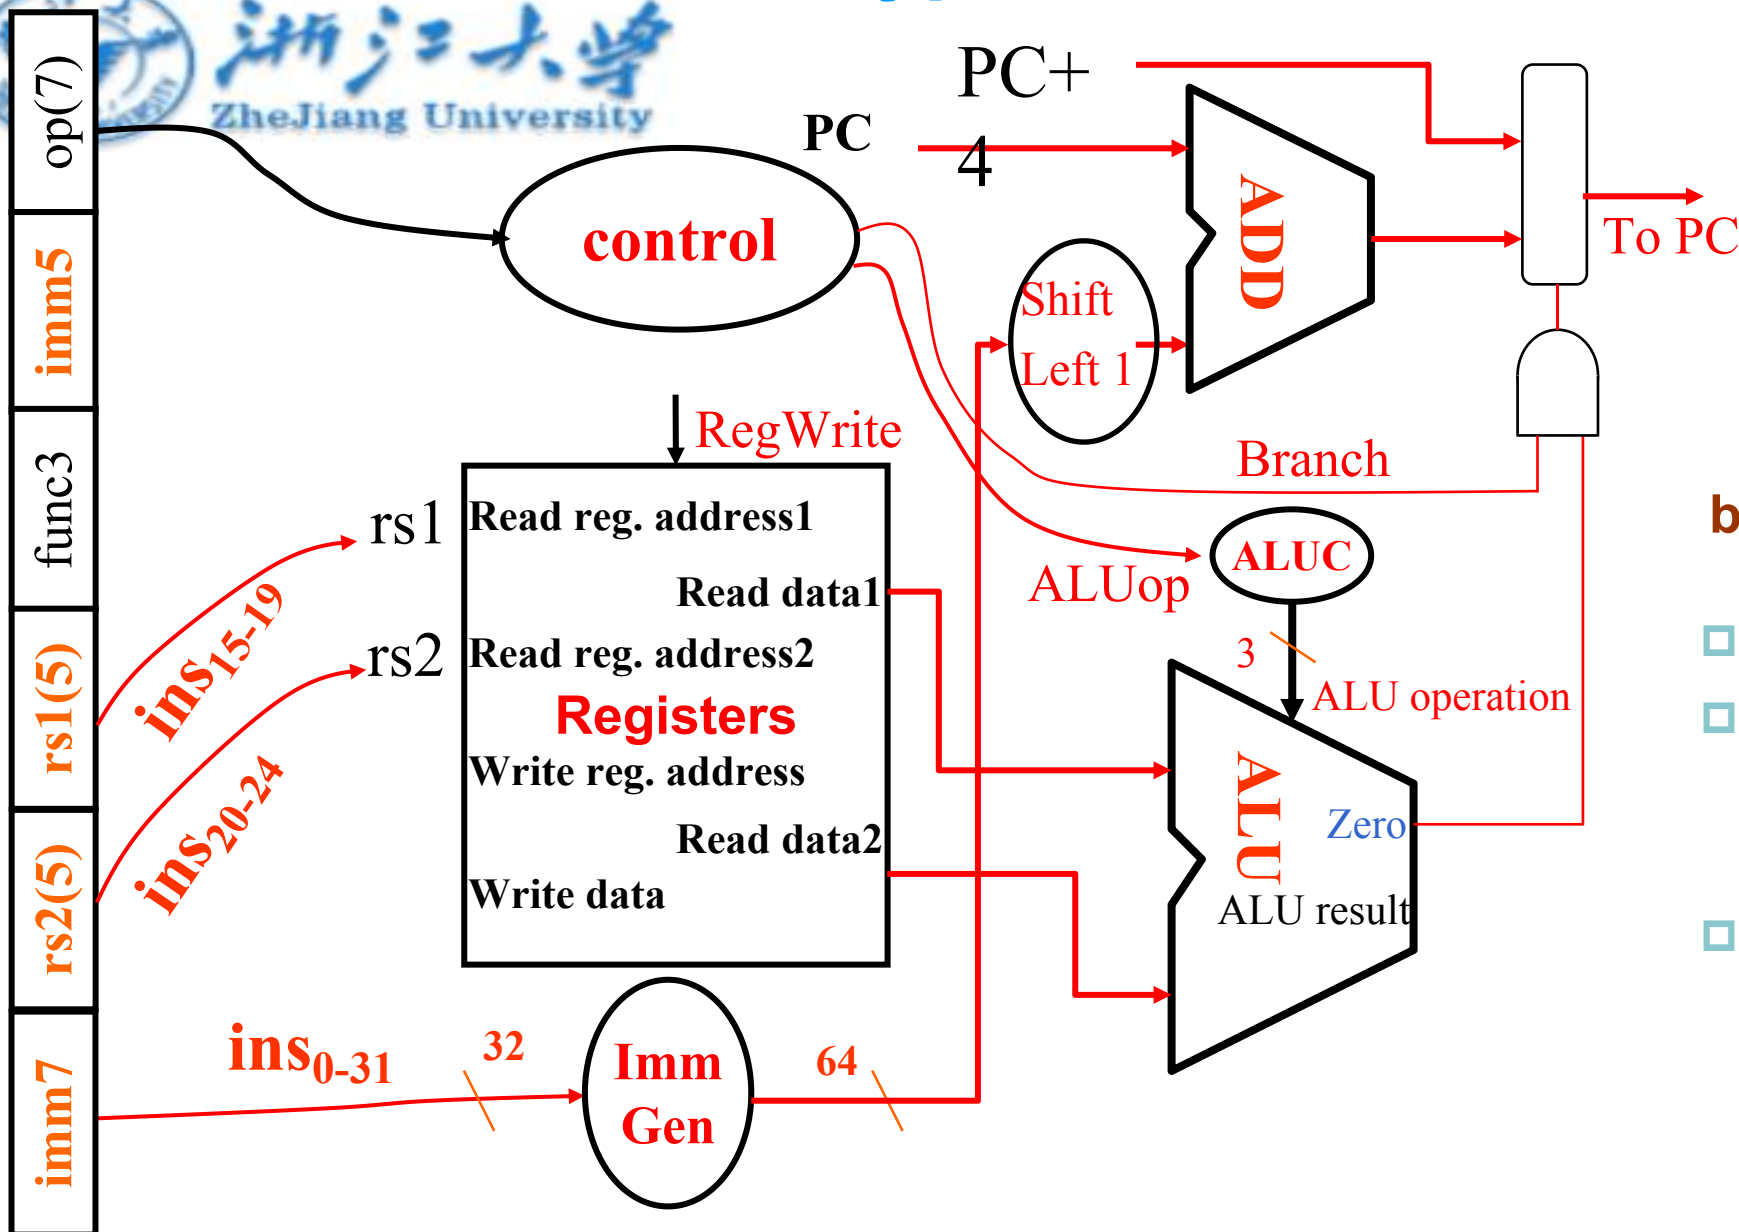
S-type (store) Instruction & Data stream

sd x1, 200(x2)

1. Read register operands
2. Calculate address using 12-bit offset
 1. Use ALU, but sign-extend offset
3. Write register value to memory



B type Instruction & Data stream of *beq*

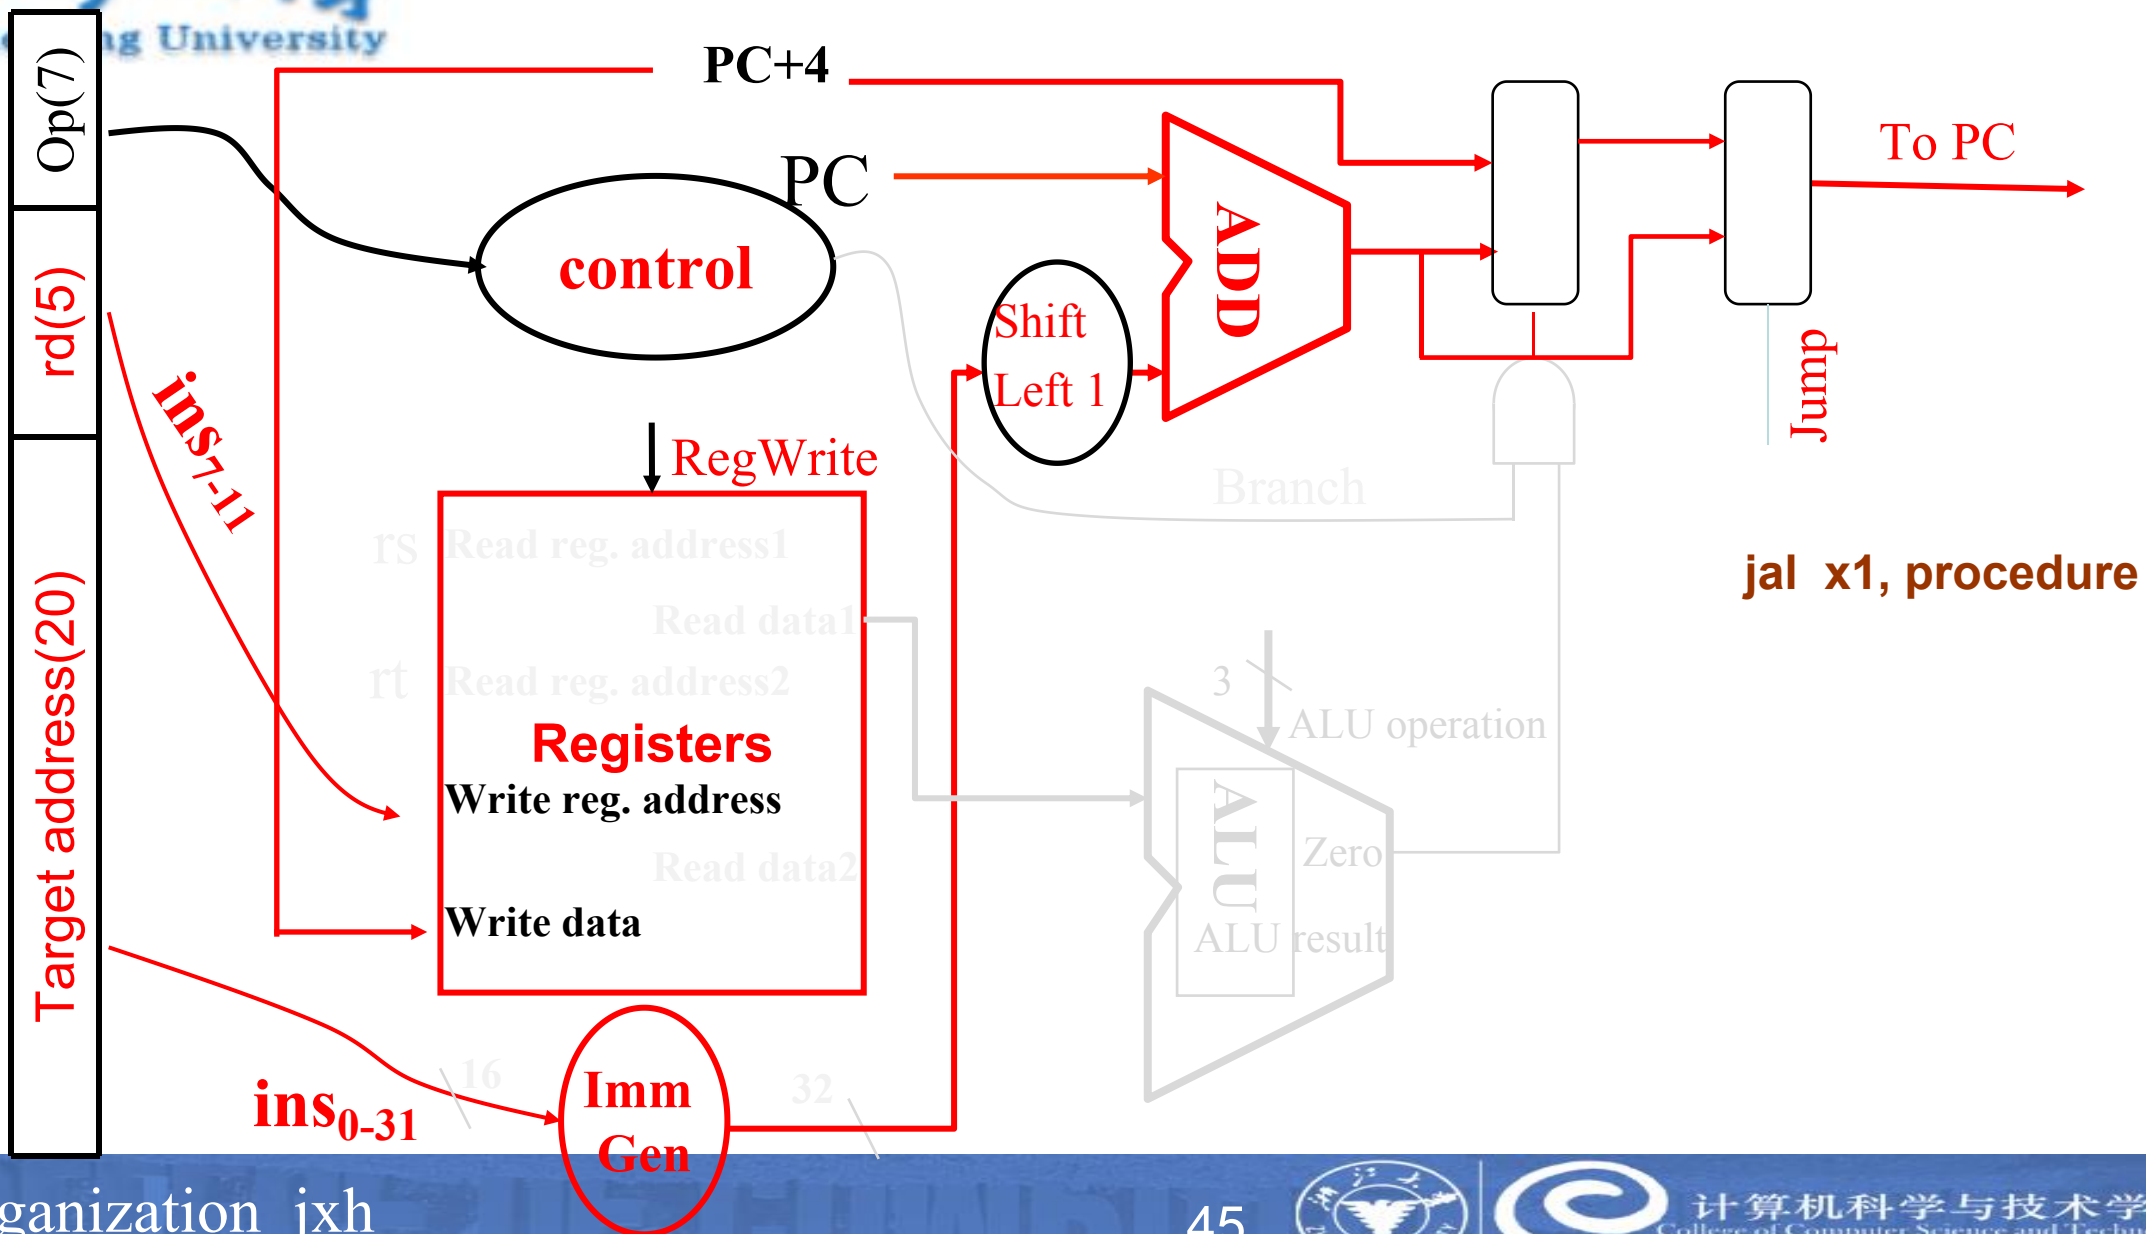


beq x1, x2, 200

- Read register operands
- Compare operands
 - Use ALU, subtract and check Zero output
- Calculate target address
 - Sign-extend displacement
 - Shift left 1 place (halfword displacement)
 - Add to PC value



J type Instruction





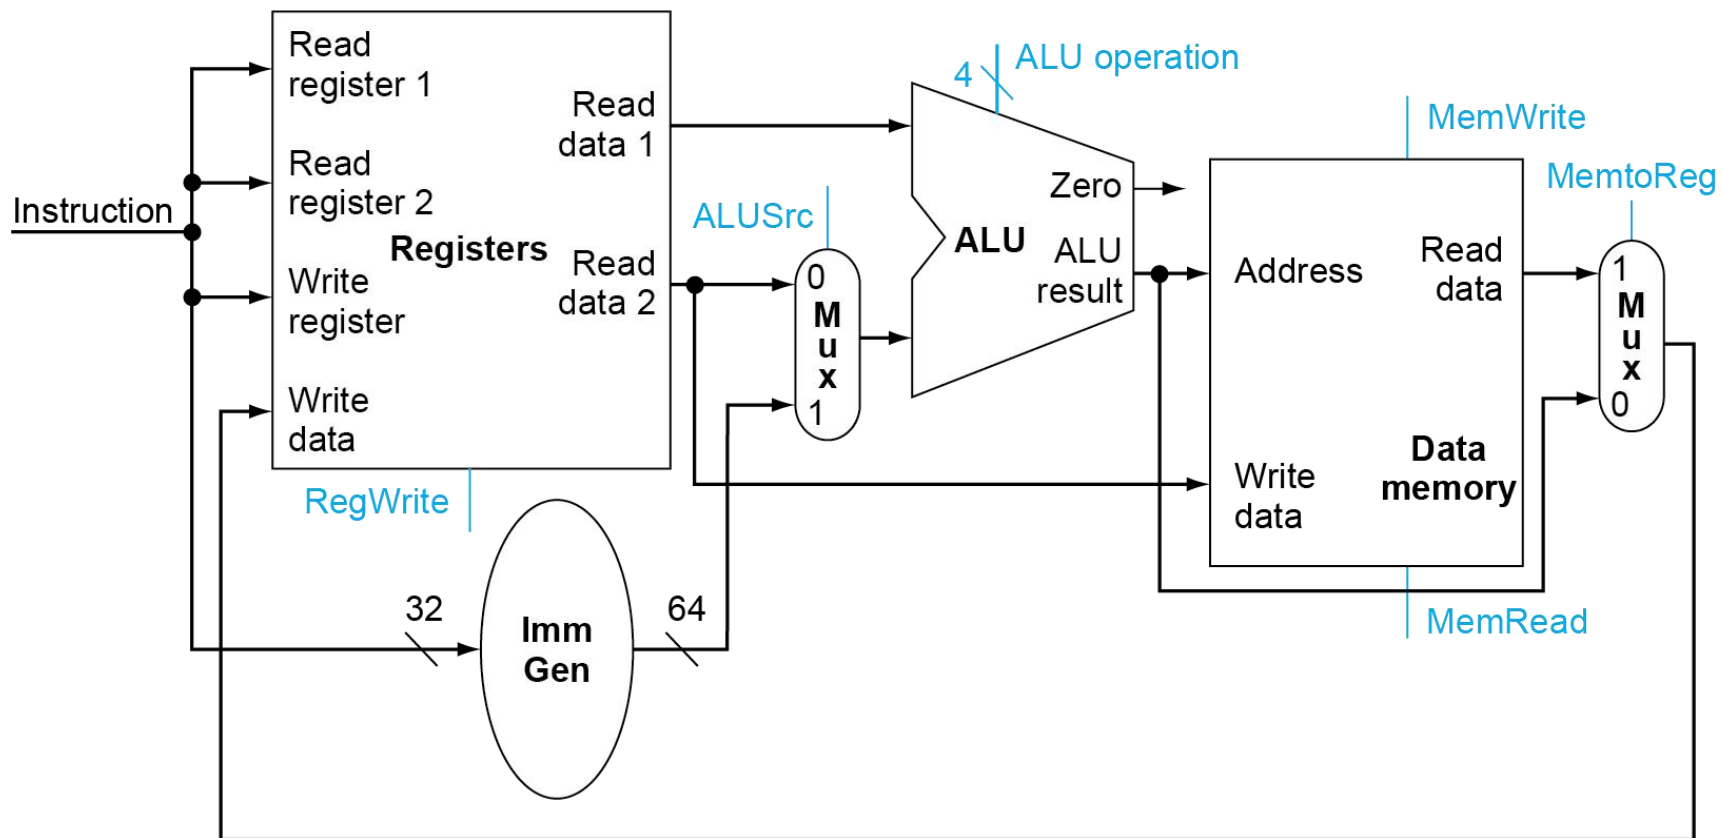
Composing the Elements

- First-cut data path does an instruction in one clock cycle
 - Each datapath element can only do one function at a time
 - Hence, we need separate instruction and data memories
- Use multiplexers where alternate data sources are used for different instructions



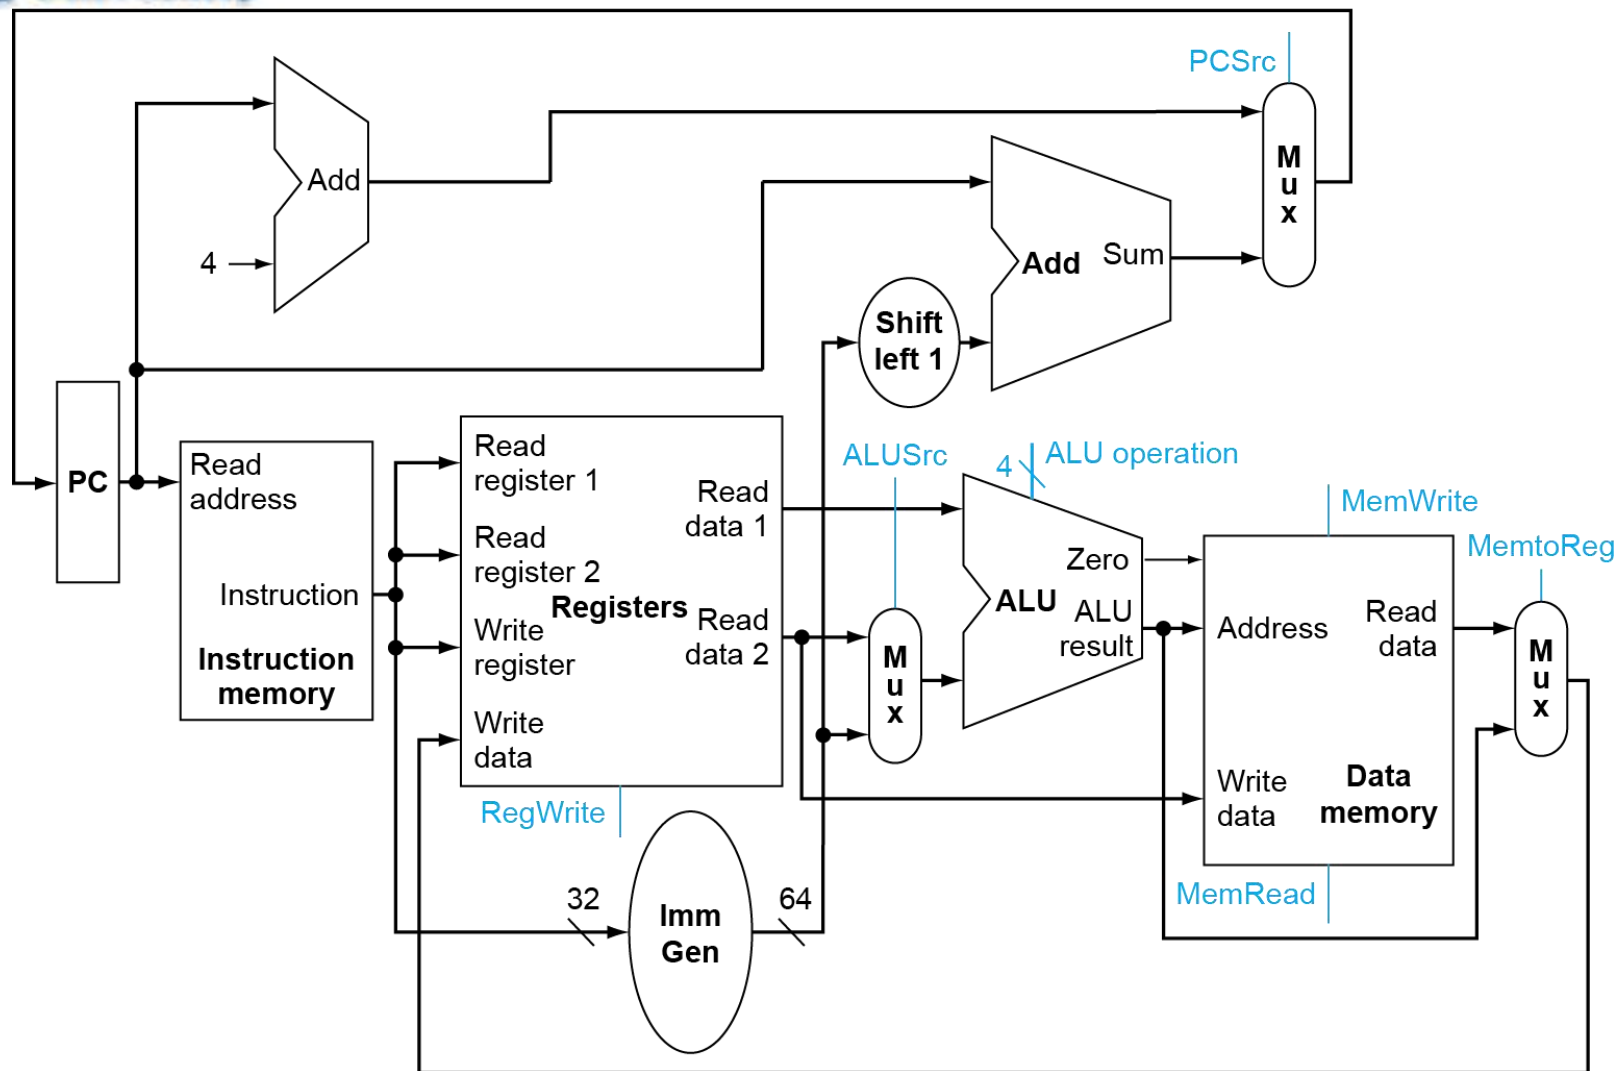


R-Type/Load/Store Datapath





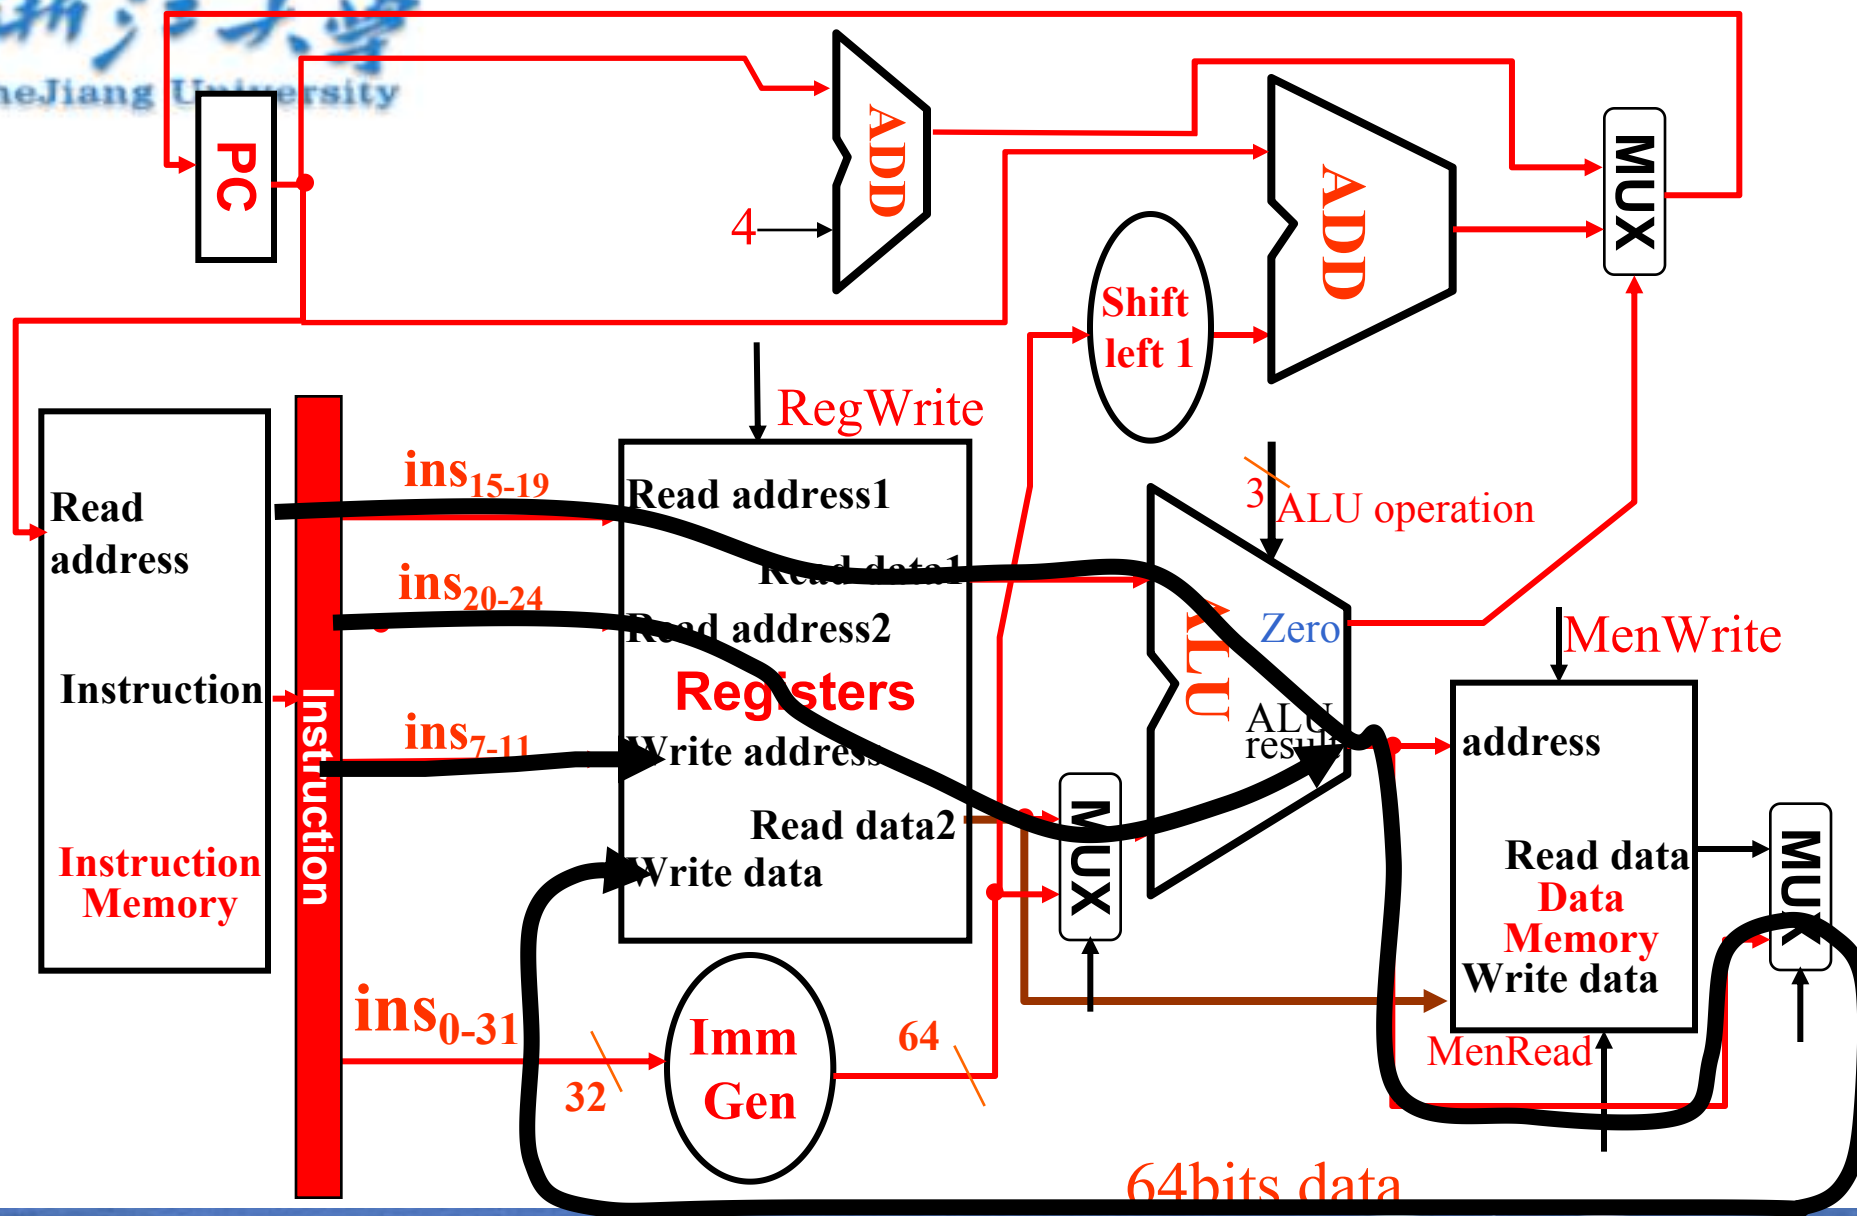
Full Datapath





Full datapath

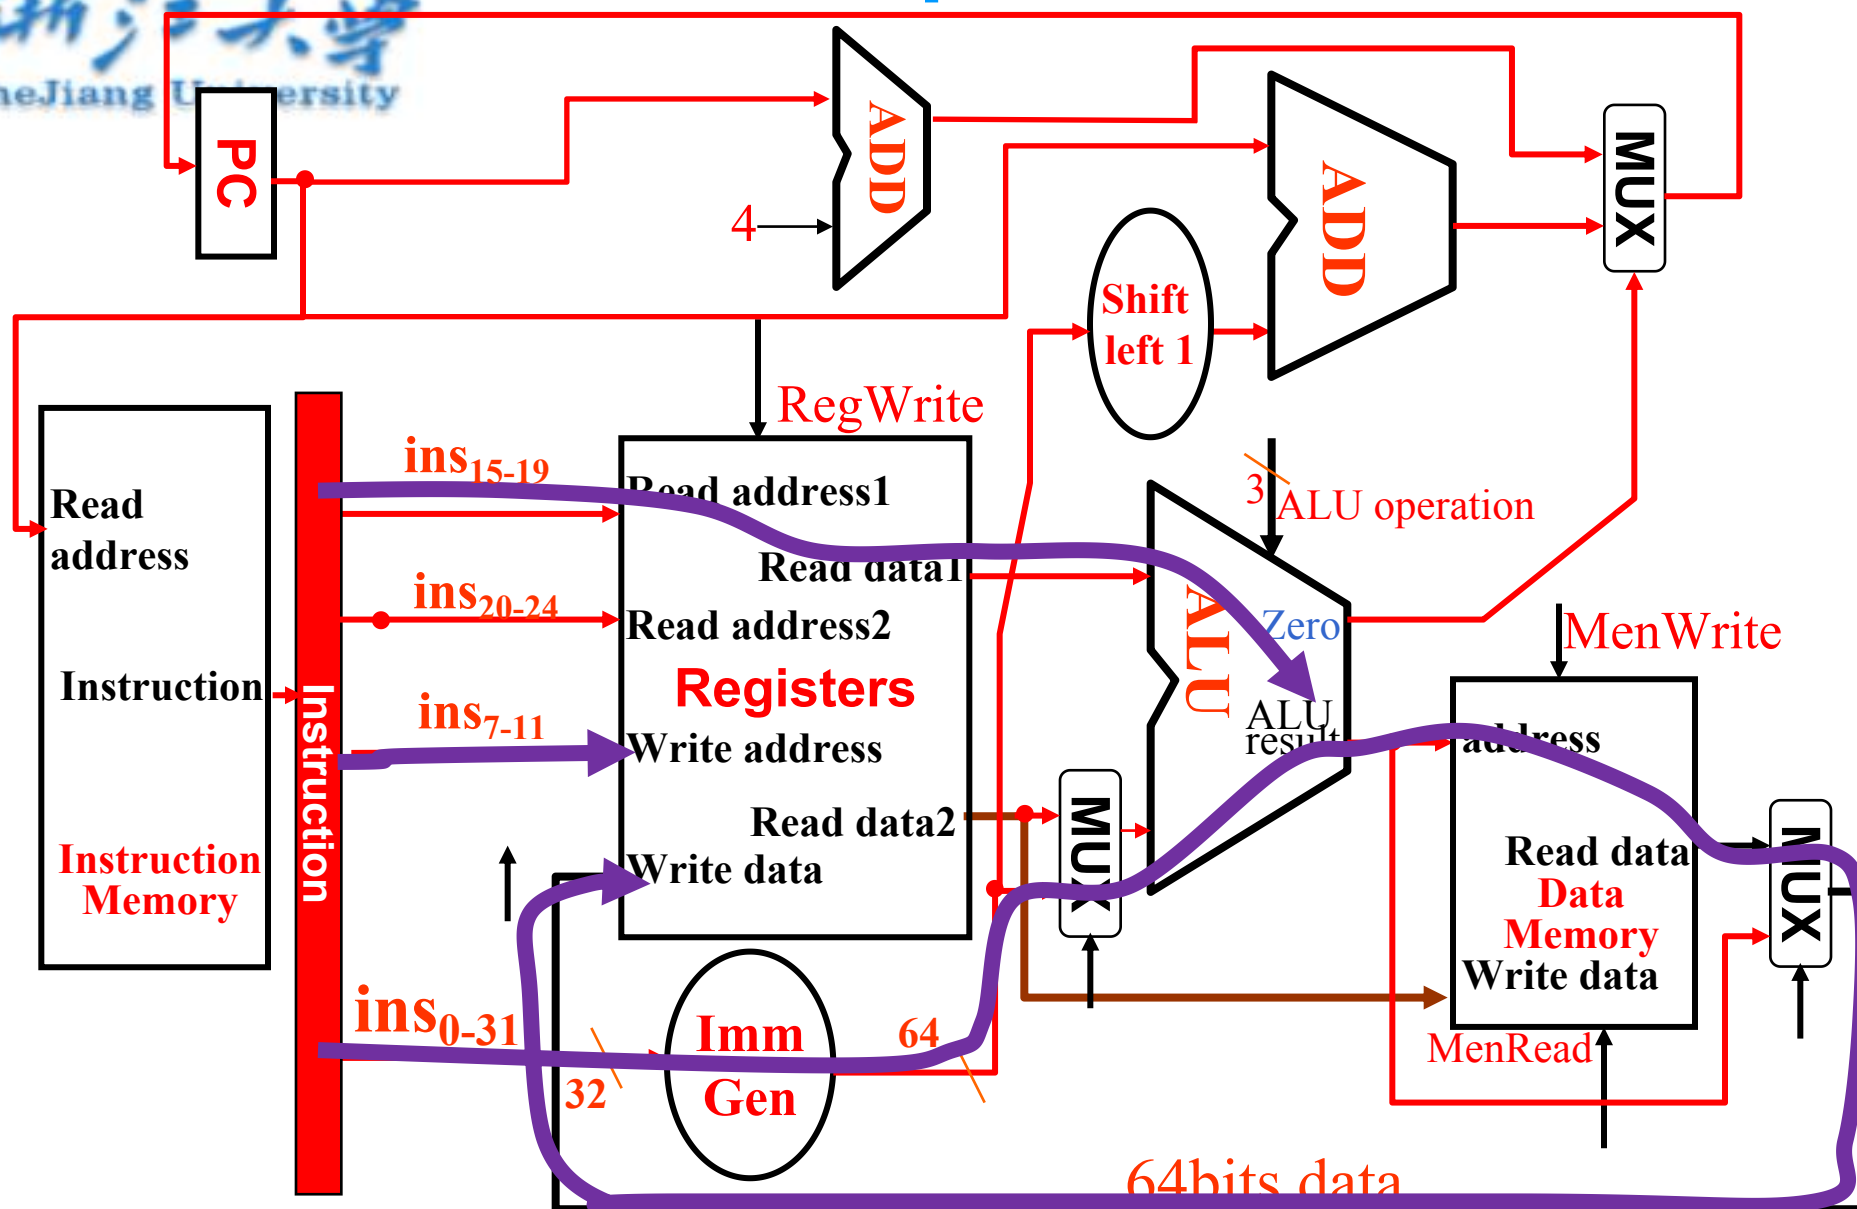
R





Full datapath

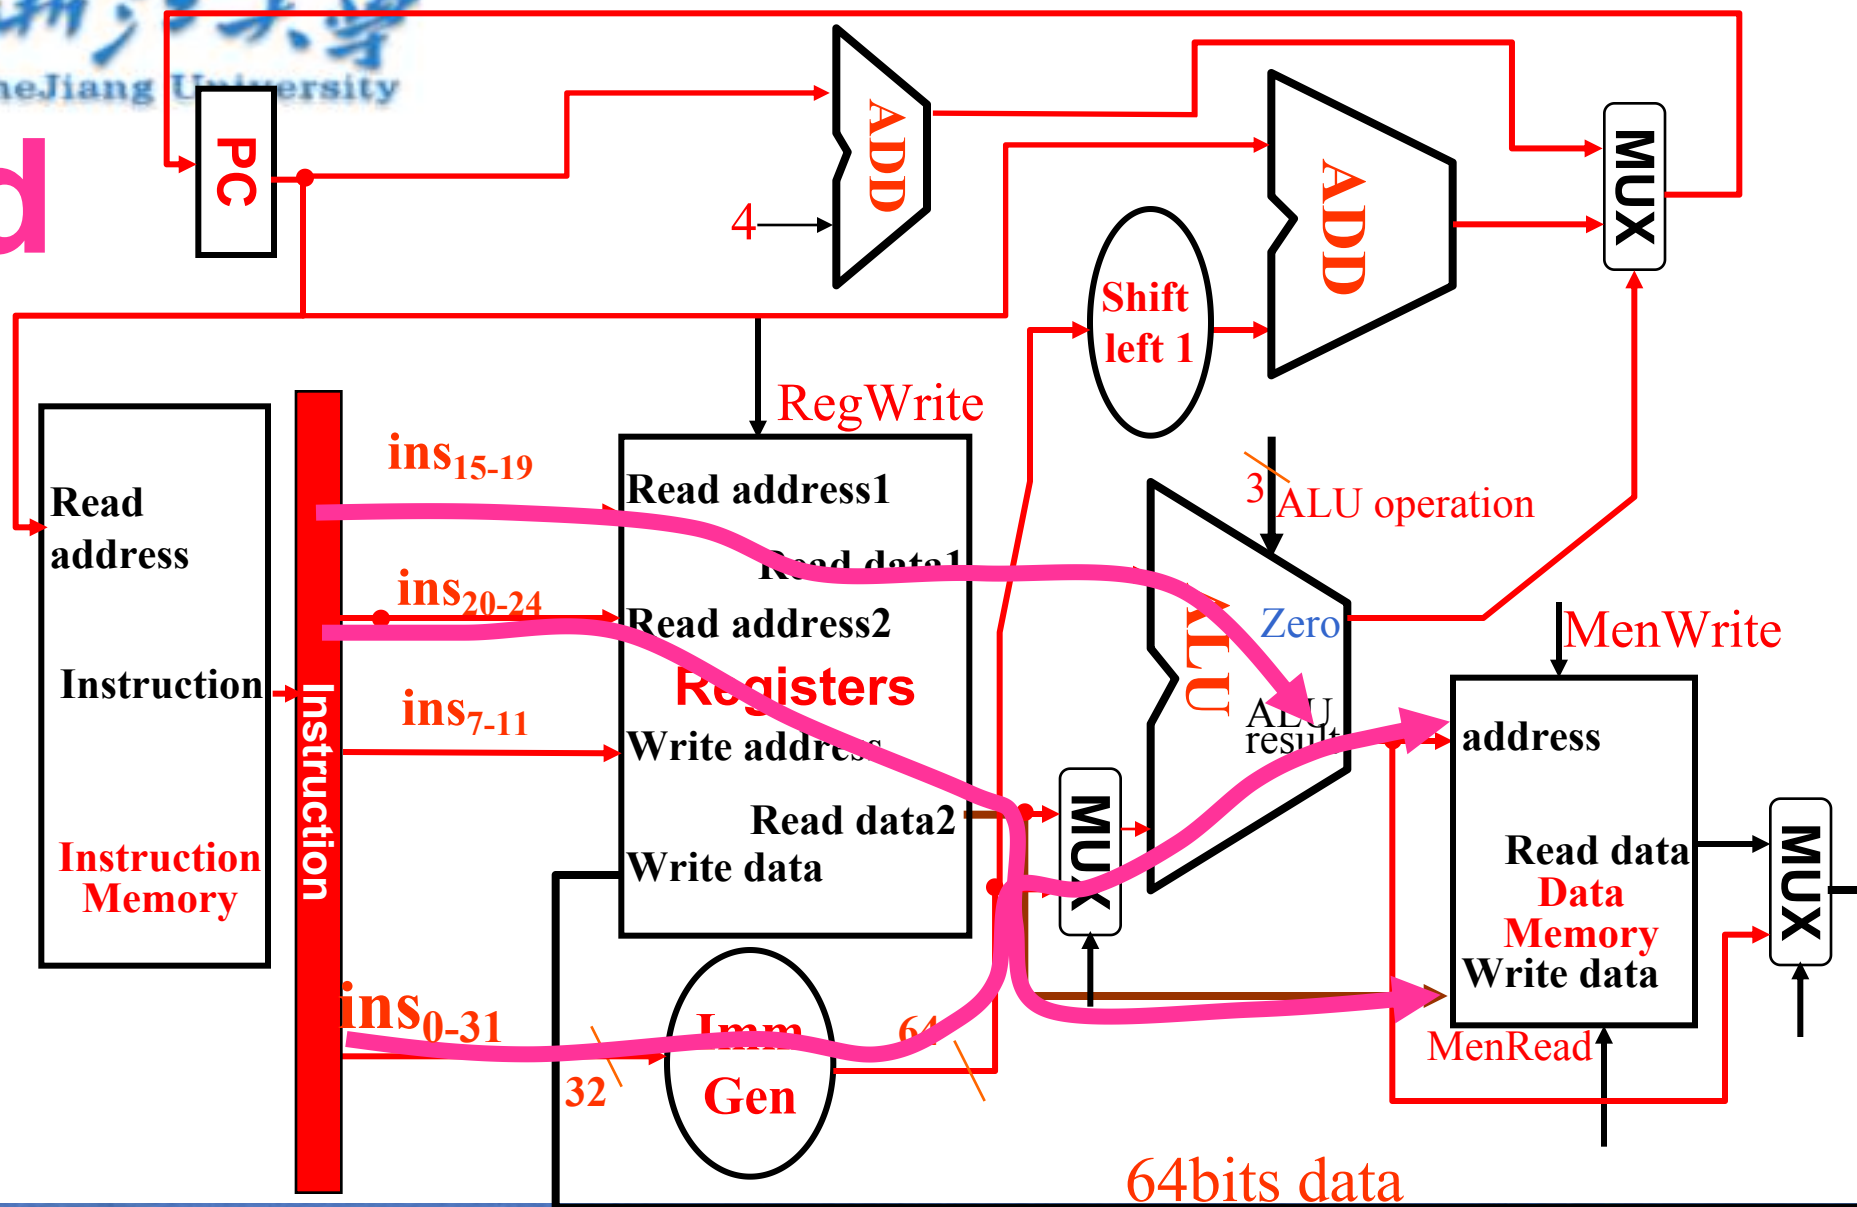
I-Id





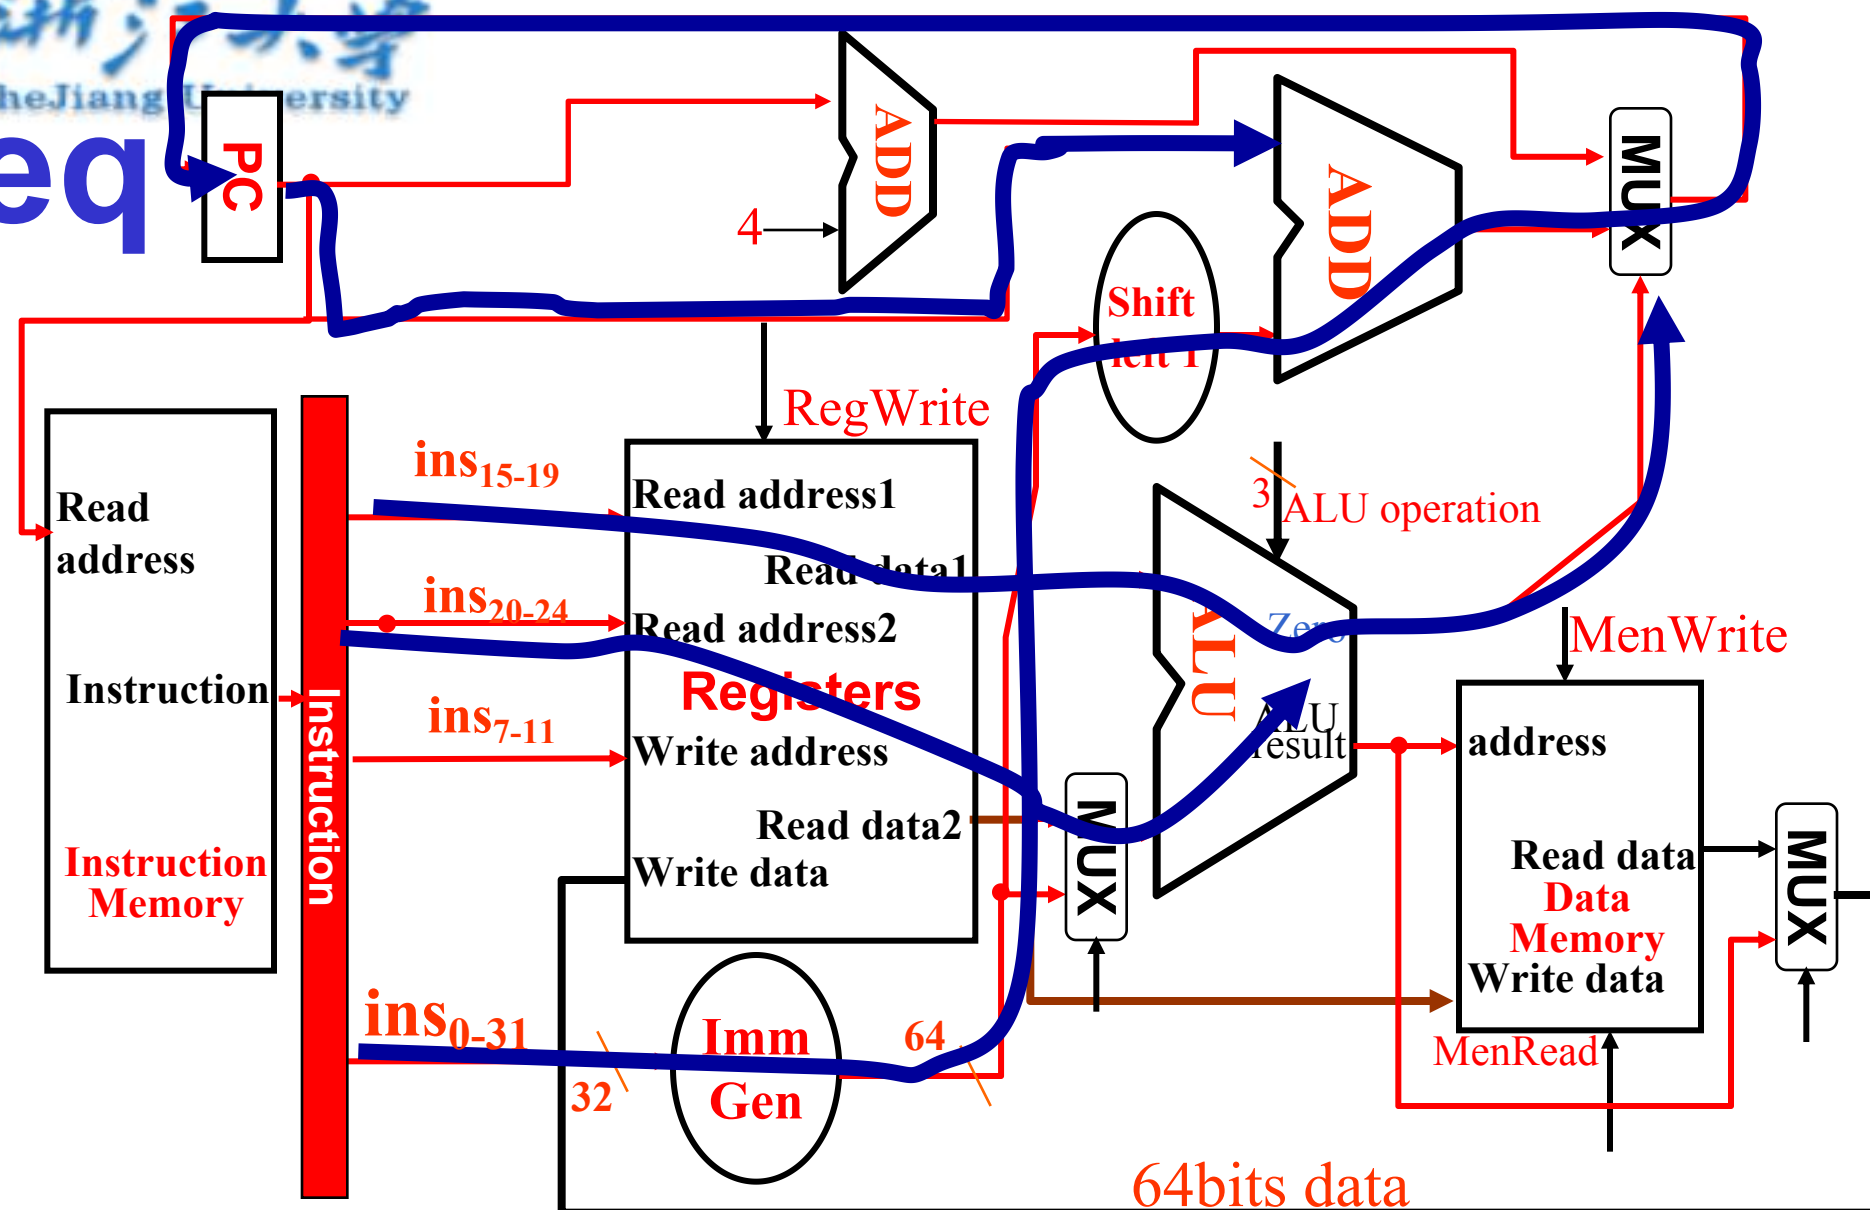
S-sd

Full datapath



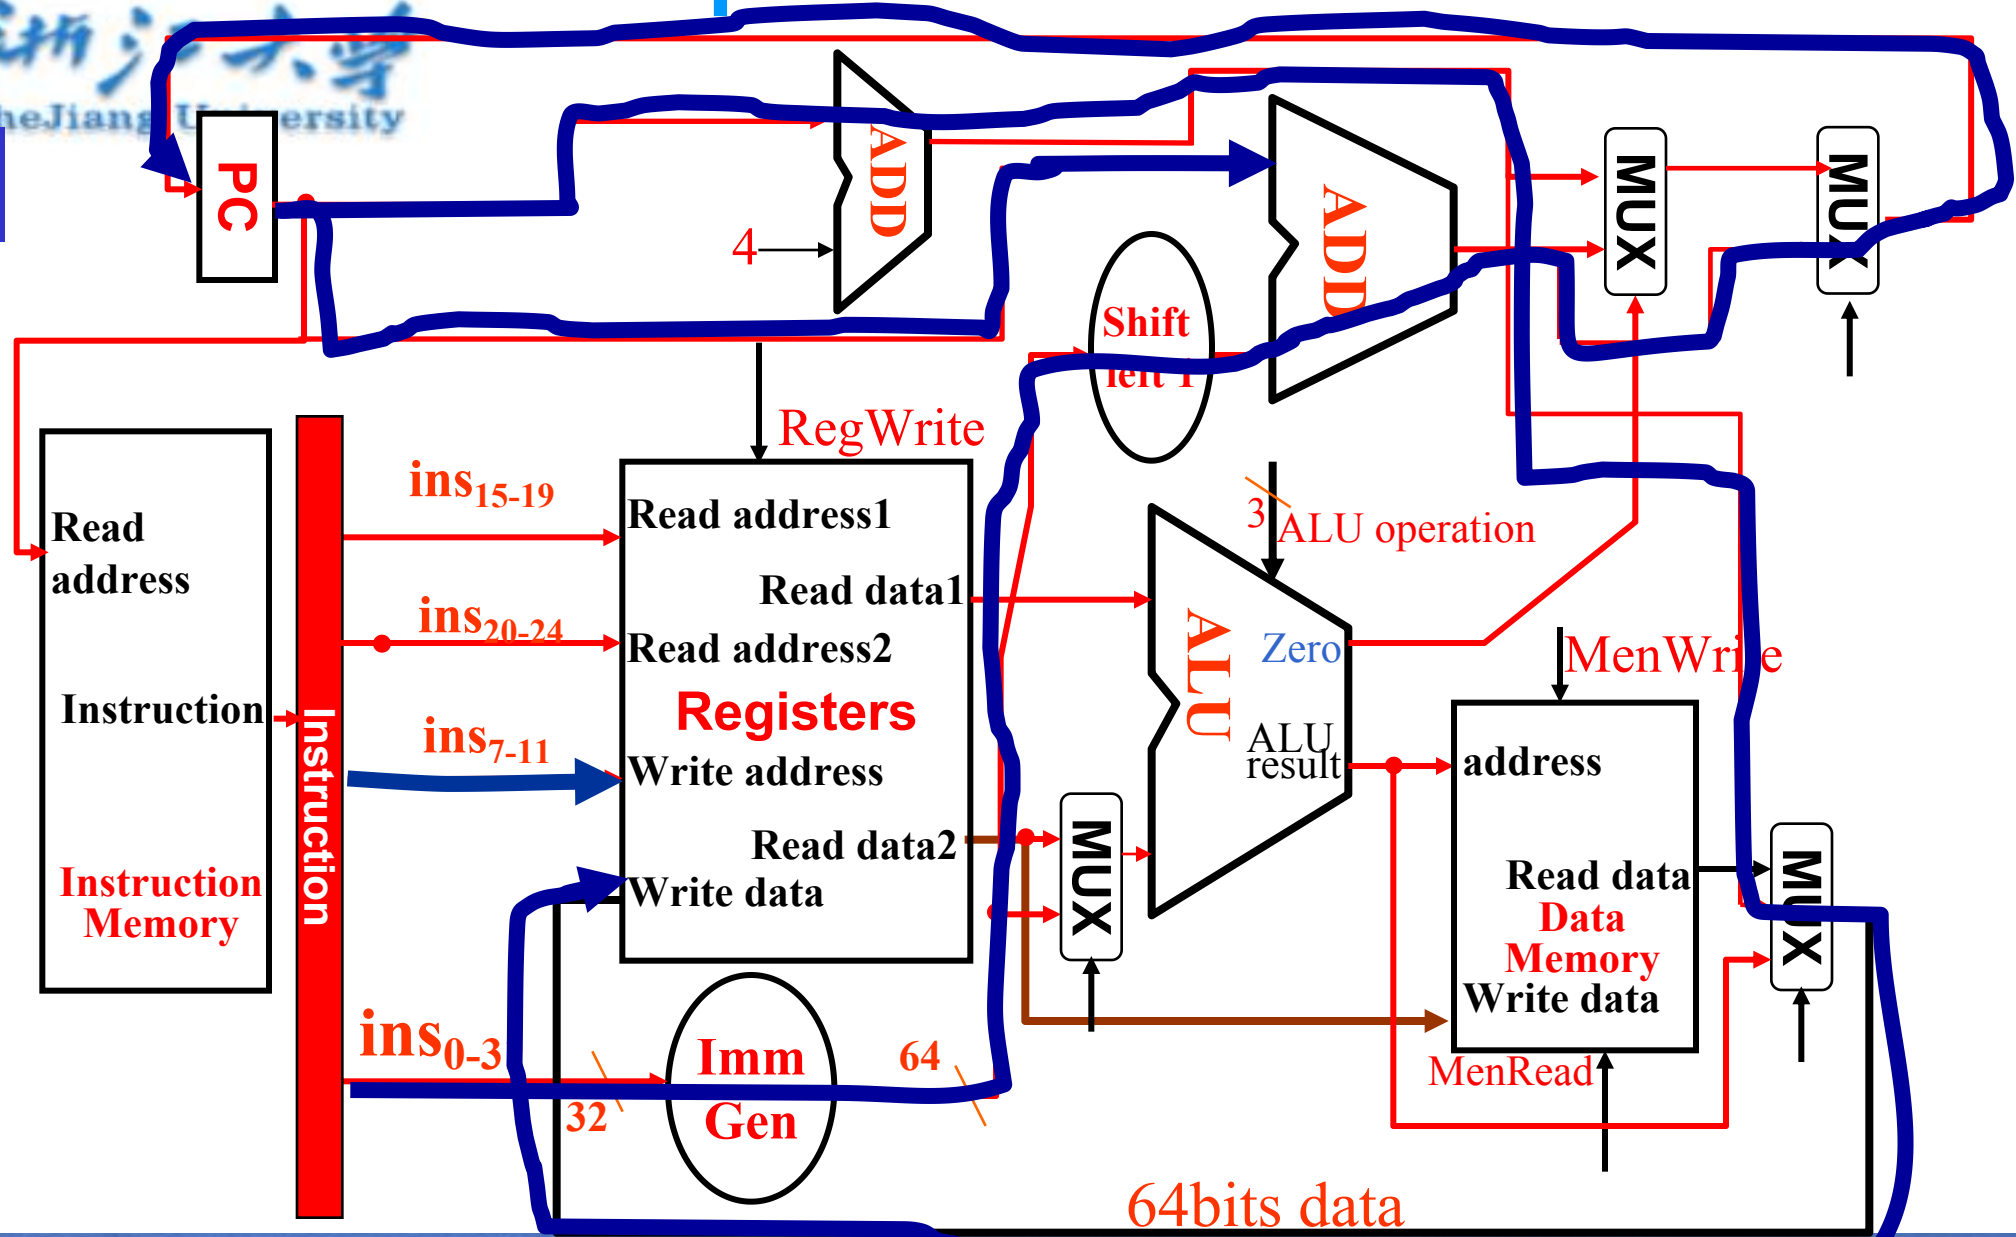
Full datapath

B-beq



Full datapath

J-jal





Contents

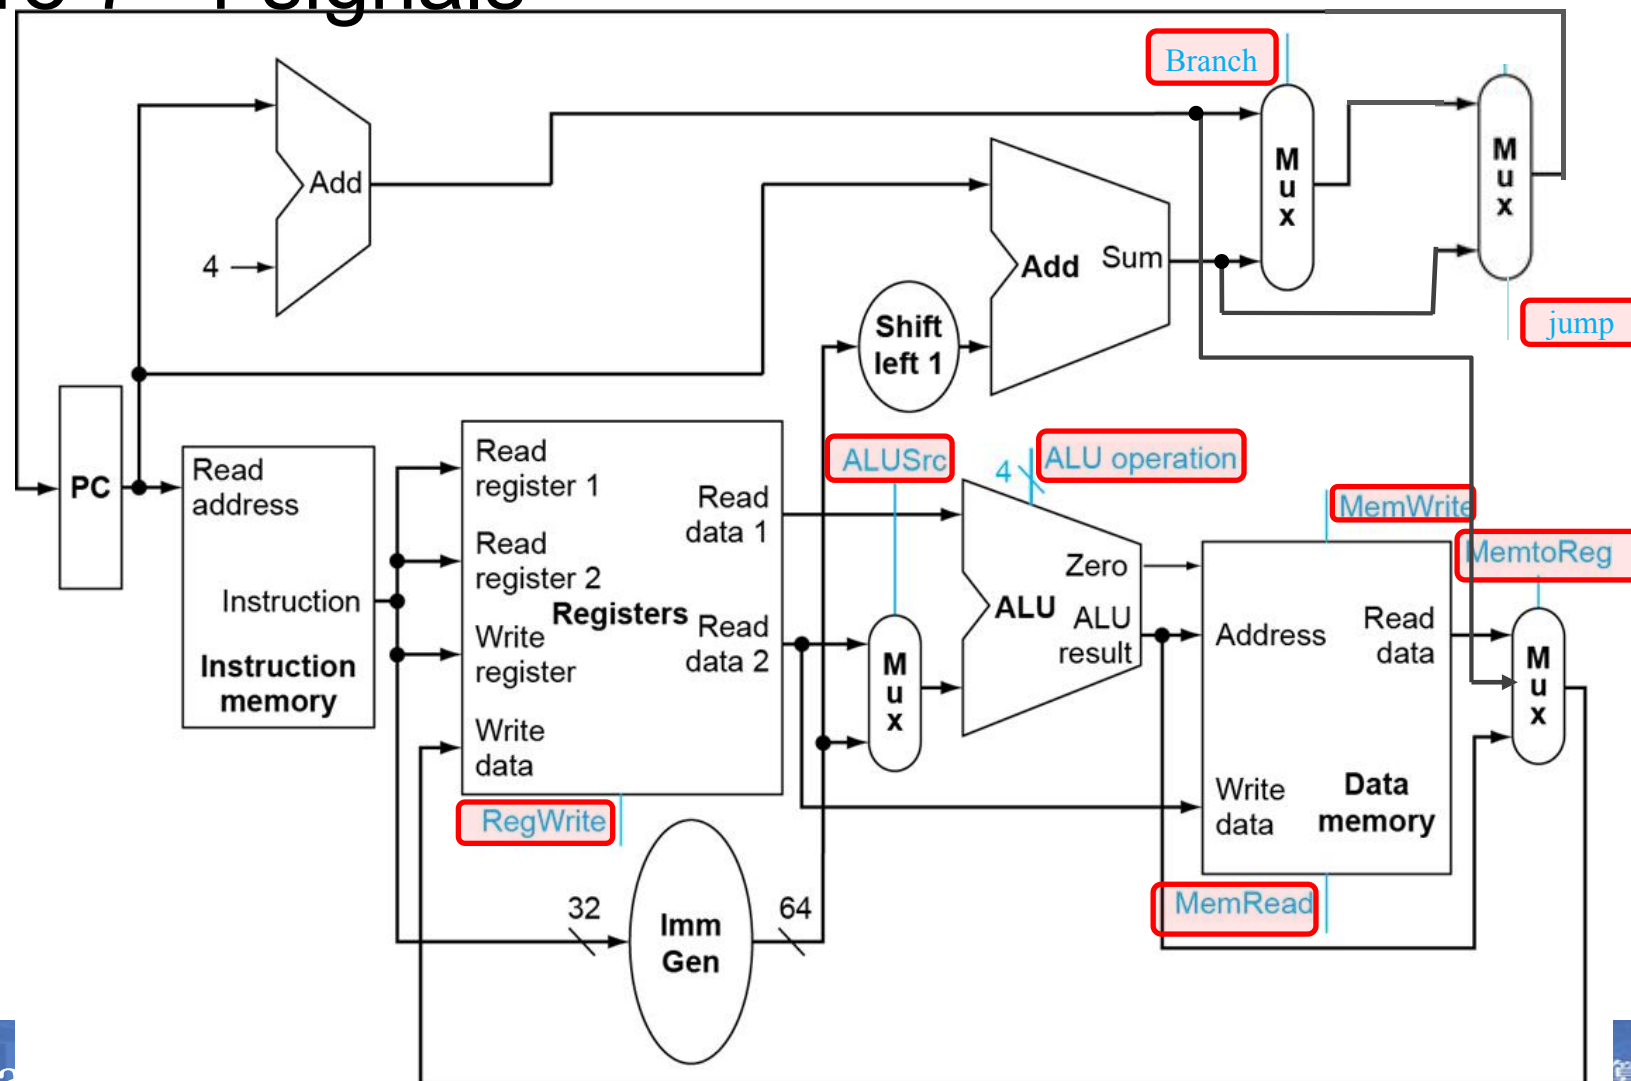
- Introduction & Logic Design Conventions
- Building a datapath
- A Simple Implementation Scheme
- Pipelining





Building the Datapath & Controller

- There are 7+4 signals





Building Controller

Analyse for cause and effect

- **Information** comes from the 32 bits of the instruction
- Selecting the **operations** to perform (ALU, read/write, etc.)
- Controlling the **flow of data** (multiplexor inputs)
- ALU's operation based on **instruction type** and **function** code

| Name | Field | | | | | | Comments |
|--------------|-----------------------------|--------|--------|--------|---------------|--------|-------------------------------|
| (Field Size) | 7 bits | 5 bits | 5 bits | 3 bits | 5 bits | 7 bits | |
| R-type | funct7 | rs2 | rs1 | funct3 | rd | opcode | Arithmetic instruction format |
| I-type | immediate[11:0] | | rs1 | funct3 | rd | opcode | Loads & immediate arithmetic |
| S-type | immed[11:5] | rs2 | rs1 | funct3 | immed[4:0] | opcode | Stores |
| SB-type | immed[12,10:5] | rs2 | rs1 | funct3 | immed[4:1,11] | opcode | Conditional branch format |
| UJ-type | immediate[20,10:1,11,19:12] | | | | rd | opcode | Unconditional jump format |
| U-type | immediate[31:12] | | | | rd | opcode | Upper immediate format |





What should ALU do ?

- ALU used for
 - Load/Store: F = add
 - Branch: F = subtract
 - R-type: F depends on opcode
- Assume 2-bit ALUOp derived from opcode
 - Combinational logic derives ALU control

| Operation | Function |
|-----------|----------|
| 000 | And |
| 001 | Or |
| 010 | Add |
| 110 | Sub |
| 111 | Slt |

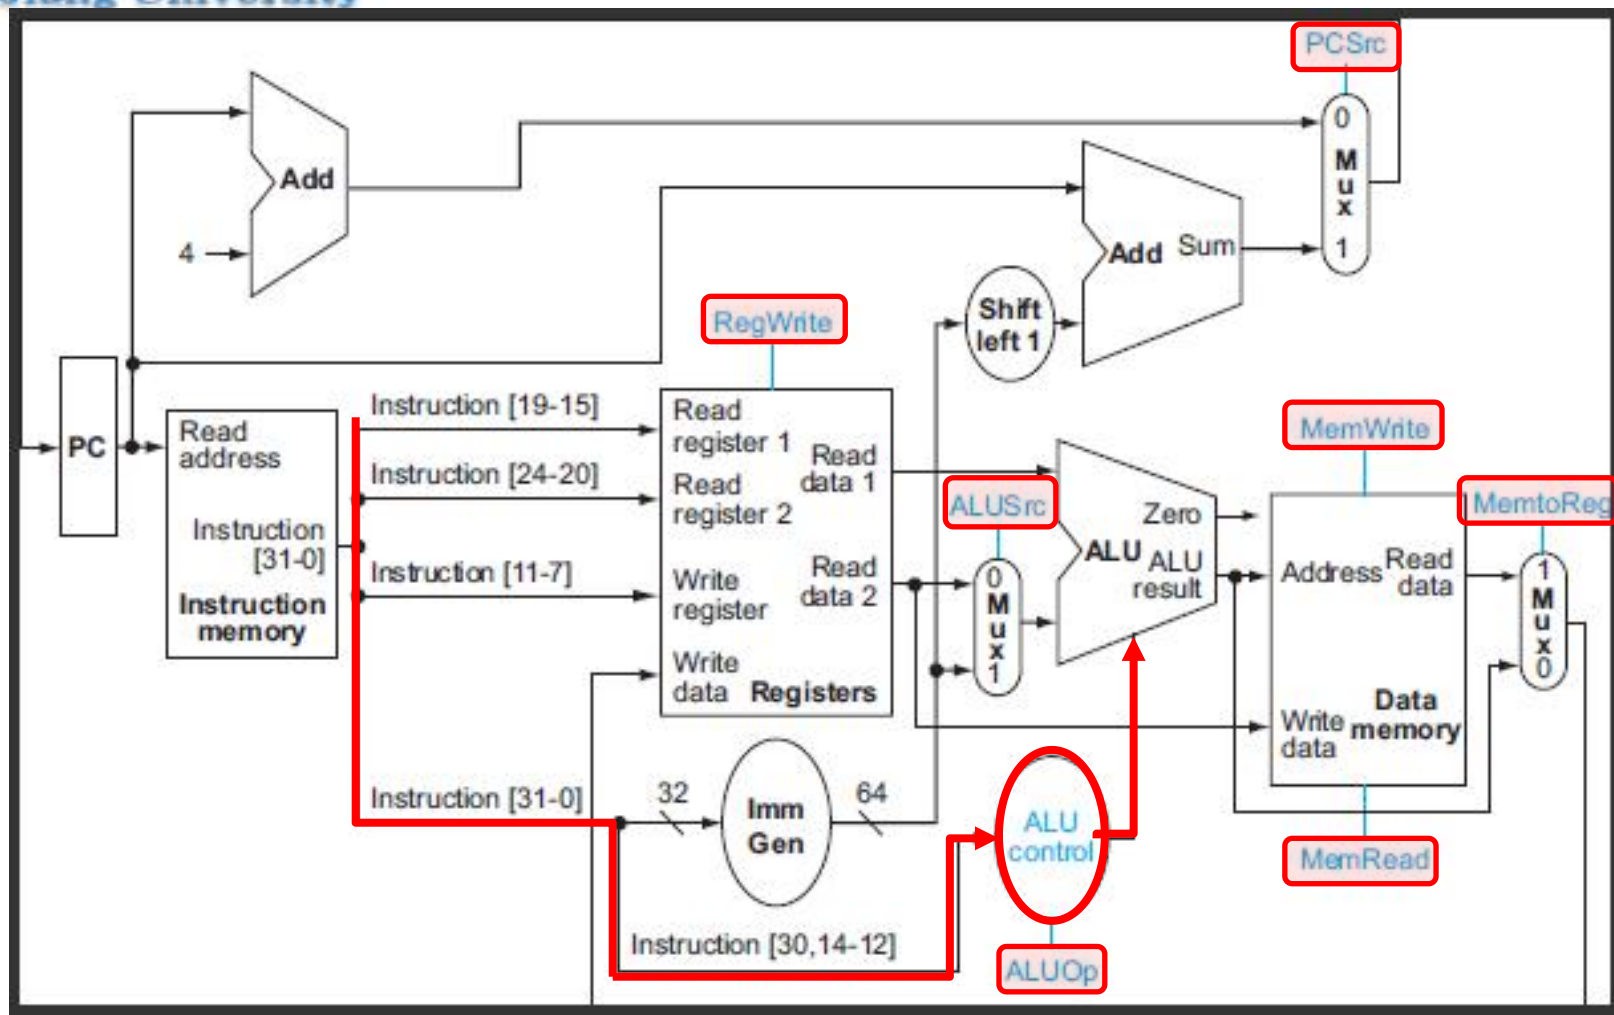
| opcode | ALUOp | Operation | Funct7 | funct3 | ALU function | ALU control |
|--------|-------|-----------------|----------|--------|--------------|-------------|
| ld | 00 | load register | XXXXXXXX | xxx | add | 0010 |
| sd | 00 | store register | XXXXXXXX | xxx | add | 0010 |
| beq | 01 | branch on equal | XXXXXXXX | xxx | subtract | 0110 |
| R-type | 10 | add | 0000000 | 000 | add | 0010 |
| | | subtract | 0100000 | 000 | subtract | 0110 |
| | | AND | 0000000 | 111 | AND | 0000 |
| | | OR | 0000000 | 110 | OR | 0001 |
| | | SLT | 0000000 | 010 | Slt | 0111 |





The ALU control is where and other signals(6)

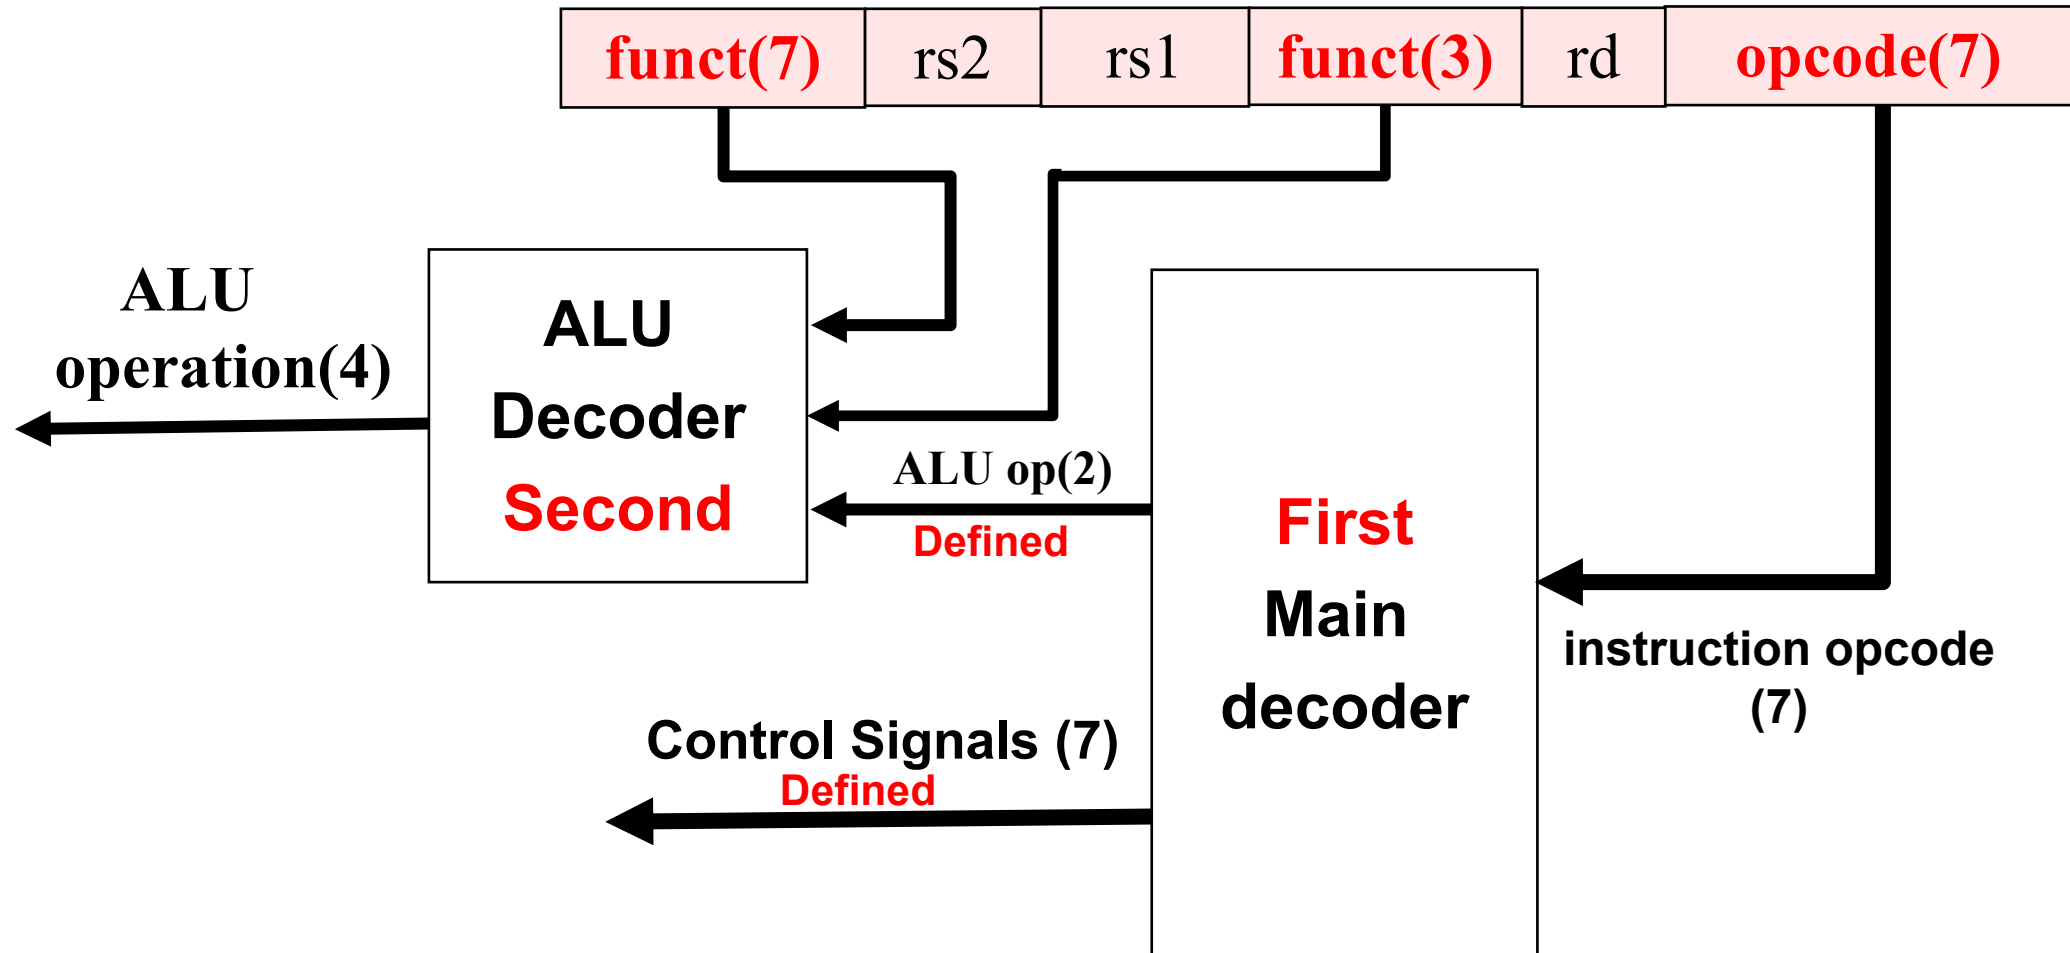
Output signals





Scheme of Controller

- 2-level decoder



Signals for datapath

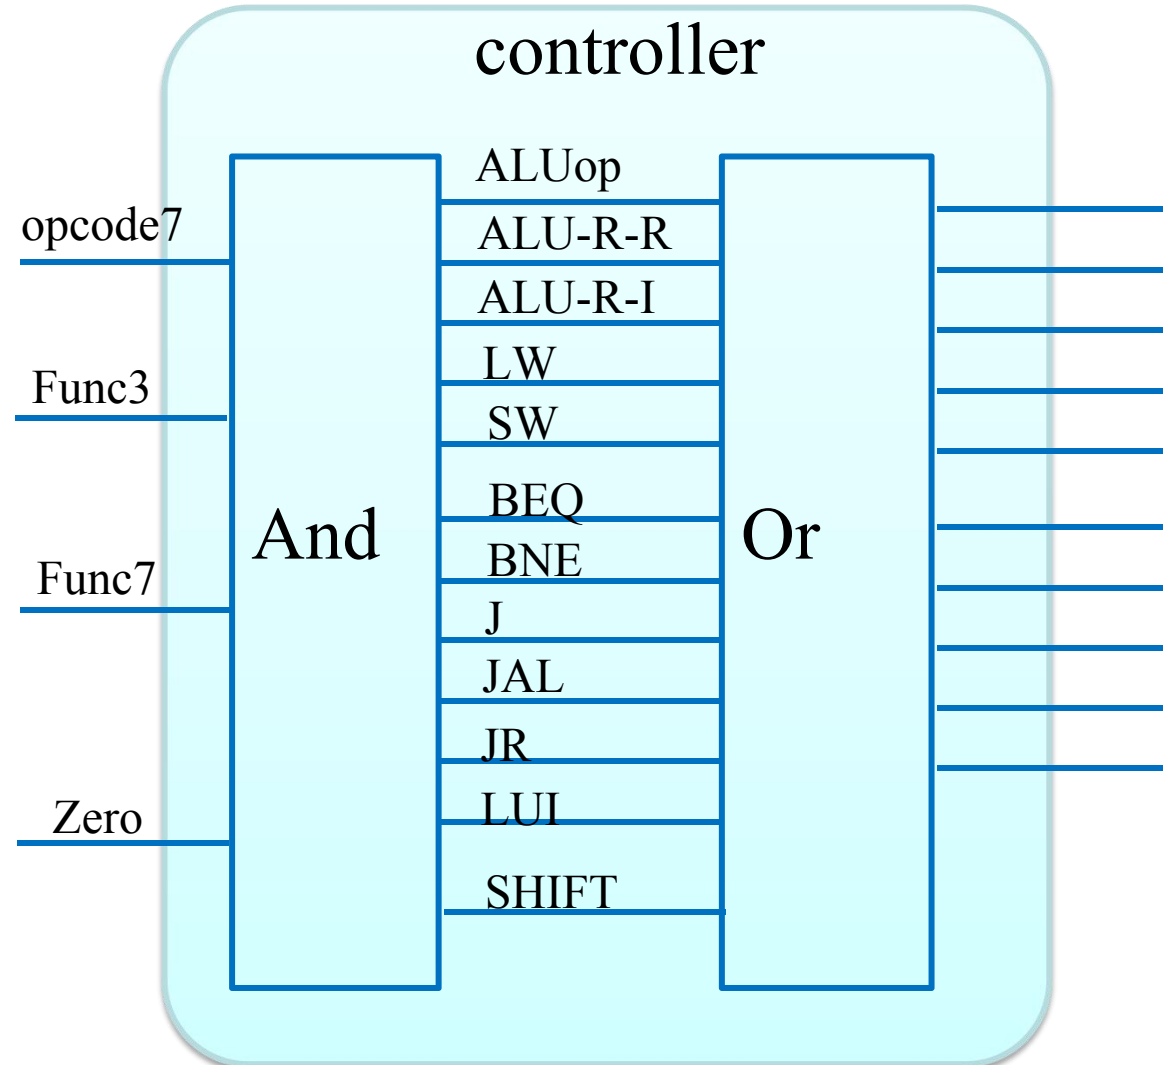
Defined 7

control signals

| Signal name | Effect when deasserted(=0) | Effect when asserted(=1) |
|-------------------|--|---|
| RegWrite | None | Register destination input is written with the value on the Write data input |
| ALUScr | The second ALU operand come from the second register file output (Read data 2) | The second ALU operand is the sign-extended lower 16 bits of the instruction.. |
| Branch (PCSrc) | The PC is replaced by the output of the adder that computers the value PC+4 | The PC is replaced by the output of the adder that computers the branch target. |
| Jump | The PC is replaced by PC+4 or branch target | The PC is updated by jump address computed by adder |
| MemRead | None | Data memory contents designated by the address input are put on the Read data output. |
| MemWrite | None | Data memory contents designated by the address input are replaced by value on the Write data input. |
| MemtoReg (2位) | 00: The value fed to register Write data input comes from the Alu | 01: The value fed to the register Write data input comes from the data memory. |
| | | 10: The value fed to the register Write data input comes from PC+4 |



Controller Implementation



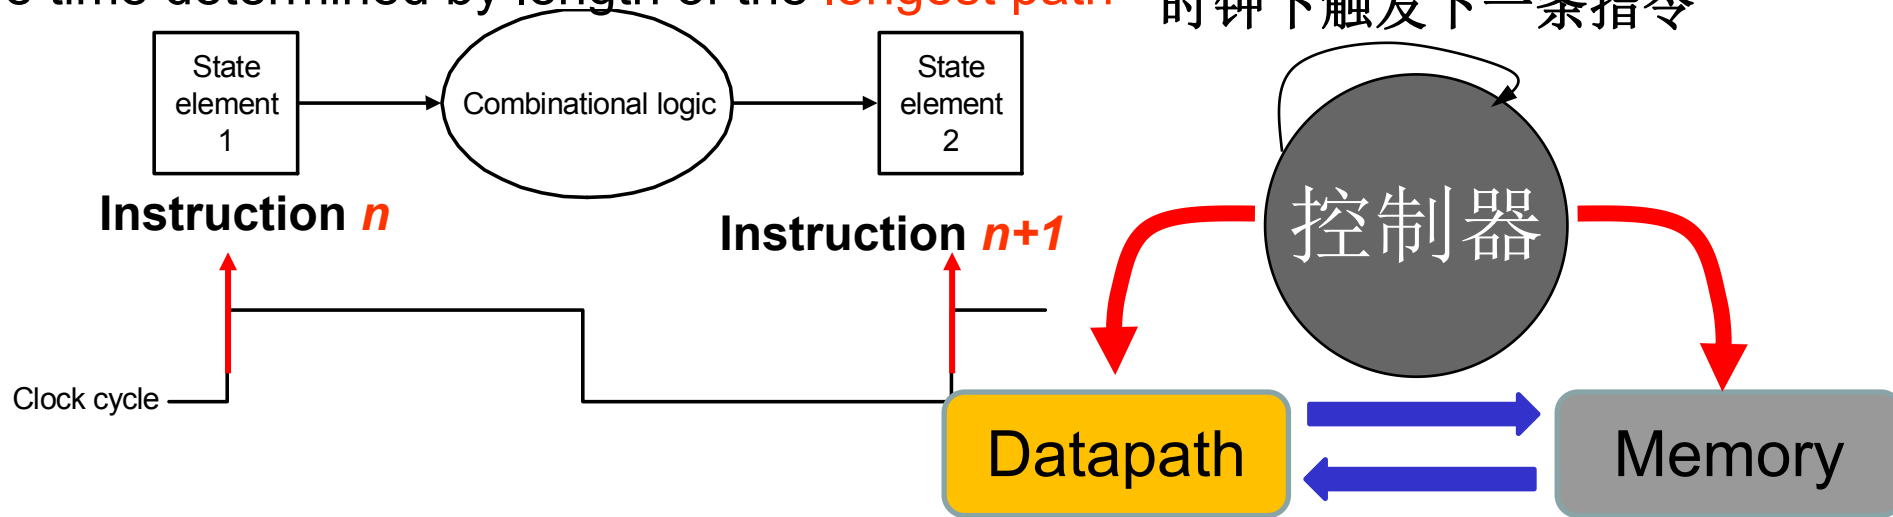
- $\text{SrcB} = \text{ALURI} + \text{LW} + \text{SW} + \text{SHIFT}$
- $\text{MRead} = \text{LW} + \text{LH} + \text{LB}$
- $\text{MWrite} = \text{SW} + \text{SH} + \text{SB}$
- $\text{M2Reg} = \text{LW} + \text{LH} + \text{LB}$
- $\text{RegDes} = \text{ALURI} + \text{LW} + \text{LUI} \quad (\text{rt})$
- $\text{Wreg} = \text{ALURR} + \text{ALURI} + \text{Load}$
- $+ \text{JAL} + \text{LUI} + \text{shift}$
- $= \sim(\text{Branch} + \text{J} + \text{JR} + \text{Store})$
- $\text{Taken} = \text{Beq} * \text{Zero} + \text{Bne} * \sim \text{Zero}$
- $\text{SrcA} = \text{SLL} + \text{SRL} = \text{shift}$
- ...





Our Simple Control Structure

- All of the logic is **combinational**
- We wait for everything to settle down, and the right thing to be done
 - ALU might not produce right answer? **right away**
 - we use write signals along with **clock** to determine when to write
- Cycle time determined by length of the **longest path** 时钟下触发下一条指令



We are ignoring some details like setup and hold times






Designing the Main Control Unit

First level

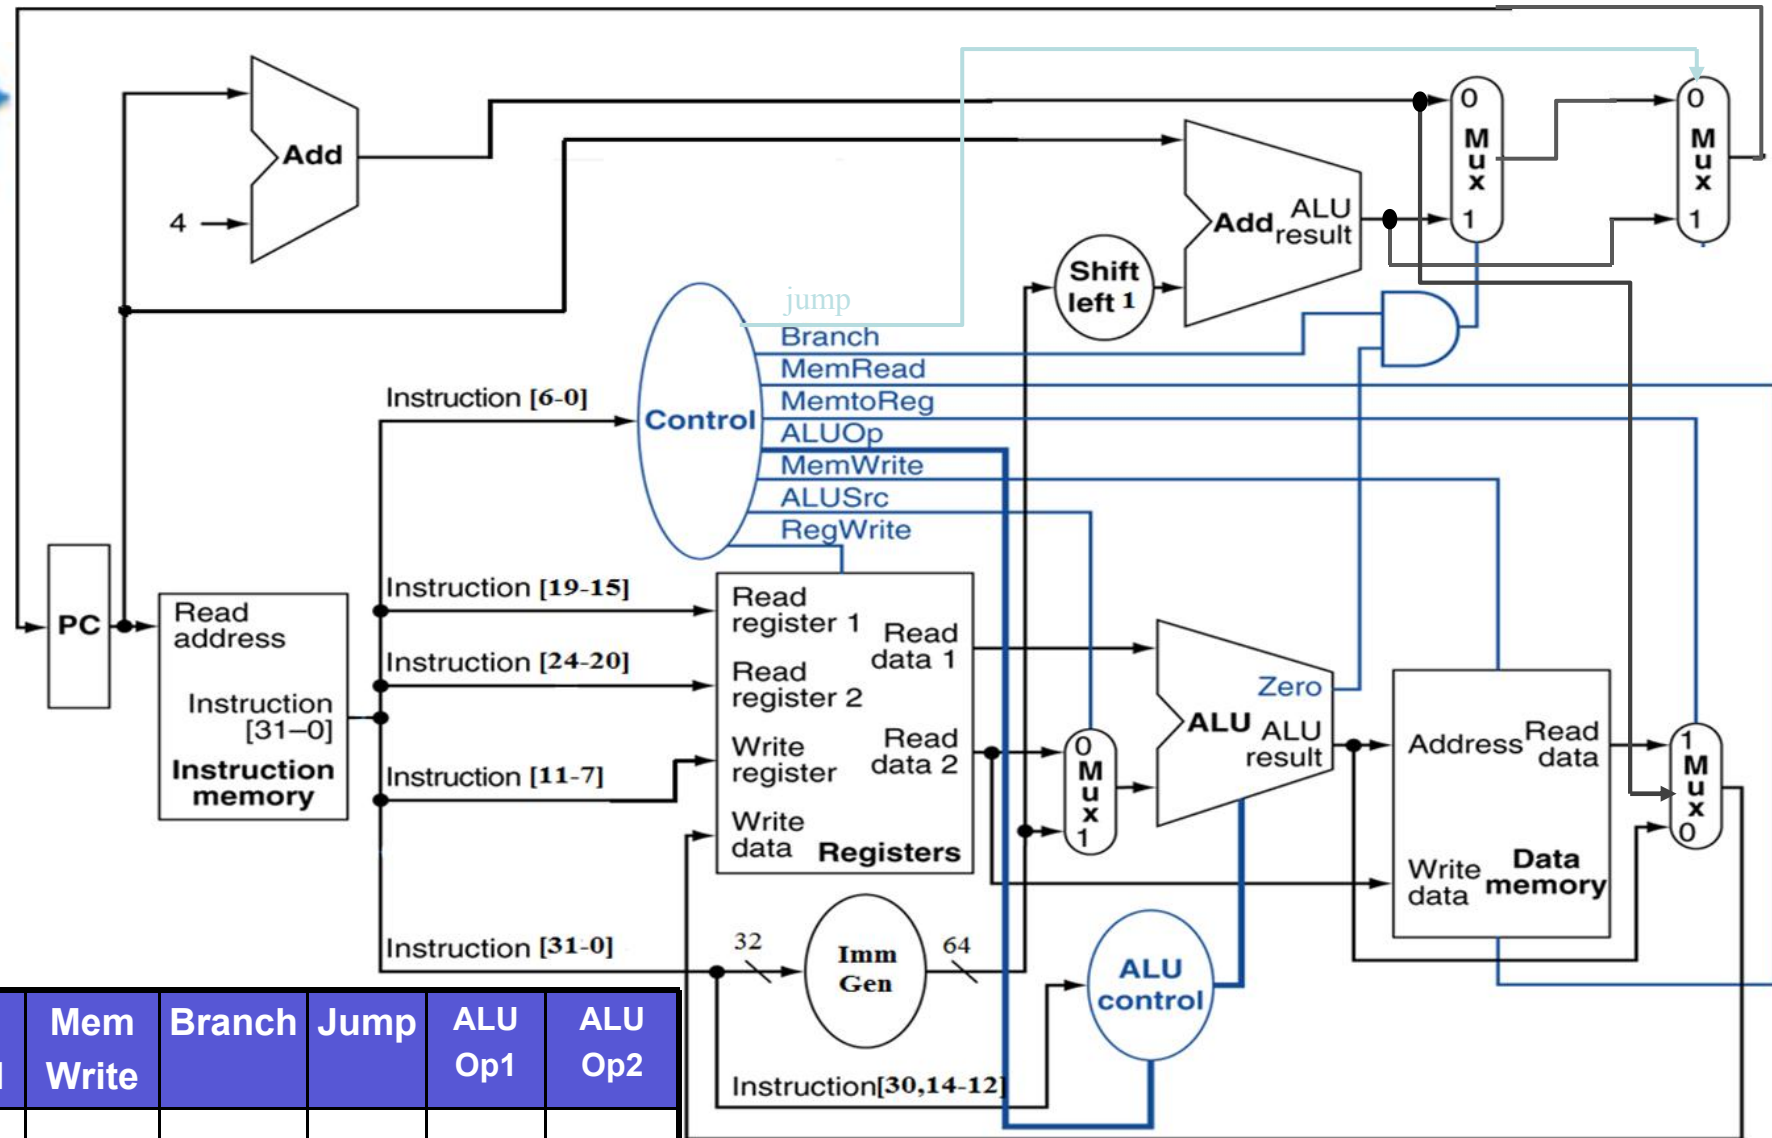
- Main Control Unit function
 - ALU op (2)
 - Divided 6 control signals into 2 groups
 - 4 Mux
 - 3 R/W





Truth Table for Main decoder

ZheJiang University



| Instruction | ALUSrc | MemtoReg | Reg Write | Mem Read | Mem Write | Branch | Jump | ALU Op1 | ALU Op2 |
|-------------|--------|----------|-----------|----------|-----------|--------|------|---------|---------|
| R-format | | | | | | | | | |
| Ld | | | | | | | | | |
| Sd | | | | | | | | | |
| beq | | | | | | | | | |
| Jal | | | | | | | | | |



Truth tables & Circuitry of main Controller

| 输入 | | 输出 | | | | | | | | |
|--------------|---------|---------|-----------|-----------|----------|-----------|--------|------|---------|---------|
| Instruction | OPCode | ALUSrcB | Memto-Reg | Reg Write | Mem Read | Mem Write | Branch | Jump | ALU Op1 | ALU Op0 |
| R-format | 0110011 | 0 | 00 | 1 | 0 | 0 | 0 | 0 | 1 | 0 |
| Ld(I-Type) | 0000011 | 1 | 01 | 1 | 1 | 0 | 0 | 0 | 0 | 0 |
| sd(S-Type) | 0100011 | 1 | X | 0 | 0 | 1 | 0 | 0 | 0 | 0 |
| beq(SB-Type) | 1100111 | 0 | X | 0 | 0 | 0 | 1 | 0 | 0 | 1 |
| Jal(UJ-Type) | 1101111 | X | 10 | 1 | 0 | 0 | 0 | 1 | X | X |



Main Controller Code

- 指令译码器参考描述

```
`define CPU_ctrl_signals {ALUSrc_B,MemtoReg,RegWR,MemWrite,Branch,Jump,ALUOp}  
  always @* begin  
    case(OPcode)  
      5'b01100: begin CPU_ctrl_signals = ?; end      //ALU  
      5'b00000: begin CPU_ctrl_signals = ?; end      //load  
      5'b01000: begin CPU_ctrl_signals = ?; end      //store  
      5'b11000: begin CPU_ctrl_signals = ?; end      //beq  
      5'b11011: begin CPU_ctrl_signals = ?; end      //jump  
      5'b00100: begin CPU_ctrl_signals = ?; end      //ALU(addi;;;)  
  
      .....  
    default: begin CPU_ctrl_signals = ?; end  
  endcase  
end
```





Design the ALU Decoder

second level

- ALU operation is decided by 2-bit ALUOp derived from opcode, and funct7 & funct3 fields of the instruction
 - Combinational logic derives ALU control

| opcode | ALUOp | Operation | Funct7 | funct3 | ALU function | ALU control |
|--------|-------|-----------------|----------|--------|--------------|-------------|
| ld | 00 | load register | XXXXXXXX | xxx | add | 0010 |
| sd | 00 | store register | XXXXXXXX | xxx | add | 0010 |
| beq | 01 | branch on equal | XXXXXXXX | xxx | subtract | 0110 |
| R-type | 10 | add | 0000000 | 000 | add | 0010 |
| | | subtract | 0100000 | 000 | subtract | 0110 |
| | | AND | 0000000 | 111 | AND | 0000 |
| | | OR | 0000000 | 110 | OR | 0001 |
| | | SLT | 0000000 | 010 | Slt | 0111 |





ALU Controller Code

- ALU Control HDL Description

```
assign Fun = {Fun3, Fun7};
```

```
always @* begin
```

```
  case(ALUOp)
```

```
    2'b10: ALU_Control = ? ;
```

//add计算地址

```
    2'b11: ALU_Control = ? ;
```

//sub比较条件

```
    2'b00:
```

```
      case(Fun)
```

```
        4'b0000: ALU_Control = 3'b010 ;    //add
```

```
        4'b0001: ALU_Control = ? ;          //sub
```

```
        4'b1110: ALU_Control = ? ;          //and
```

```
        4'b1100: ALU_Control = ? ;          //or
```

```
        4'b0100: ALU_Control = ? ;          //slt
```

```
        4'b1010: ALU_Control = ? ;          //srl
```

```
        4'b1000: ALU_Control = ? ;          //xor
```

```
      default: ALU_Control=3'bx;
```

```
    endcase
```

```
  2'b01:
```

```
    case(Fun3)
```

```
  .....
endcase
```

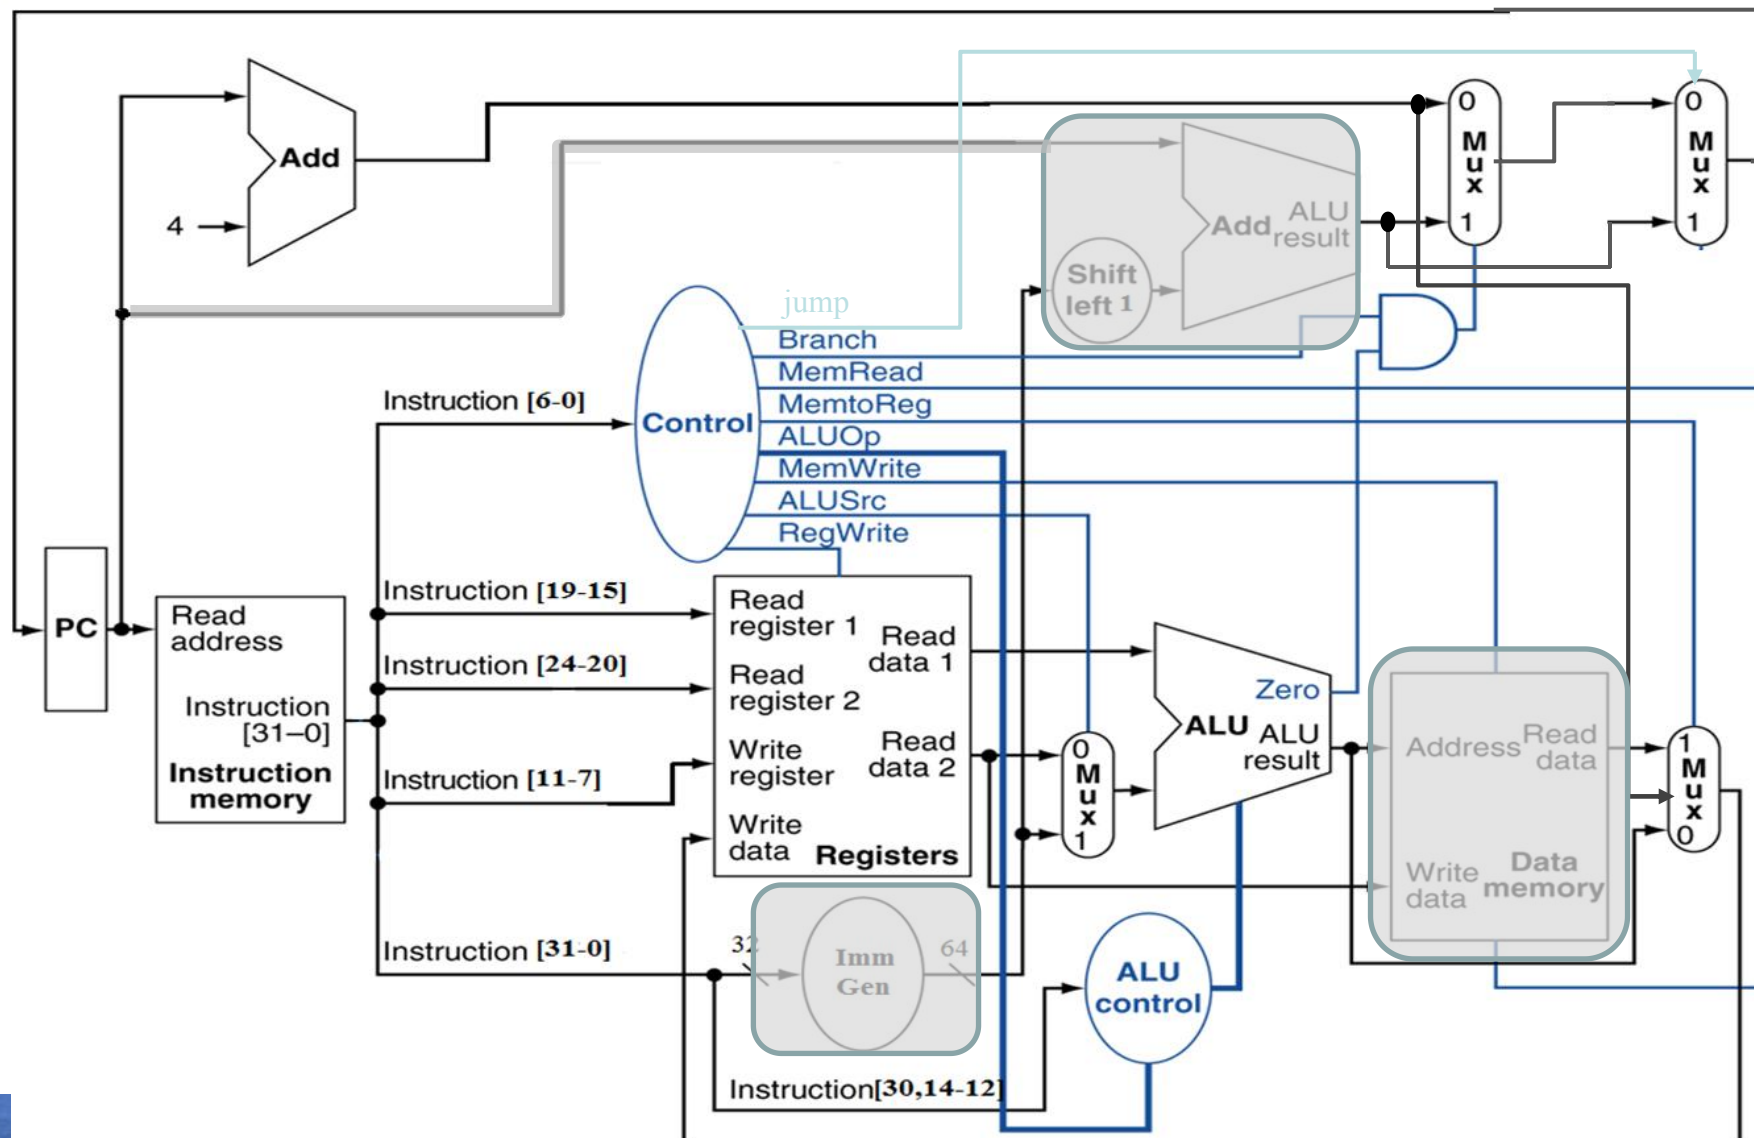




| funct7 | rs2 | rs1 | funct3 | rd | opcode |
|--------|--------|--------|--------|--------|--------|
| 7 bits | 5 bits | 5 bits | 3 bits | 5 bits | 7 bits |

add x9, x20, x21

1. Read two register operands
2. Perform arithmetic/logical operation
3. Write register result

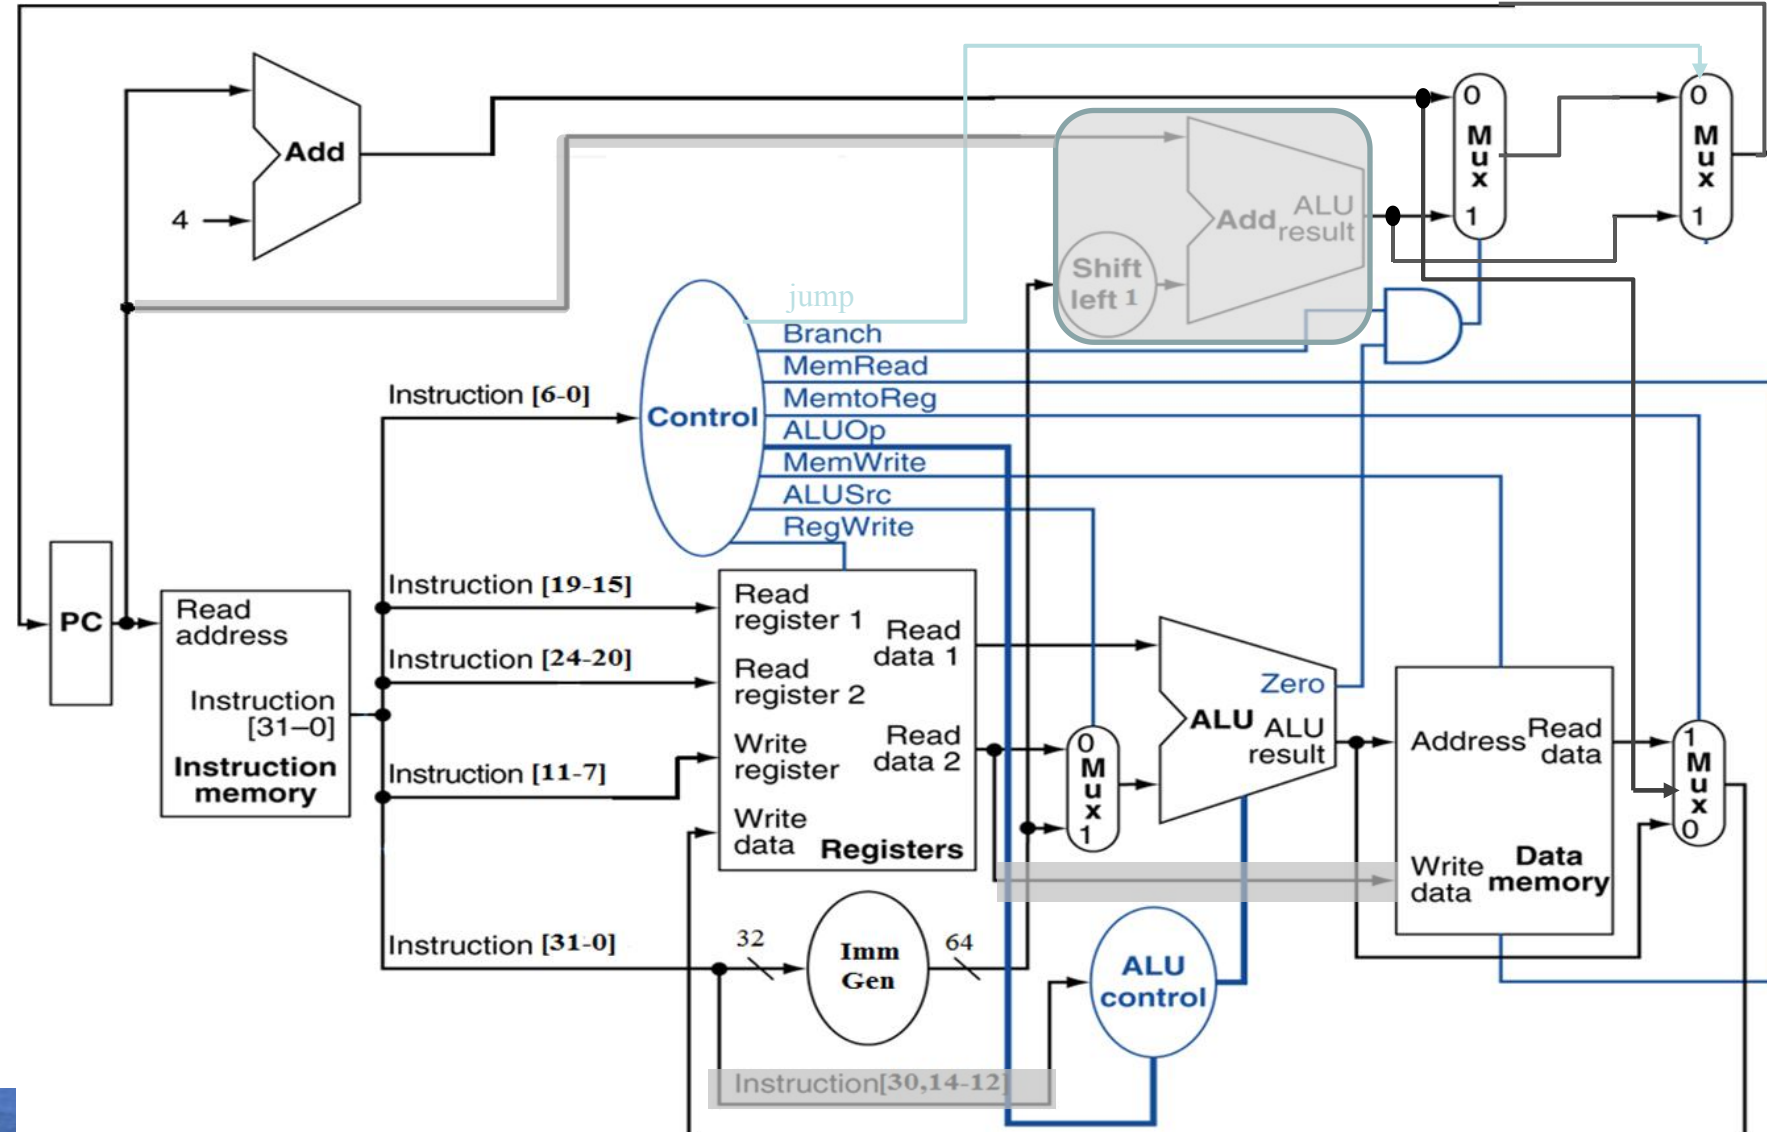




ld x1, 200(x2)

1. Read register operands
2. Calculate address using 12-bit offset
 1. Use ALU, but sign-extend offset
3. Read memory and update register

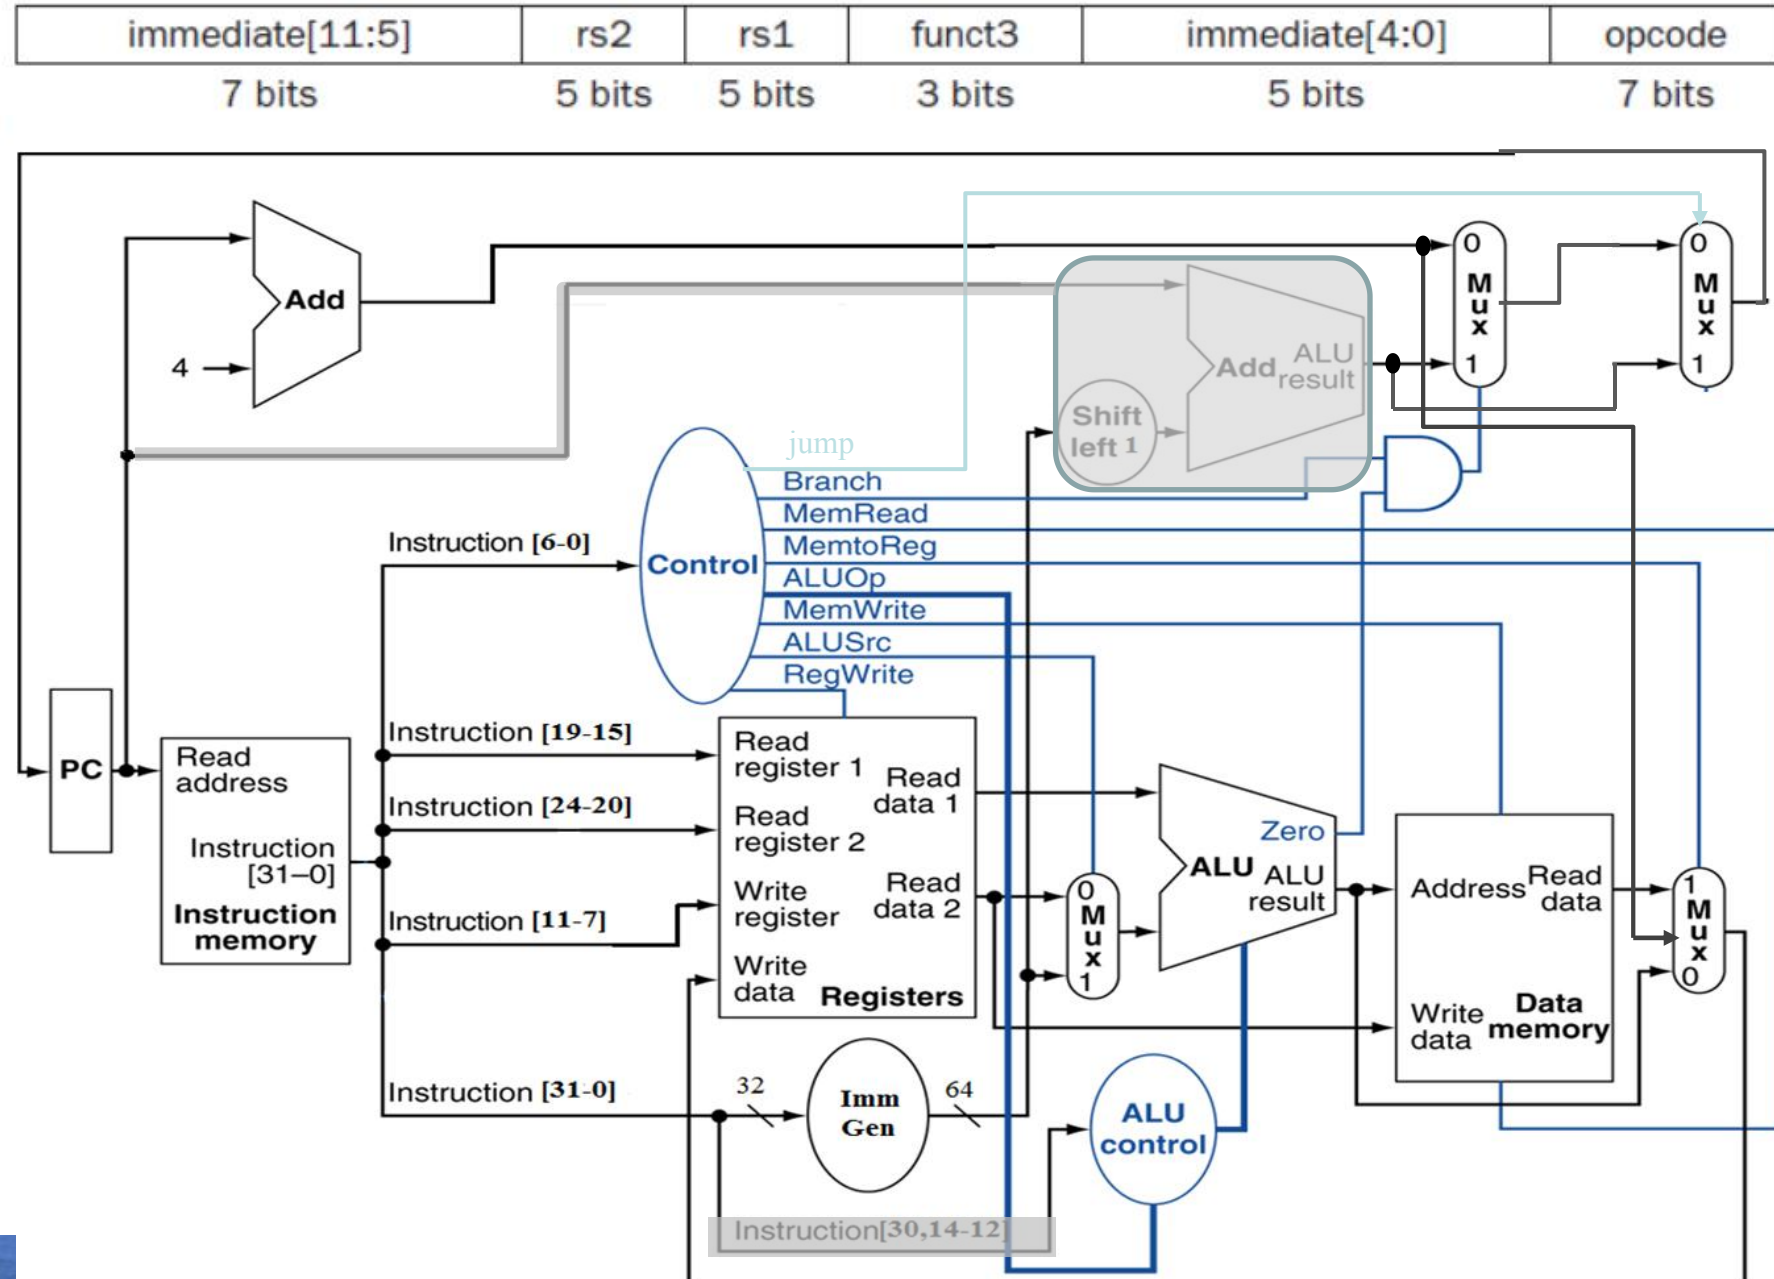
| immediate | rs1 | funct3 | rd | opcode |
|-----------|--------|--------|--------|--------|
| 12 bits | 5 bits | 3 bits | 5 bits | 7 bits |





sd x1, 200(x2)

1. Read register operands
2. Calculate address using 12-bit offset
 1. Use ALU, but sign-extend offset
3. Write register value to memory





BEQ Instruction

| imm[12] | imm[10:5] | rs2 | rs1 | funct3 | imm[4:1] | imm[11] | opcode |
|---------|-----------|-----|-----|--------|----------|---------|--------|
|---------|-----------|-----|-----|--------|----------|---------|--------|

beq x1, x2, 200

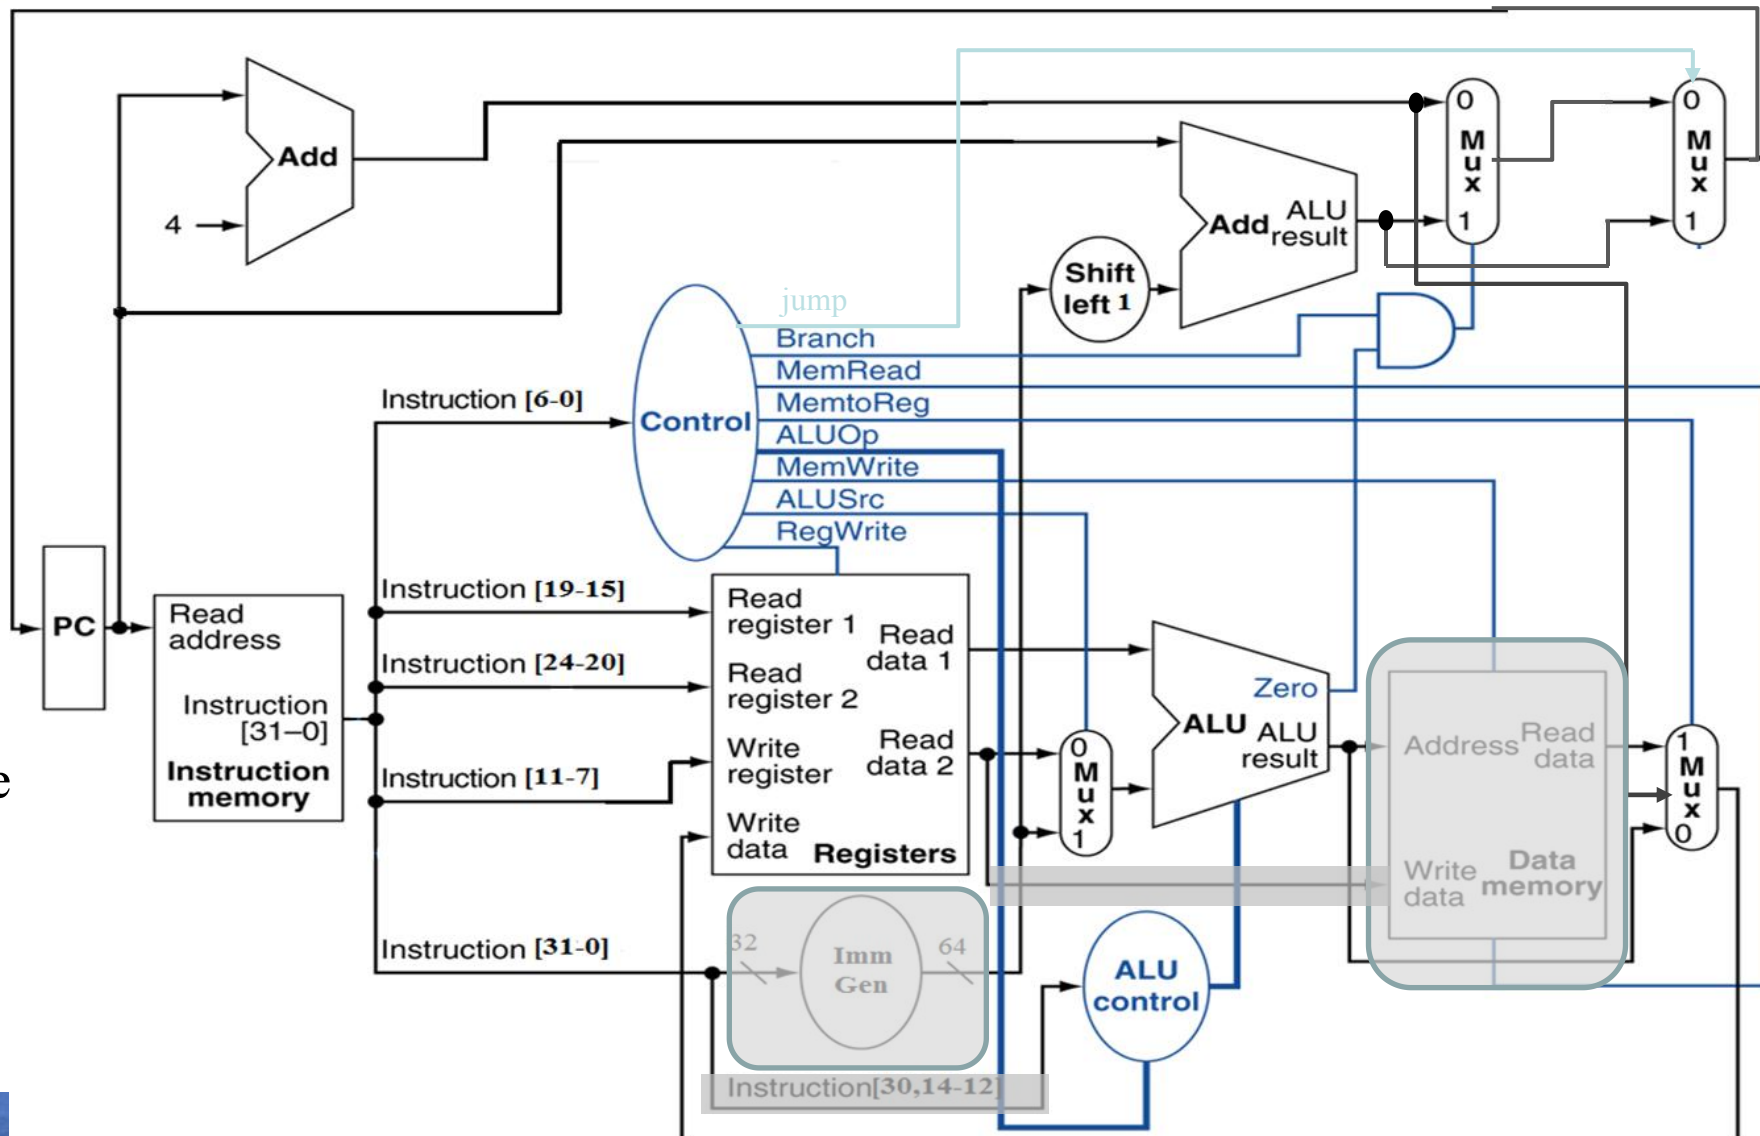
Read register operands

Compare operands

- Use ALU, subtract and check Zero output

Calculate target address

- Sign-extend displacement
- Shift left 1 place (halfword displacement)
- Add to PC value and update PC

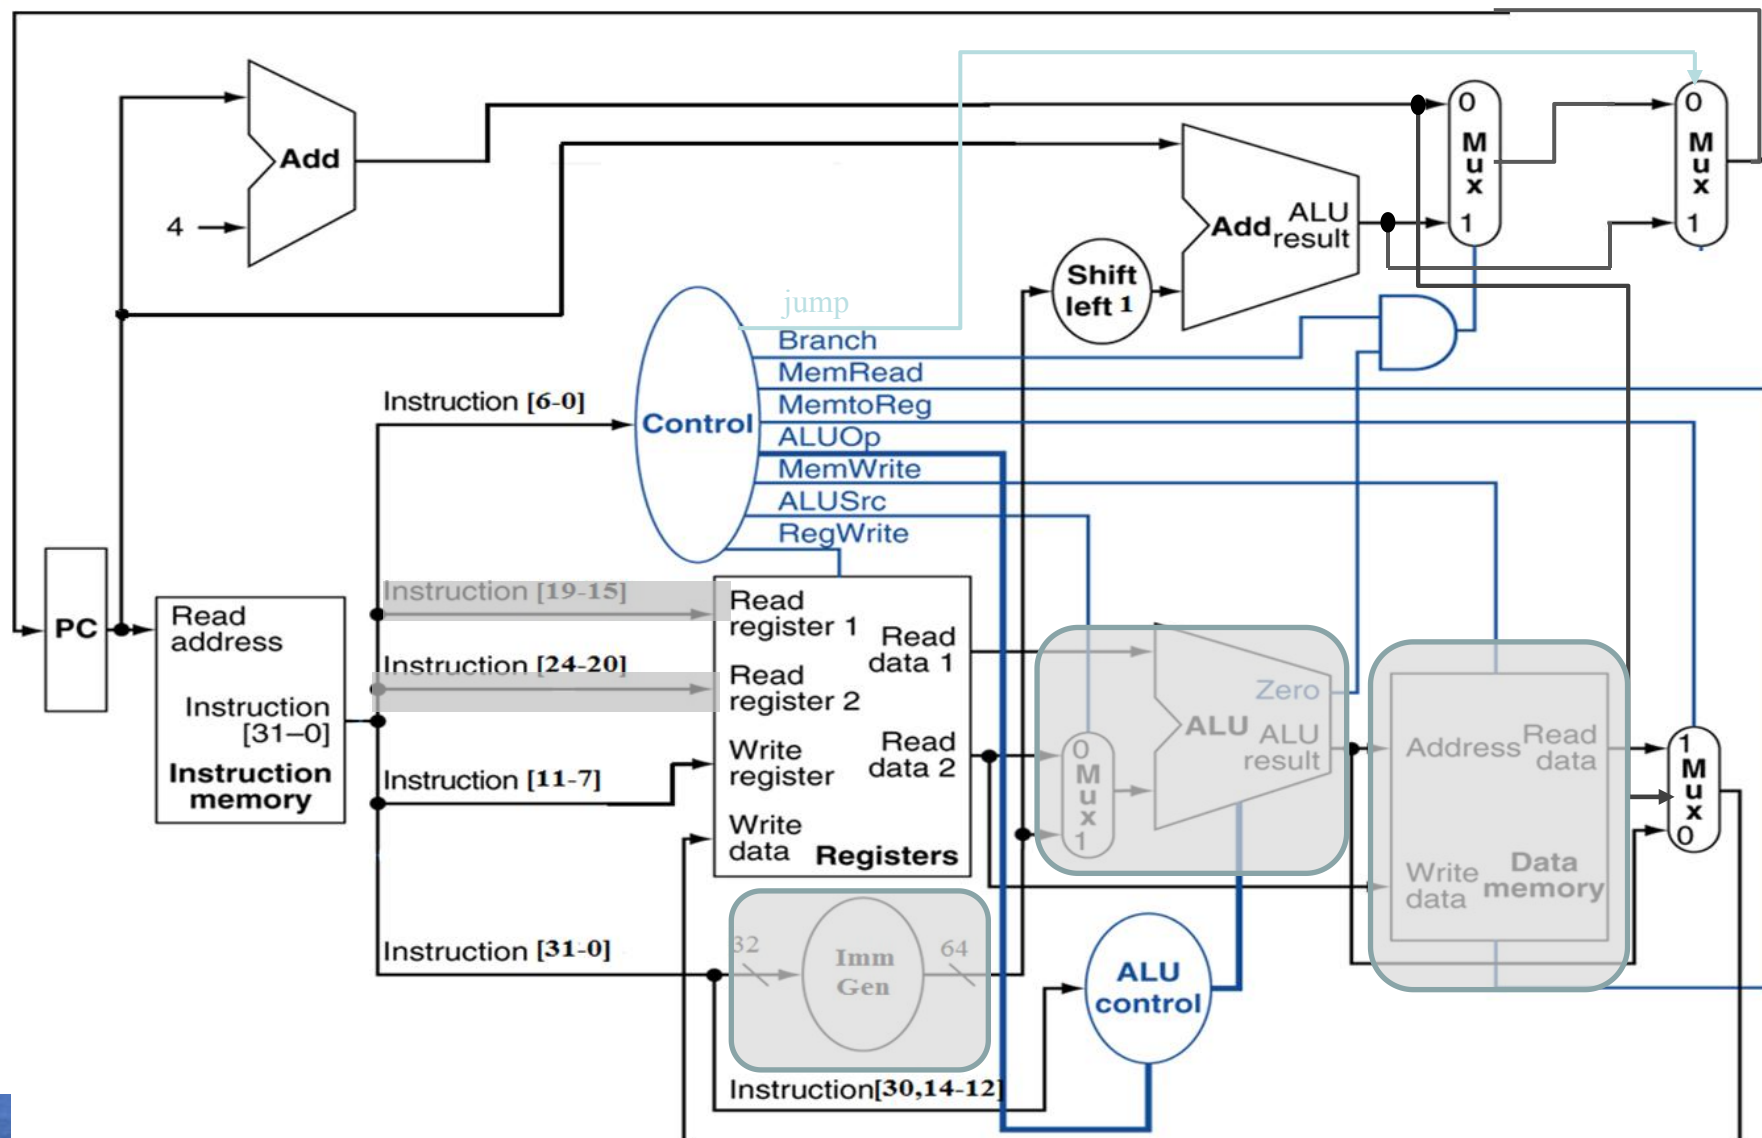




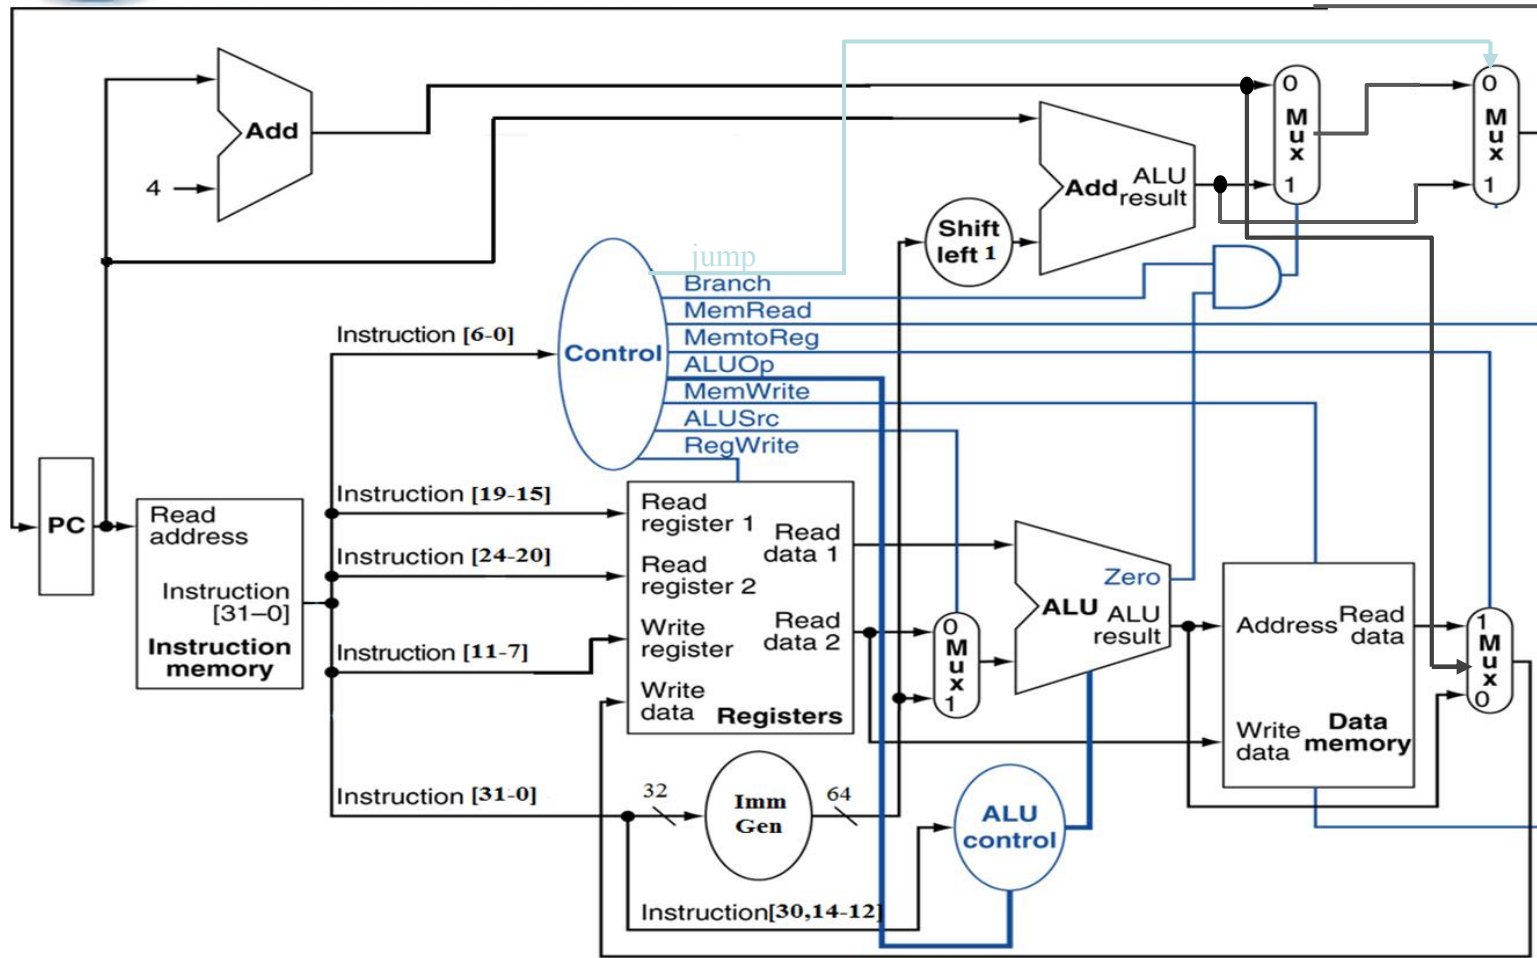
jal x1, procedure

- Write PC+4 to rd
- Calculate target address
 - Sign-extend displacement
 - Shift left 1 place (halfword displacement)
 - Add to PC value and update PC

| | | | | | |
|---------|-----------|---------|------------|----|--------|
| imm[20] | imm[10:1] | imm[11] | imm[19:12] | rd | opcode |
|---------|-----------|---------|------------|----|--------|



Single Cycle Implementation performance for lw



200ps

100+100=200ps

200ps

200ps

- Calculate cycle time assuming negligible delays except:
 - memory (200ps), ALU and adders (200ps), register file access (100ps)



Performance in Single Cycle Implementation

- Let's see the following table:

| Instr | Instr fetch | Register read | ALU op | Memory access | Register write | Total time |
|----------|-------------|---------------|--------|---------------|----------------|------------|
| ld | 200ps | 100 ps | 200ps | 200ps | 100 ps | 800ps |
| sd | 200ps | 100 ps | 200ps | 200ps | | 700ps |
| R-format | 200ps | 100 ps | 200ps | | 100 ps | 600ps |
| beq | 200ps | 100 ps | 200ps | | | 500ps |

❖ The conclusion:

Different instructions needs different time.

The clock cycle must meet the need of the slowest instruction.

So, some time will be wasted.





Performance Issues

- Longest delay determines clock period
 - Critical path: load instruction
 - Instruction memory → register file → ALU → data memory → register file
- Wasteful of area. If the instruction needs to use some functional unit multiple times.
 - E.g., the instruction 'mult' needs to use the ALU repeatedly. So, the CPU will be very large.
- Violates design principle
 - Making the common case fast
- We will improve performance by pipelining





4.6 Interruption & Exception

- What's Interruption & Exception ?
- Why need interruption & exception?
- How to deal with interruption & Exception in RISC V ?
 - Transfer control to exception handler & return from exception
 - Control status registers
 - CSR instructions
 - How to write an exception handler ?





Interrupt & Exception

- The cause of changing CPU's work flow :
 - Control instructions in program (bne/beq, jal , etc)
It is **foreseeable** in programming flow
 - Something happen suddenly (Exception and Interruption)
It is **unpredictable**
 - Call Instructions triggered by hardware
- Exception
 - Arises within the CPU when execute instruction
 - e.g., overflow, undefined opcode, syscall, ...
- Interrupt
 - From an external I/O controller
- Dealing with them without sacrificing performance is hard

| Type of event | From where? | RISC-V terminology |
|---|-------------|--------------------|
| System reset | External | Exception |
| I/O device request | External | Interrupt |
| Invoke the operating system from user program | Internal | Exception |
| Using an undefined instruction | Internal | Exception |
| Hardware malfunctions | Either | Either |





Why we need interrupt ?

- When you double click the mouse
- When a network package arrives
- When you want to print a sentence on screen

Event external to the running program can interrupt the processor : ex. Interruption driven I/O

- polling ----waste a lot of processor time
- **Interruption driven I/O**
- DMA ---- direct memory access



Why we need exception ?

- Processor can be interrupted by exceptional events that occur while the program is running that are caused by the program itself.
- Example:
 - Page fault:
 - need OS to load the page into the memory from disk, then resume the program
 - Memory address fault (segmentation fault)
 - Undefined opcode
 - The OS will stop the program and then transfer to other process.





Handling Exceptions

- **Save PC** of offending (or interrupted) instruction
 - In RISC-V: **Supervisor Exception Program Counter (SEPC)** (P316, 7th line)
- **Save indication of the problem**
 - In RISC-V: **Supervisor Exception Cause Register (SCAUSE)**
 - 64 bits, but most bits unused
 - Exception code field: 2 for undefined opcode, 12 for hardware malfunction, ...
- **Jump to handler**
 - Assume at 0000 0000 1C09 0000_{hex}
 - Entry address in a special register : **Supervisor Trap Vector (STVEC)**, which can be loaded by OS.





An Alternate Mechanism

- Vectored Interrupts

- Handler address determined by the cause

- Exception vector address to be added to a vector table base register:

- Undefined opcode 00 0100 0000_{two}

- Hardware malfunction: 01 1000 0000_{two}

- ...: ...

- Instructions either

- Deal with the interrupt, or

- Jump to real handler





Handler Actions

- Read cause, and transfer to relevant handler
- Determine action required
- If restartable
 - Take corrective action
 - use SEPC to return to program (mret)
- Otherwise
 - Terminate program
 - Report error using SEPC, SCAUSE, ...
 - OS make the choice to transfer to another ready process





Exceptions in RISC V

- Transfer control to exception handler & return from exception
- Control status registers
- CSR instructions
- How to write an exception handler ?



Privileged Architecture

- Software stack

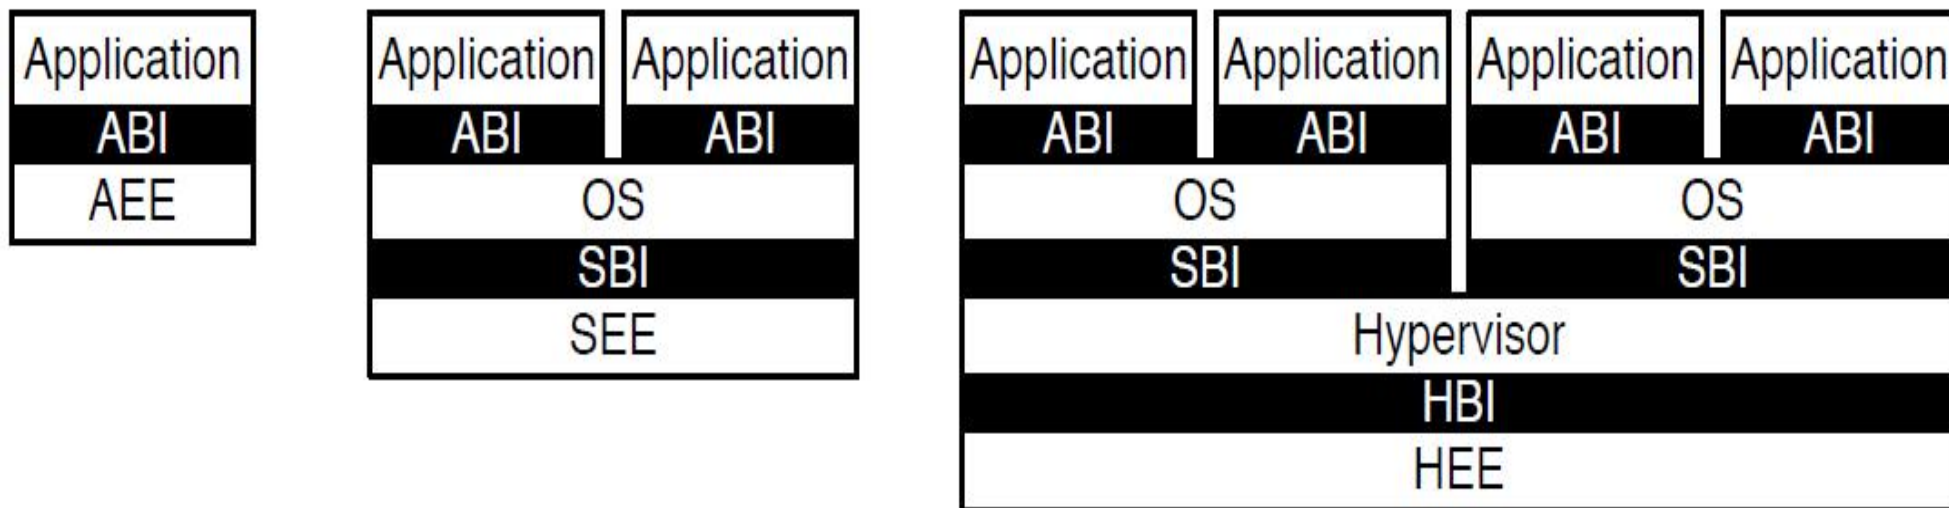


Figure 1.1: Different implementation stacks supporting various forms of privileged execution.





Privileged level

- Privilege Level in RISC V
 - Provide **protection** between different components of the software stack

| Level | Encoding | Name | Abbreviation |
|-------|----------|------------------|--------------|
| 0 | 00 | User/Application | U |
| 1 | 01 | Supervisor | S |
| 2 | 10 | <i>Reserved</i> | |
| 3 | 11 | Machine | M |

Table 1.1: RISC-V privilege levels.

| Number of levels | Supported Modes | Intended Usage |
|------------------|-----------------|---|
| 1 | M | Simple embedded systems |
| 2 | M, U | Secure embedded systems |
| 3 | M, S, U | Systems running Unix-like operating systems |

Table 1.2: Supported combinations of privilege modes.



CSR Address Mapping Convention

- CSR[11..0] 4096
- 12 bit encoding space



The lowest privileged level that can access CSR

11: read only

00/01/10: read/write

| Machine CSRs | | | | |
|--------------|----|------|-------------|--------------------------------|
| 00 | 11 | XXXX | 0x300-0x3FF | Standard read/write |
| 01 | 11 | 0XXX | 0x700-0x77F | Standard read/write |
| 01 | 11 | 100X | 0x780-0x79F | Standard read/write |
| 01 | 11 | 1010 | 0x7A0-0x7AF | Standard read/write debug CSRs |
| 01 | 11 | 1011 | 0x7B0-0x7BF | Debug-mode-only CSRs |
| 01 | 11 | 11XX | 0x7C0-0x7FF | Custom read/write |
| 10 | 11 | 0XXX | 0xB00-0xB7F | Standard read/write |
| 10 | 11 | 10XX | 0xB80-0xBBF | Standard read/write |
| 10 | 11 | 11XX | 0xBC0-0xBFF | Custom read/write |
| 11 | 11 | 0XXX | 0xF00-0xF7F | Standard read-only |
| 11 | 11 | 10XX | 0xF80-0xFBF | Standard read-only |
| 11 | 11 | 11XX | 0xFC0-0xFFF | Custom read-only |



8 important CSR for exception handling

■ 8 CSR :

- mtvec (Machine Trap Vector): jump to this address when exception
- mepc (Machine Exception PC): the instruction raise the exception
- mcause (Machine exception Cause): which kind of exception(cause)
- mie (Machine Interrupt Enable): which exception can be handled or neglected
- mip (Machine interrupt pending): pending interruptions (read only register)
- mtval (Machine trap value): error address , illegal instruction, or 0
- mscratch (Machine Scratch):
- mstatus (Machine status) : processor status





| base[31..2] | mode |
|-------------|------|
|-------------|------|

- Store the interruption handler entrance address
- The base can be explained according to mode code
 - 00: $PC \leftarrow \text{base}$
 - x1: vector mode, $PC \leftarrow \text{mtval} - 1 + 4x$ (not required)





- Save the instruction address when exception raised or interruption happens.
 - the PC indicate the instruction that raise the exception
 - the instruction need to be executed after back from interruption
 - next instruction



Exception Causes

Asynchronization exception

- Software interrupt
- Timer interrupt
- External interrupt

Synchronization exception

- Address misaligned
- Access fault
- illegal instruction
- Breakpoint
- Environment call

| Interrupt / Exception cause[XLEN-1] | Exception Code mcause[XLEN-2:0] | Description |
|--|------------------------------------|--------------------------------|
| 1 | 1 | Supervisor software interrupt |
| 1 | 3 | Machine software interrupt |
| 1 | 5 | Supervisor timer interrupt |
| 1 | 7 | Machine timer interrupt |
| 1 | 9 | Supervisor external interrupt |
| 1 | 11 | Machine external interrupt |
| 0 | 0 | Instruction address misaligned |
| 0 | 1 | Instruction access fault |
| 0 | 2 | Illegal instruction |
| 0 | 3 | Breakpoint |
| 0 | 4 | Load address misaligned |
| 0 | 5 | Load access fault |
| 0 | 6 | Store address misaligned |
| 0 | 7 | Store access fault |
| 0 | 8 | Environment call from U-mode |
| 0 | 9 | Environment call from S-mode |
| 0 | 11 | Environment call from M-mode |
| 0 | 12 | Instruction page fault |
| 0 | 13 | Load page fault |
| 0 | 15 | Store page fault |

图 10.3: RISC-V 异常和中断的原因。中断时 mcause 的最高有效位置 1，同步异常时置 0，且低有效位标识了中断或异常的具体原因。只有在实现了监管者模式时才能处理监管者模式中断和页面错误异常



Machine Status Register

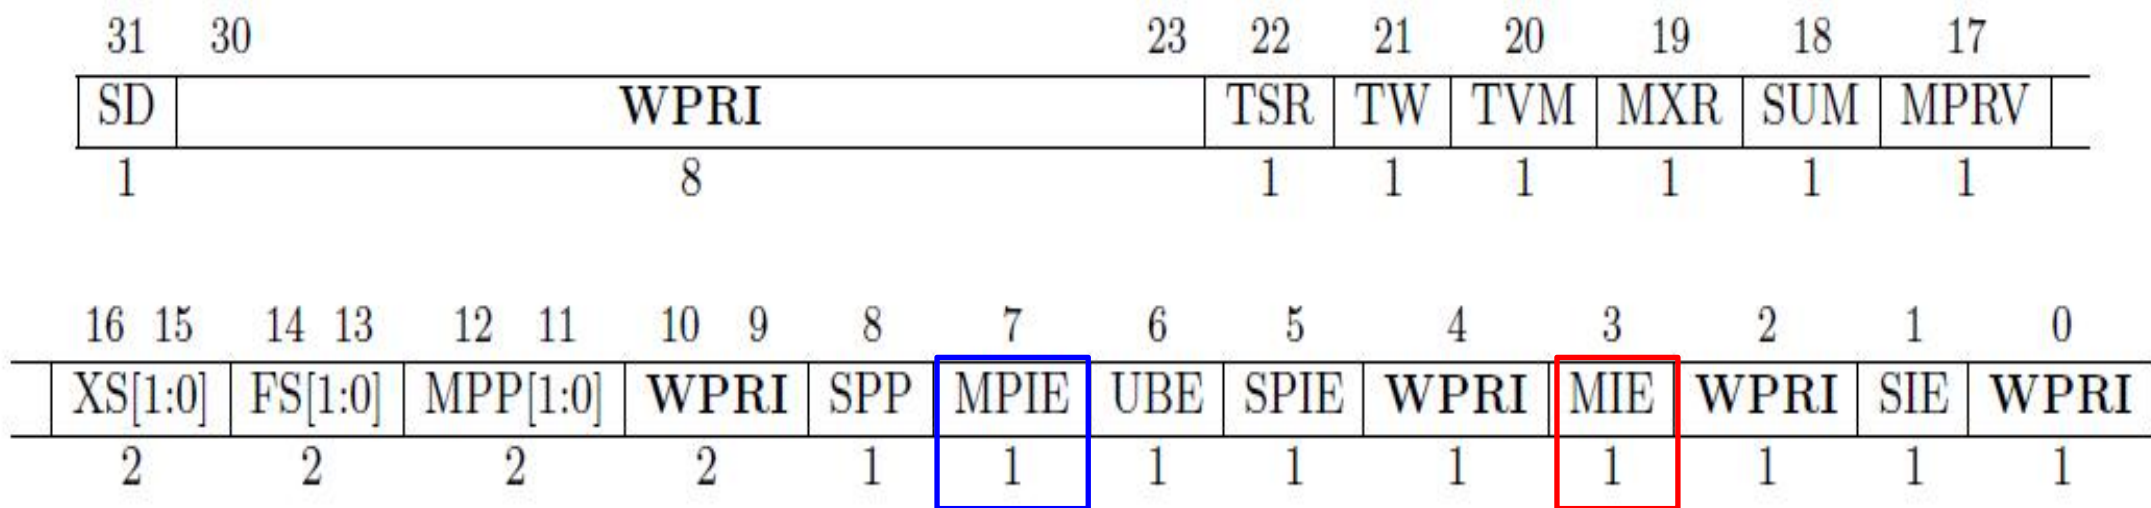


Figure 3.7: Machine-mode status register (`mstatus`) for RV32.





CSR instructions

- csrrw rd, csr, zimm[4..0]: $t \leftarrow \text{CSRs}[\text{csr}], \text{CSRs}[\text{csr}] \leftarrow x[\text{rs1}], x[\text{rd}] \leftarrow t$
- csrrs rd, csr, rs1: $t \leftarrow \text{CSRs}[\text{csr}], \text{CSRs}[\text{csr}] \leftarrow t \mid x[\text{rs1}], x[\text{rd}] \leftarrow t$
- csrrc rd, csr, rs1: $t \leftarrow \text{CSRs}[\text{csr}], \text{CSRs}[\text{csr}] \leftarrow t \ \& \ \sim x[\text{rs1}], x[\text{rd}] \leftarrow t$
- csrrwi rd, csr, zimm[4..0]: $x[\text{rd}] \leftarrow \text{CSRs}[\text{csr}], \text{CSRs}[\text{csr}] \leftarrow \text{zimm}$
- cssrrsi rd, csr, zimm[4..0]: $t \leftarrow \text{CSRs}[\text{csr}], \text{CSRs}[\text{csr}] \leftarrow t \mid \text{zimm}; x[\text{rd}] \leftarrow t$
- csrrci rd, csr, zimm[4..0]: $t \leftarrow \text{CSRs}[\text{csr}], \text{CSRs}[\text{csr}] \leftarrow t \ \& \ \sim \text{zimm}; x[\text{rd}] \leftarrow t$
- Mret:





How Control Checks for Exceptions (Hardware)

- Add test logic
 - illegal instruction, load address misaligned, store address misaligned
- add control signal
 - CauseWrite for mcause
 - EPCWrite for mepc
 - TVALWrite for mtval
- process of control
 - $mepc \leftarrow PC(\text{exception / interruption})$
 - $mcause \leftarrow \text{set correspondent bit}$
 - $mtval \leftarrow \text{memory address or illegal instruction}$
 - $mstatus.mpie \leftarrow Mstatus.mie; mstatus.mie \leftarrow 0; mstatus.mpp \leftarrow mp; mp \leftarrow 11$
 - $PC \leftarrow \text{address of process routine (mtvec, ex. 1c090000)}$





When jump to exception handler ?

- Jump to handler
 - Assume at 0000 0000 1C09 0000_{hex}
- jump when
 $mstatus.MIE = 1 \ \&\& \ mie[i] = 1 \ \&\& \ mip[i] = 1$





Ex. Exception handler

save registers

csrrw a0, mscratch, a0 # save a0; set a0 = &temp storage

sw a1, 0(a0) # save a1

sw a2, 4(a0) # save a2

sw a3, 8(a0) # save a3

sw a4, 12(a0) # save a4

decode interrupt cause

csrr a1, mcause # read exception cause

bgez a1, exception # branch if not an interrupt

andi a1, a1, 0x3f # isolate interrupt cause

li a2, 7 # a2 = timer interrupt cause

bne a1, a2, otherInt # branch if not a timer interrupt





Ex. Exception handler (cont.)

```
# handle timer interrupt by incrementing time comparator
la a1, mtimecmp          # a1 = &time comparator
lw a2, 0(a1)             # load lower 32 bits of comparator
lw a3, 4(a1)             # load upper 32 bits of comparator
addi a4, a2, 1000        # increment lower bits by 1000 cycles
sltu a2, a4, a2          # generate carry-out
add a3, a3, a2           # increment upper bits
sw a3, 4(a1)             # store upper 32 bits
sw a4, 0(a1)             # store lower 32 bits

# restore registers and return
lw a4, 12(a0)            # restore a4
lw a3, 4(a0)             # restore a3
lw a2, 4(a0)            # restore a2
lw a1, 0(a0)            # restore a1
csrrw a0, mscratch, a0   # restore a0; mscratch = &temp storage
mret                    # return from handler
```





How to back to the exception breakpoint ?

- mret
 - $PC \leftarrow CSRs[mepc]$
 - $mstatus.MIE \leftarrow mstatus.MPIE$
 - $mp \leftarrow mstatus.MPP$

| | | | | | | |
|---------|-------|-------|-------|-------|---------|---|
| 31 | 25 24 | 20 19 | 15 14 | 12 11 | 7 6 | 0 |
| 0011000 | 00010 | 00000 | 000 | 00000 | 1110011 | |





浙江大学
ZheJiang University

END

